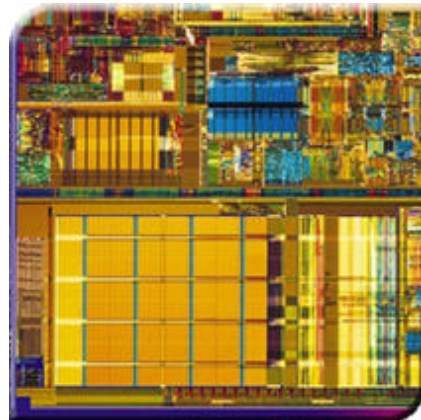
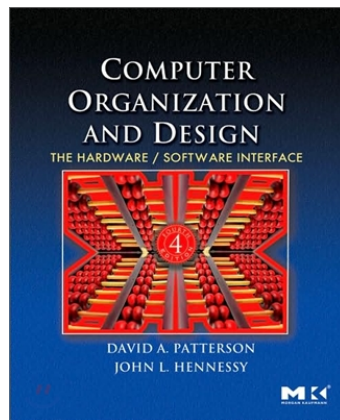


Computer Architecture

Lecture 1 Introduction

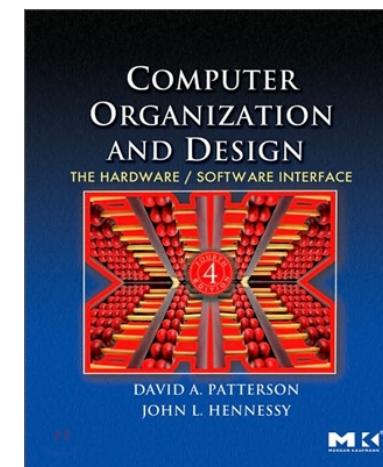


Prof. Jongmyon Kim

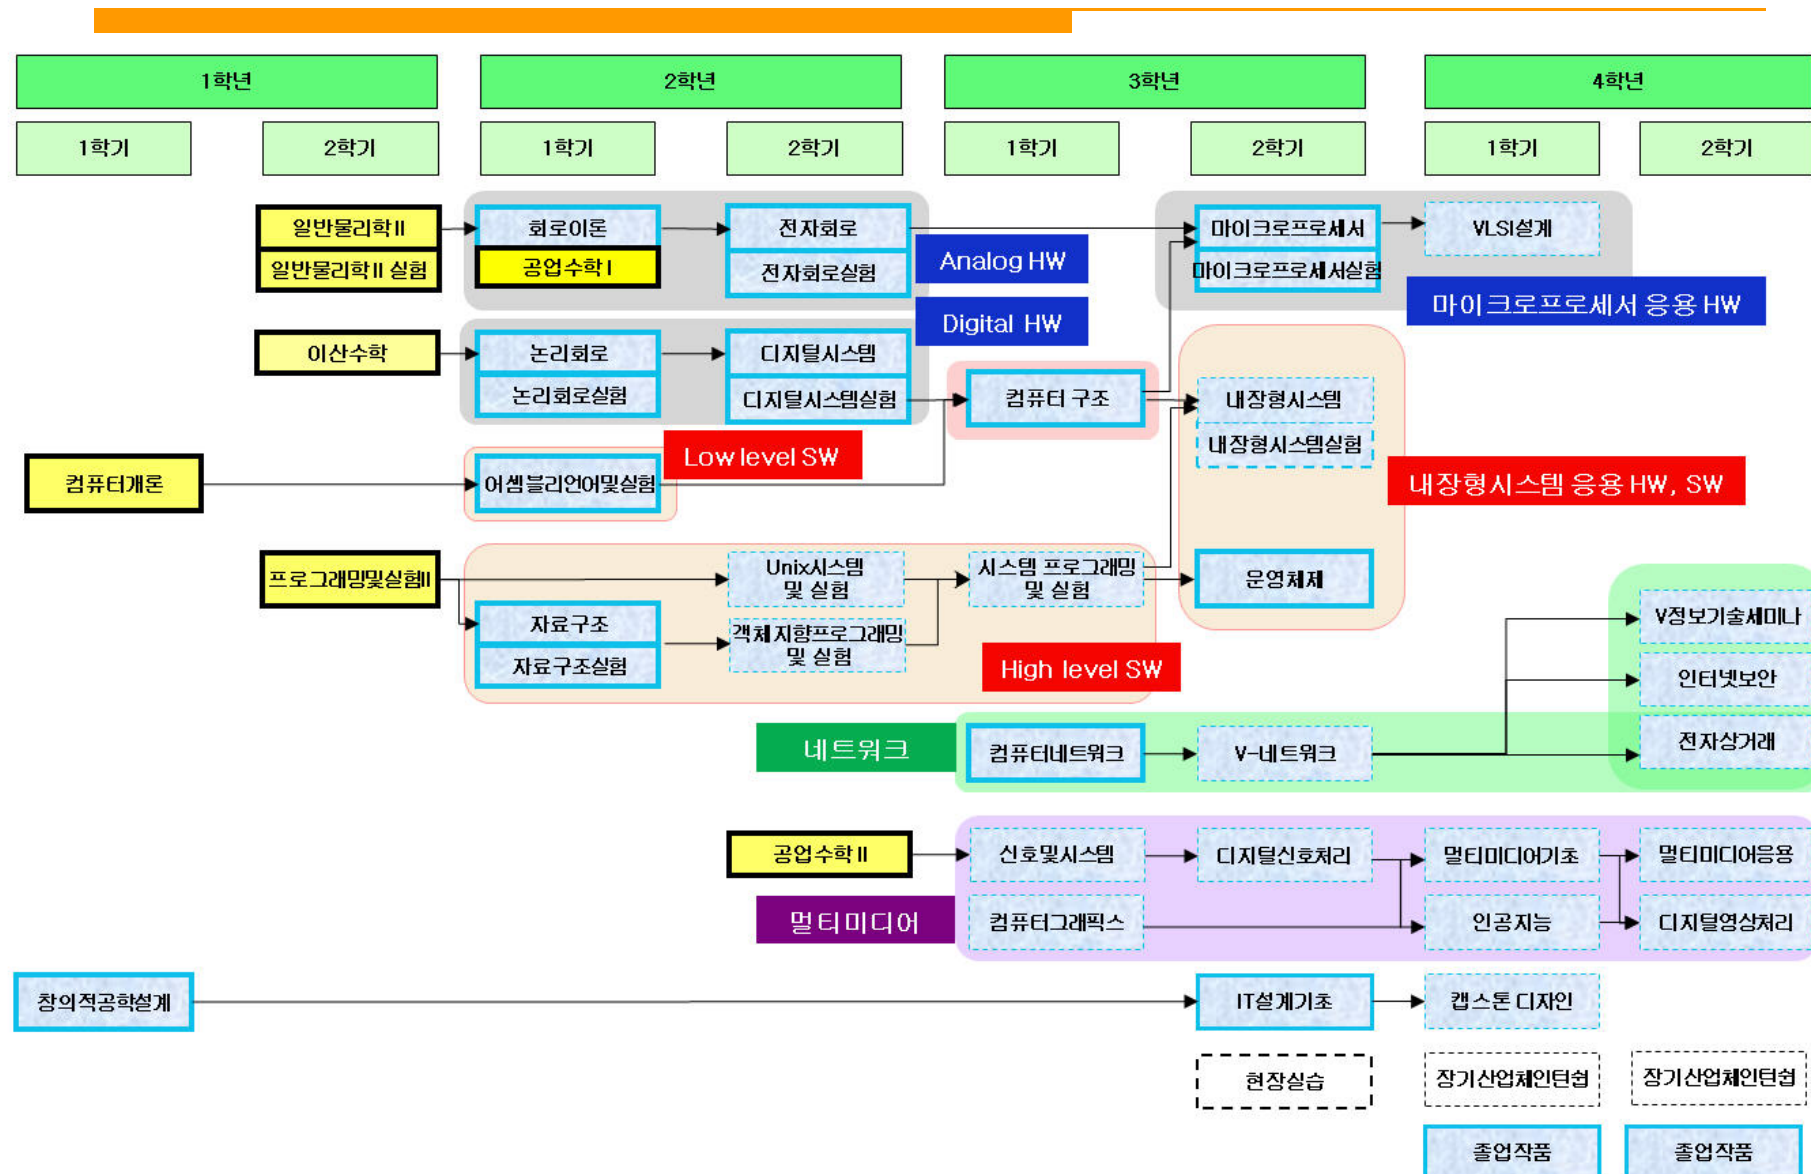


Course Information

- ☐ Instructor
 - ☐ Prof. Jong-Myon Kim (jongmyon.kim@gmail.com)
- ☐ Class page:
 - ☐ <http://uclass.ulsan.ac.kr/>
 - ☐ Constantly updated, check it out regularly
- ☐ Prerequisite: Logic circuit or the equivalent
- ☐ Meeting time: Monday 2,3, Wednesday 3
- ☐ Meeting place : 7-613
- ☐ Textbook
 - ☐ Hennessy and Patterson, *Computer Organization and Design* (4th edition), Morgan Kaufmann, 2008.
- ☐ Other teaching materials
 - ☐ Slides & Lectures



Text Book



Weekly Plan & Reading

- Week 1: Introduction and Instruction Set Architecture
 - Week 2: MIPS Instructions, Software Conversion, Stack
 - Week 3: Arithmetic for Computers
 - Week 4: Arithmetic for Computers
 - Week 5: CPU Performance
 - Week 6: Single-Cycle Datapath Processor Design
 - Week 7: Single-Cycle Datapath Processor Design & Term project 1 (tentative)
 - Week 8: *Mid-Term Exam*
 - Week 9: Multi-Cycle Datapath Processor Design
 - Week 10: Multi-Cycle Datapath Implementation – Exception & Microprogramming
 - Week 11: Pipelining Processor : Overview & Datapath
 - Week 12: Pipelining Processor : Control, Hazard, Exceptions
 - Week 13: Memory Hierarchy : Caches
 - Week 14: Memory Hierarchy : Cache performance
 - Week 15: Storage, Networks, and Other Peripherals
 - Week 16: *Final Exam*
-

Grading Policy

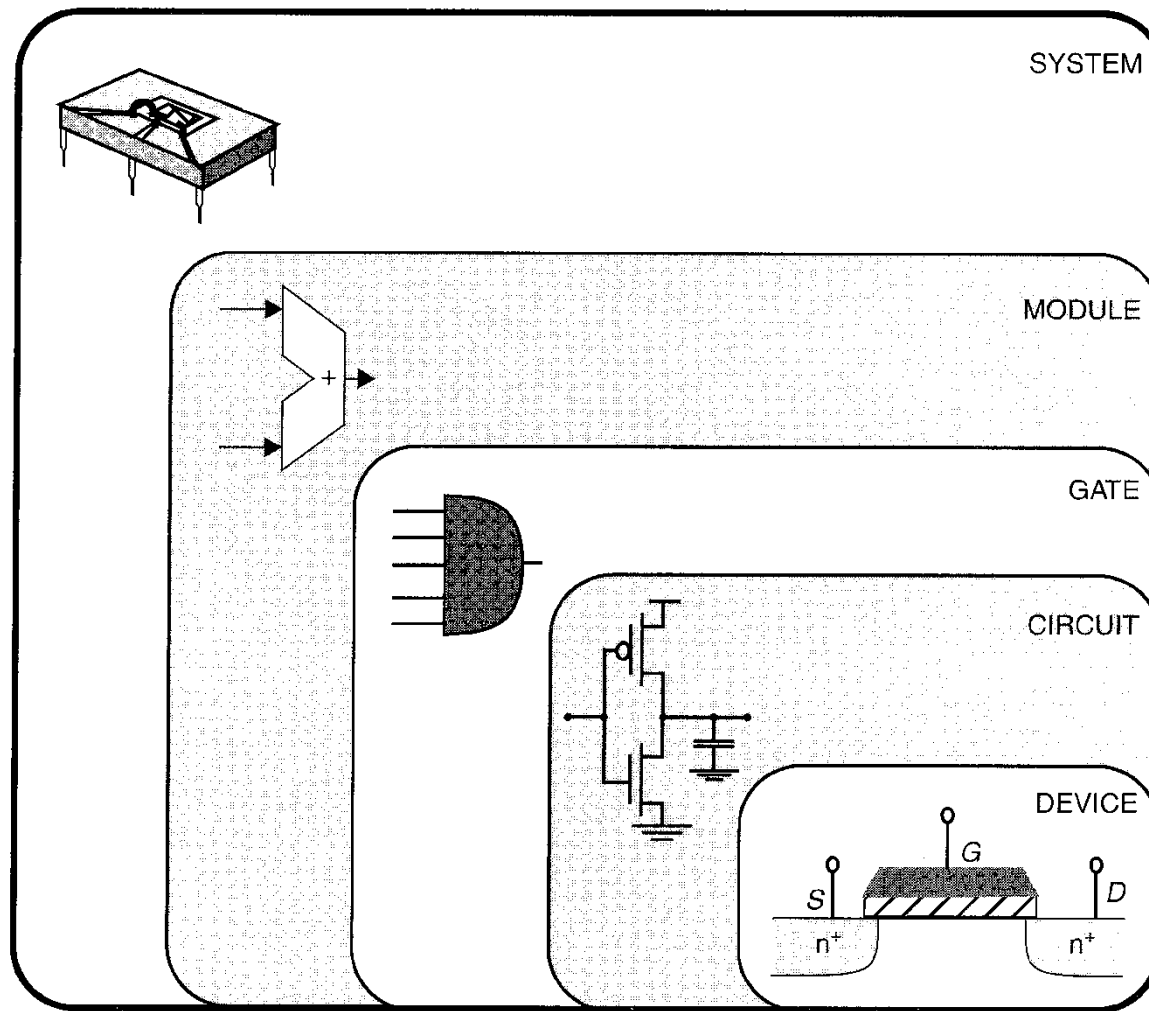
- Term Projects: 20%
 - Individual work, no collaboration unless otherwise announced

- Exams: 70%
 - Mid-Term Exam: 30%
 - Final Exam: 40%

- Class Participation: 10%

- Final Grade is relative to your peer in class

Design Steps



[Jan Rabaey's Digital Circuit Design 중에서]

Objectives — To Learn

- Core concepts of microprocessor architecture
 - ISA
 - Pipelining
 - Hazards
 - Cache/Memory hierarchy
 - Memory management

Introduction

- This course is all about how computers work

- But what do we mean by a computer?
 - Different types: desktop, servers, embedded devices
 - Different uses: automobiles, graphics, finance, genomics...
 - Different manufacturers: Intel, Apple, IBM, Microsoft, Sun...
 - Different underlying technologies and different costs!

- Best way to learn:
 - Focus on a specific instance and learn how it works
 - While learning general principles and historical perspectives

Why learn this stuff?

- You want to call yourself a “computer scientist”
- You want to build software people use (need performance)
- You need to make a purchasing decision or offer “expert” advice
- Both Hardware and Software affect performance:
 - Algorithm determines number of source-level statements
 - Language/Compiler/Architecture determine machine instructions
 - Processor/Memory determine how fast instructions are executed
- Assessing and Understanding Performance

Moore's Law

IBM latest POWER5 has
276 million transistors

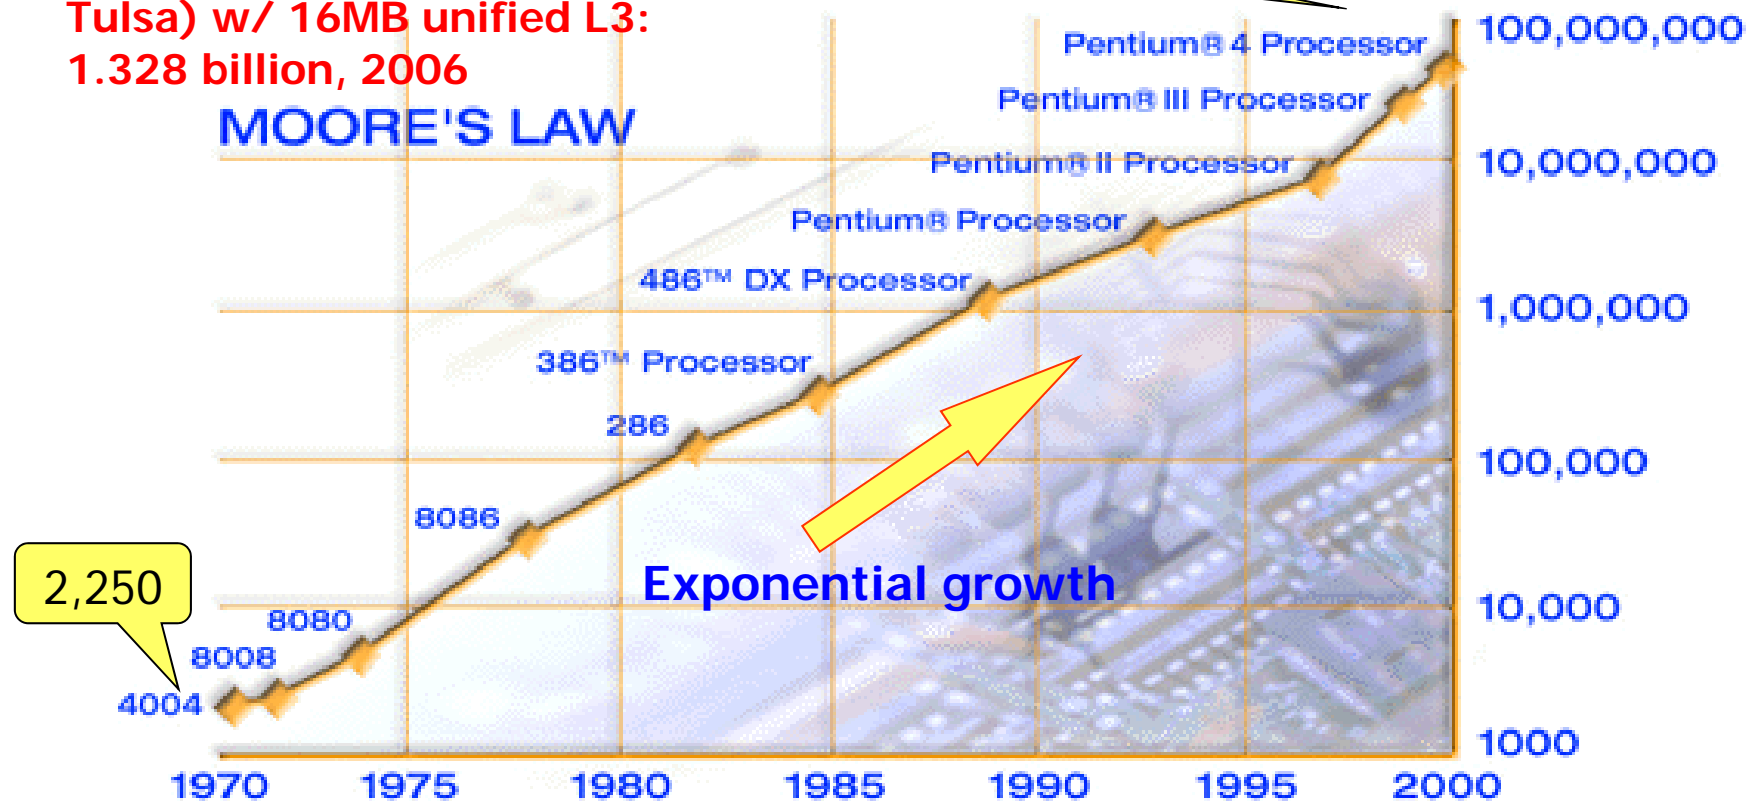
Intel Dual-Core Xeon (P4-based
Tulsa) w/ 16MB unified L3:
1.328 billion, 2006

P4 Extreme Ed.
178 millions w/ 2MB L3

Core 2 Duo (Conroe)
291 millions, July
2006

42 millions

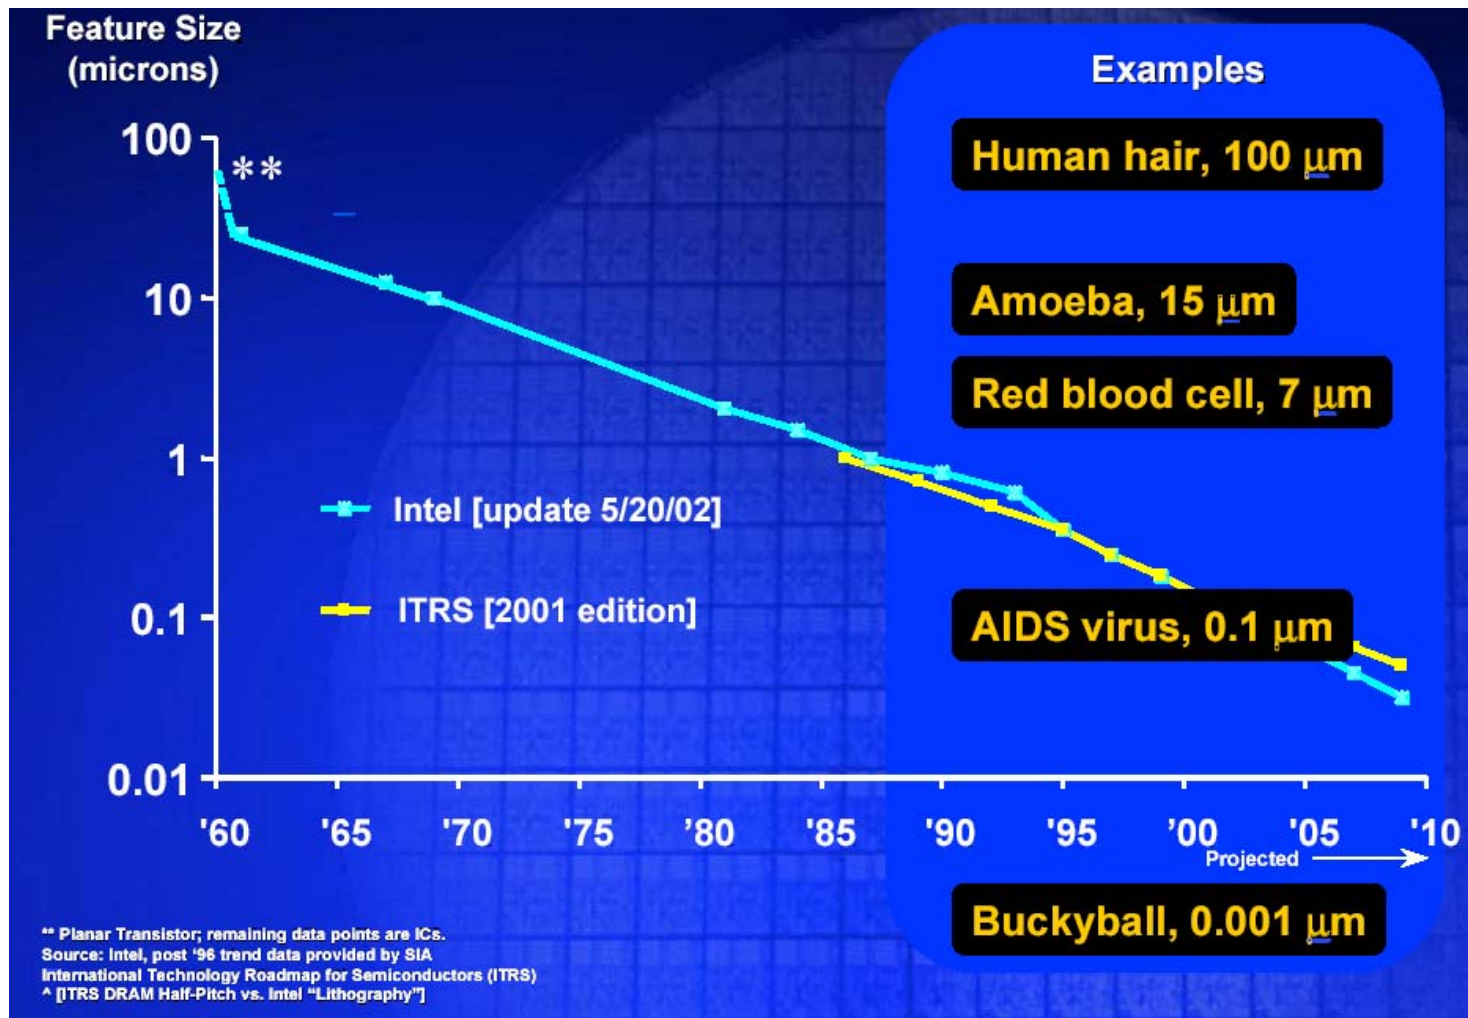
transistors



Transistor count will be doubled every 18 months

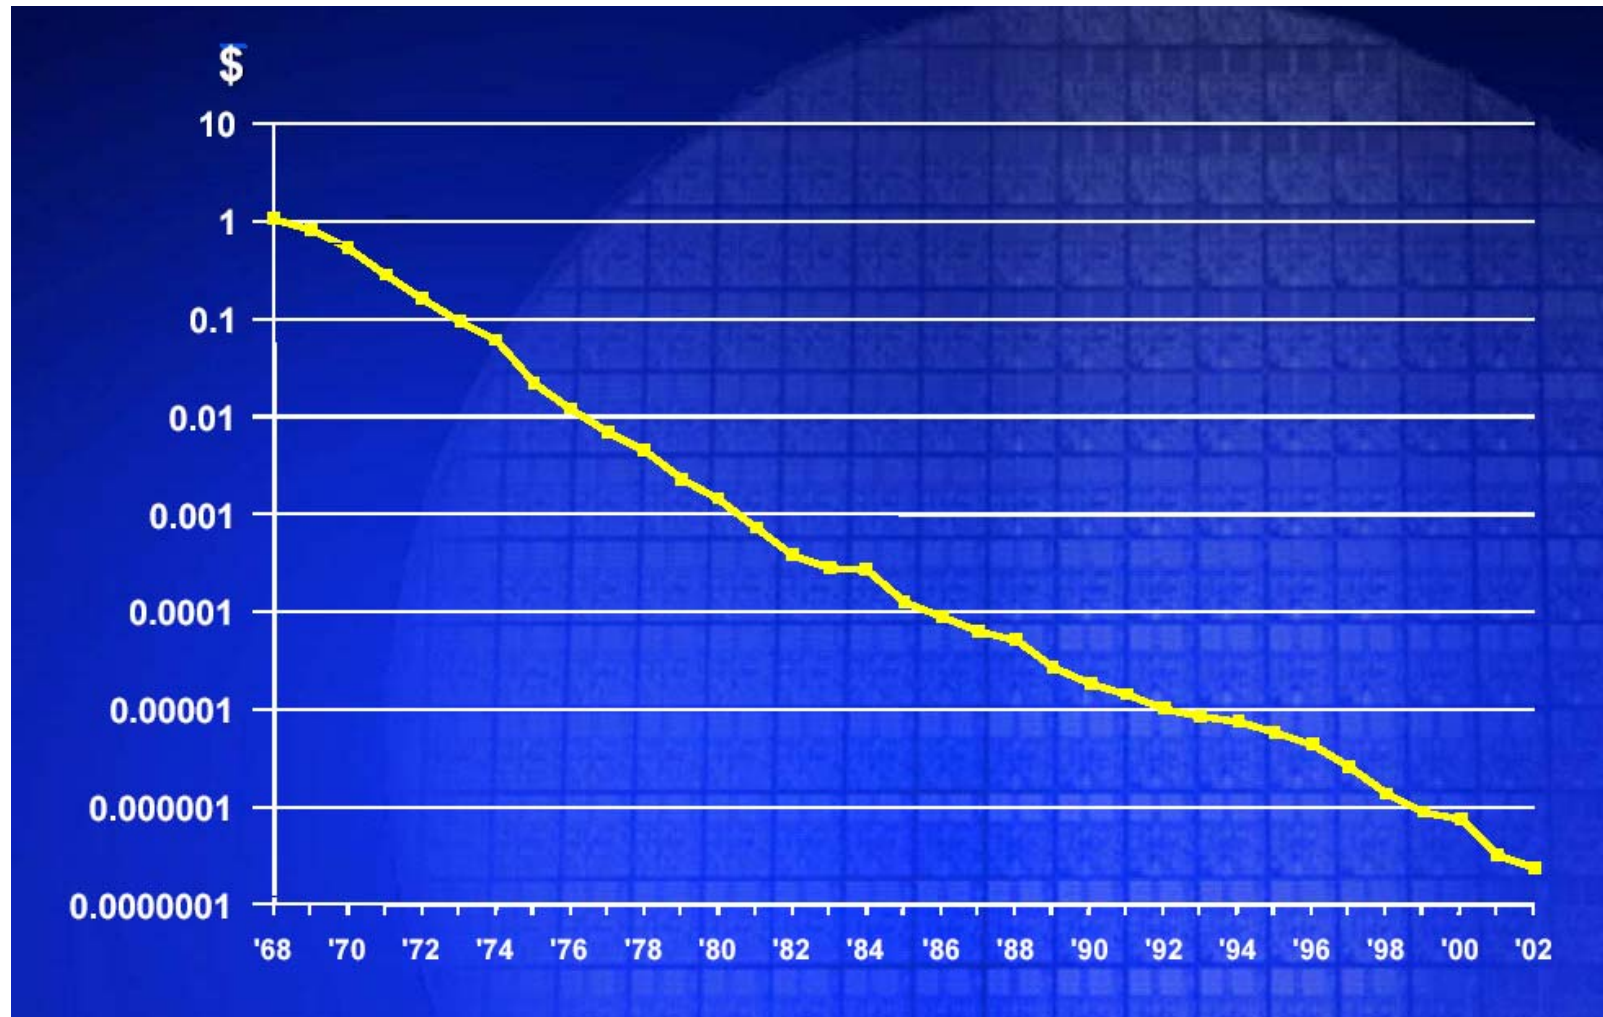
— Gordon Moore, Intel co-founder

Feature Size



We are currently at 0.09 μm and moving towards 0.025 μm

Average Transistor Cost Per Year



What is a computer?

- Components:
 - Processor(s)
 - Co-processors (graphics, security)
 - Memory (disk drives, DRAM, SRAM, CD/DVD)
 - input (mouse, keyboard, mic)
 - output (display, printer)
 - network
- Our primary focus: the processor (datapath and control)
 - implemented using millions of transistors
 - Impossible to understand by looking at each transistor
 - We need...

Abstraction

High Level
Language

```
main() {
    int i,b,c,a[10];
    for (i=0; i<10; i++)...
        a[2] = b + c*i;
}
```

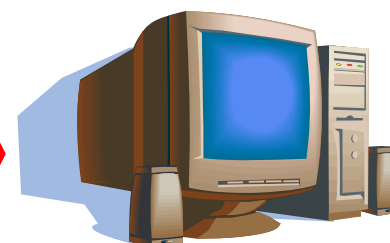
Compiler

ISA

```
...
lw  r2, mem[r7]
add r3, r4, r2
st  r3, mem[r8]
```

Assembler

- Delving into the depths reveals more information
- An abstraction omits unneeded detail, helps us cope with complexity
- *What are some of the details that appear in these familiar abstractions?*

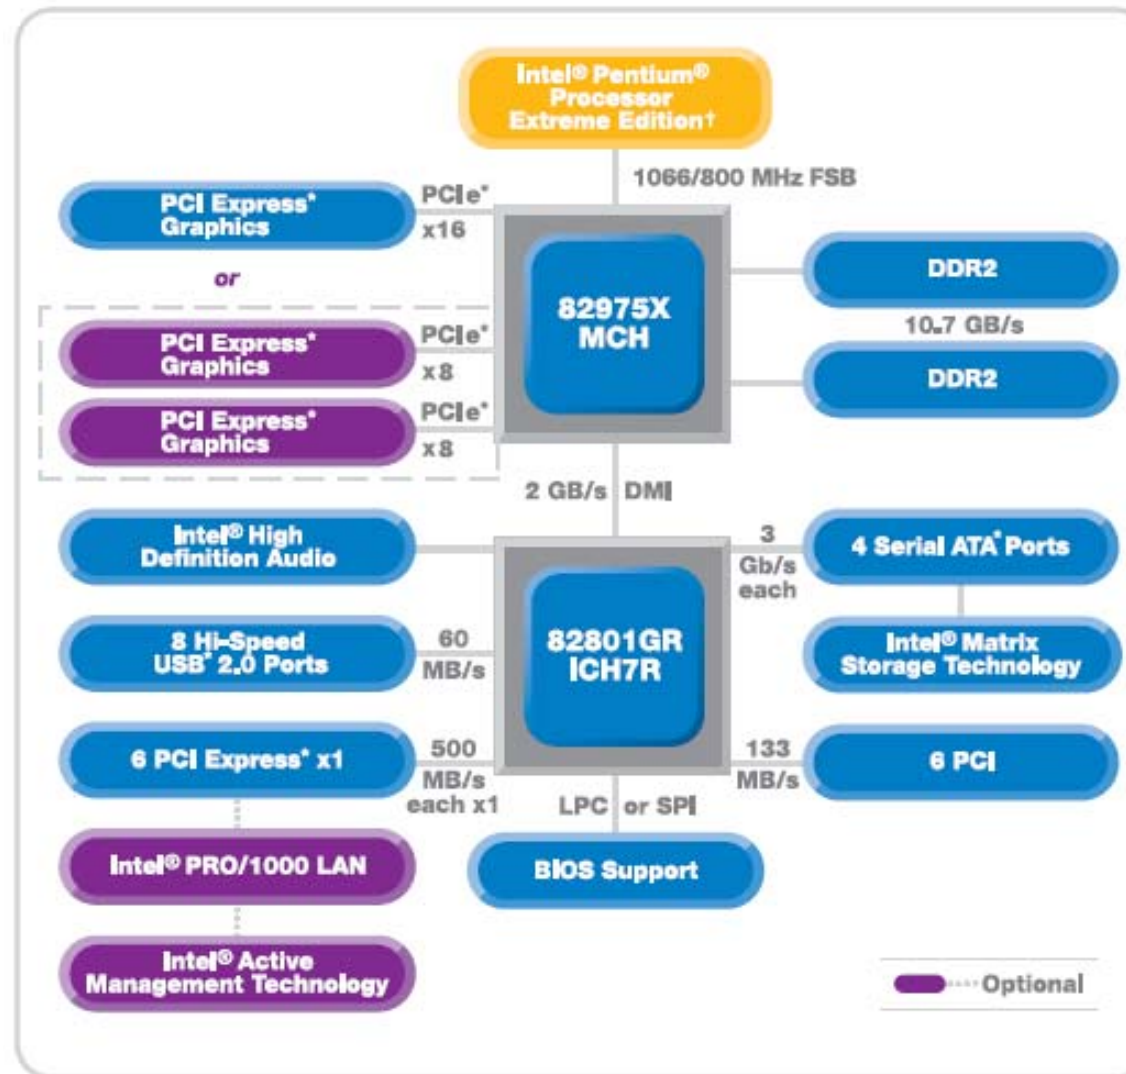


How do computers work?

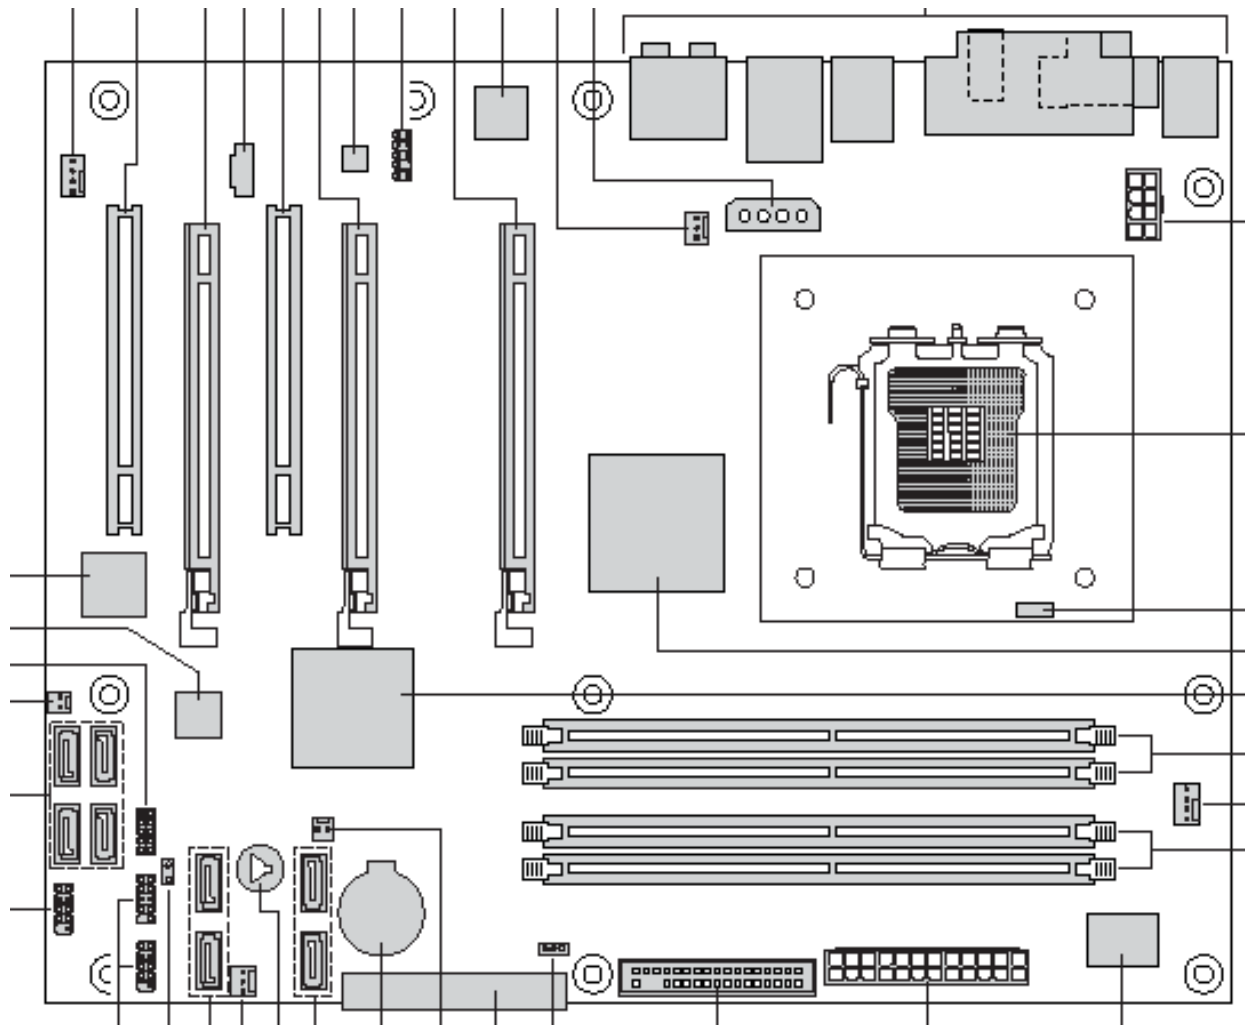
- Need to understand abstractions such as:
 - Applications software
 - Systems software
 - Assembly Language
 - Machine Language
 - Architectural Issues: i.e., Caches, Virtual Memory, Pipelining
 - Sequential logic, finite state machines
 - Combinational logic, arithmetic circuits
 - Boolean logic, 1s and 0s
 - Transistors used to build logic gates (CMOS)
 - Semiconductors/Silicon used to build transistors
 - Properties of atoms, electrons, and quantum dynamics

 - So much to learn!
-

A Typical PC System Architecture



A Typical PC Motherboard (D975XBX)



A Typical PC Motherboard (D975XBX)



Instruction Set Architecture

- A very important abstraction
 - interface between hardware and low-level software
 - standardizes instructions, machine language bit patterns, etc.
 - advantage: *different implementations of the same architecture*
 - disadvantage: *sometimes prevents using new innovations*

 - Modern instruction set architectures:
 - IA-32, PowerPC, MIPS, SPARC, ARM, and others
-

Historical Perspective

- ENIAC built in World War II was the first general purpose computer
 - Used for computing artillery firing tables
 - 80 feet long by 8.5 feet high and several feet wide
 - Each of the twenty 10 digit registers was 2 feet long
 - Used 18,000 vacuum tubes
 - Performed 1900 additions per second



–Since then:

Moore's Law:

**transistor capacity doubles
every 18-24 months**

Where we are headed

- A specific instruction set architecture
- Performance issues
- Arithmetic and how to build an ALU
- Constructing a processor to execute our instructions
- Pipelining to improve performance
- Memory: caches and virtual memory
- I/O