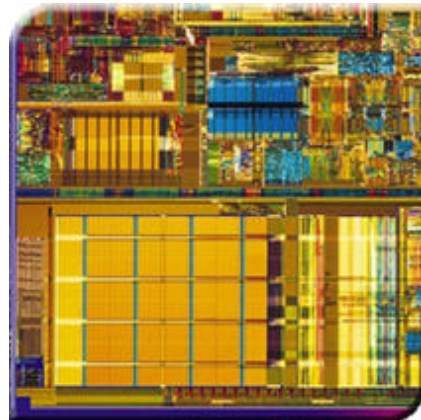



Computer Architecture

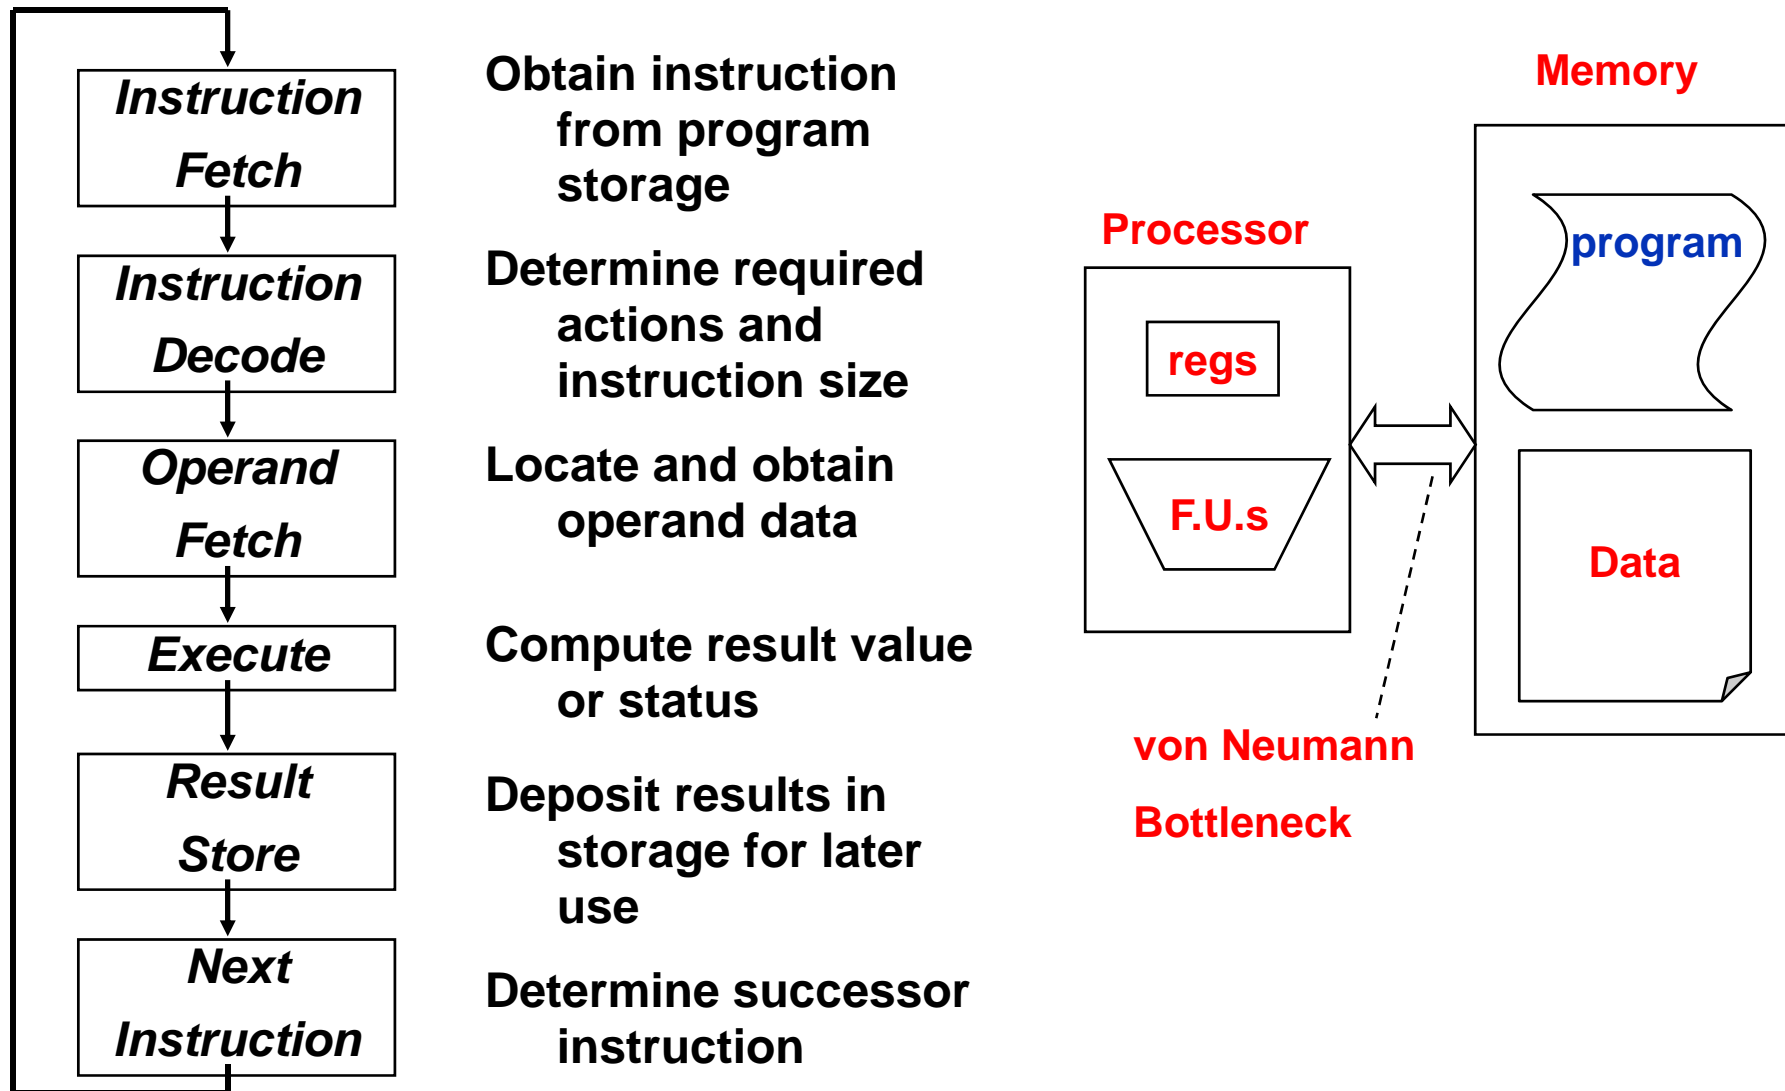
Lecture 7 Pipelining



Prof. Jongmyon Kim



Fundamental Execution Cycle





Why Pipelining?

Instruction Class	Instruction Fetch	Register Read	ALU Operation	Data Access	Register Write	Total Time
Load word	2ns	1ns	2ns	2ns	1ns	8ns
Store word	2ns	1ns	2ns	2ns		7ns
R-format	2ns	1ns	2ns		1ns	6ns
Branch	2ns	1ns	2ns			5ns

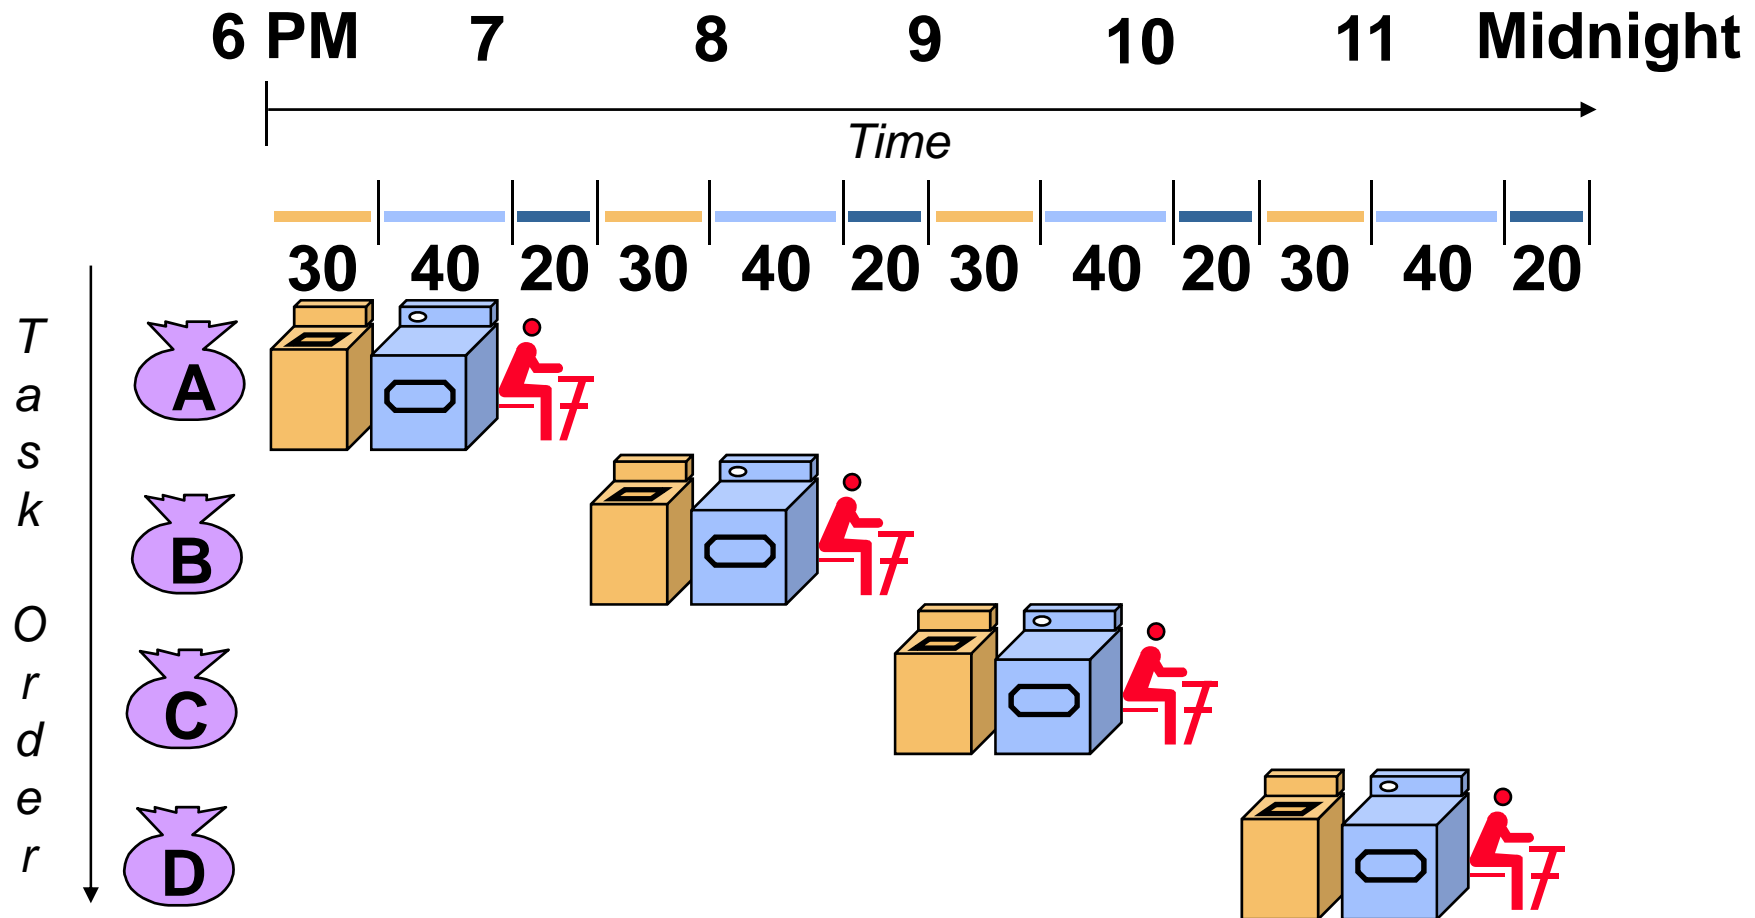
- Single cycle datapath
 - Design for the worst case
 - Need to make the cycle time = 8ns per cycle
- Multi-cycle datapath
 - Design for each individual instruction class
 - For the above example: cycle time = 2ns
 - Lw=10ns (5 cycles), sw=8ns (4 cycles), R-format=8ns(4 cycles), beq=6ns (3 cycles)
- Can we do better?

Pipelining: It's Natural!

- Laundry Example
- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold
- Washer takes 30 minutes
- Dryer takes 40 minutes
- “Folder” takes 20 minutes

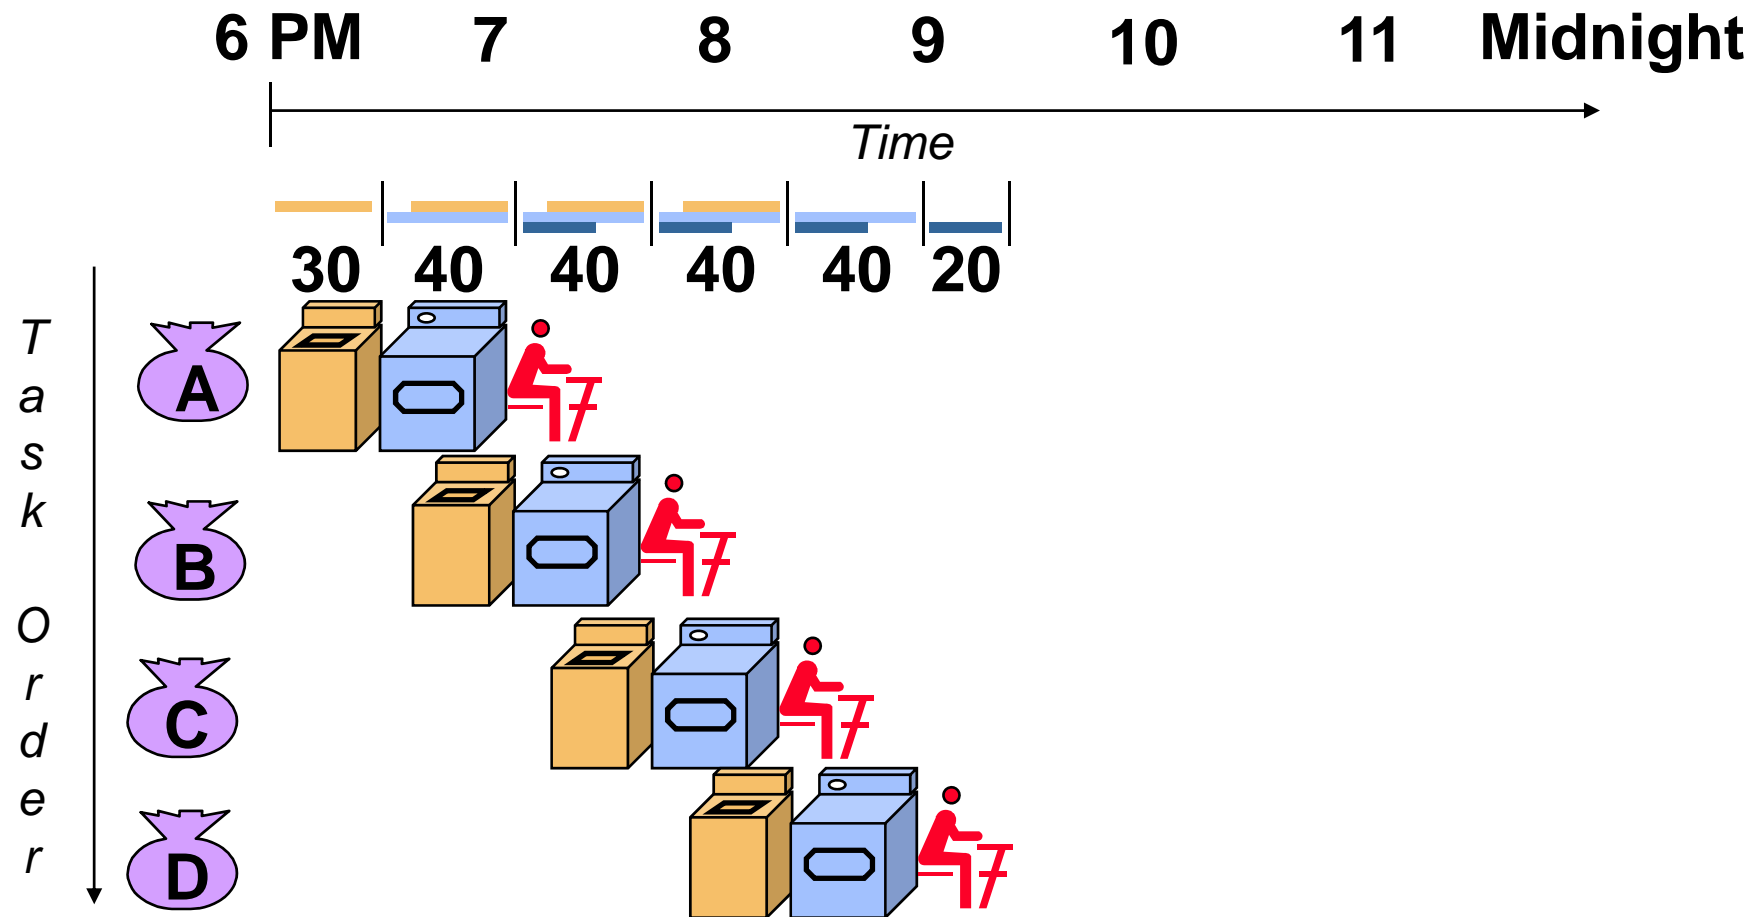


Sequential Laundry



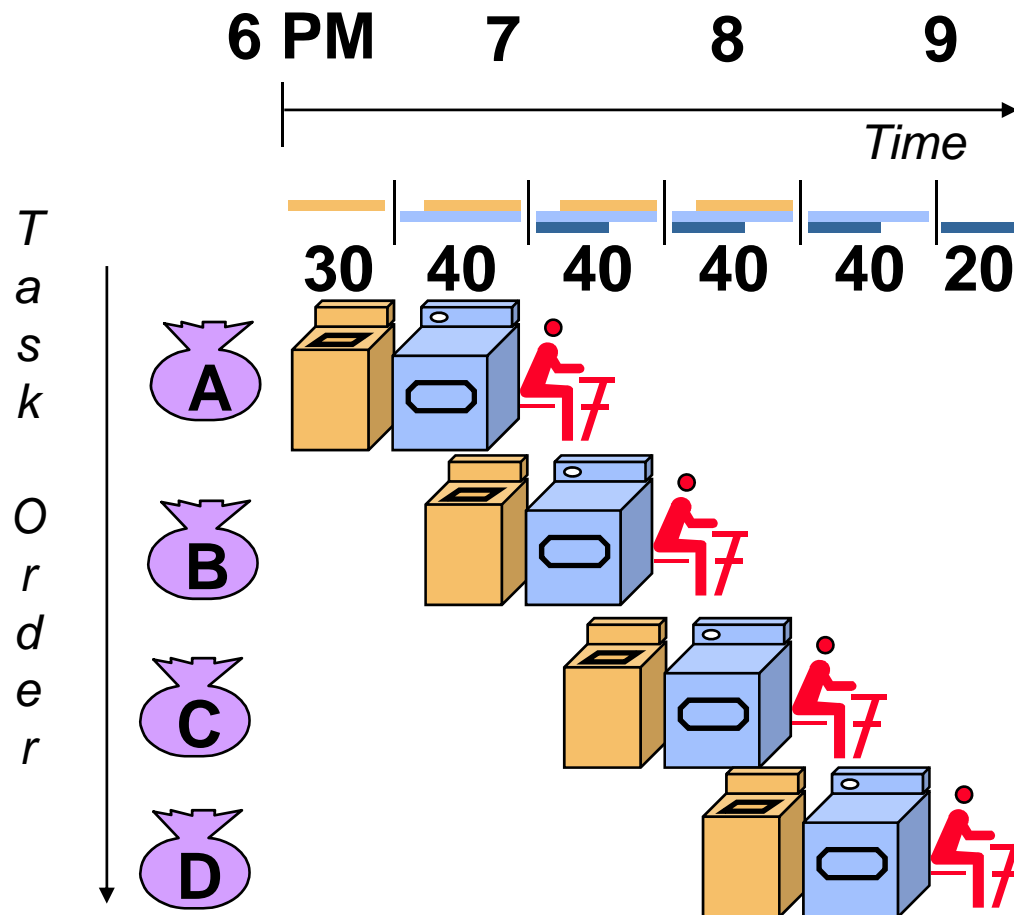
- Sequential laundry takes 6 hours for 4 loads
- If they learned pipelining, how long would laundry take?

Pipelined Laundry: Start work ASAP



□ Pipelined laundry takes 3.5 hours for 4 loads

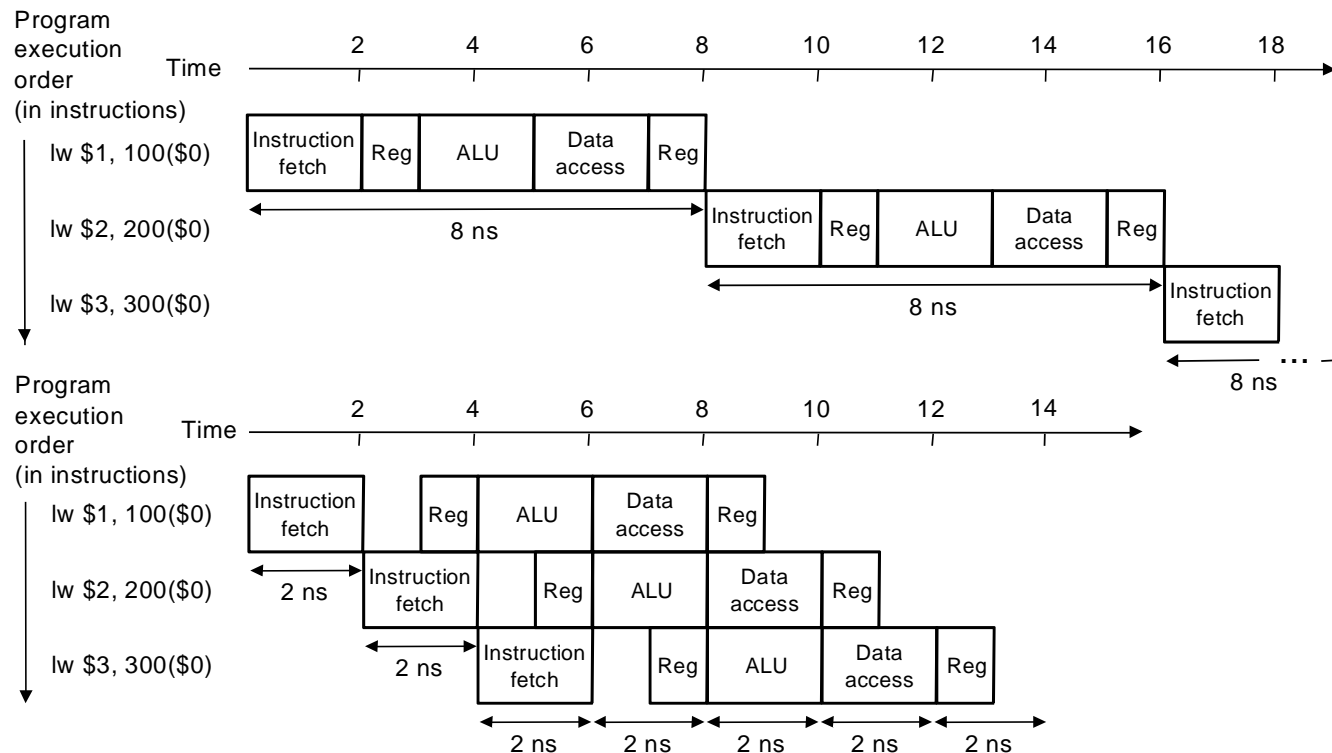
Pipelining Lessons



- Pipelining doesn't help **latency** of a single task, it helps **throughput** of entire workload
- Pipeline rate limited by **slowest** pipeline stage
- **Multiple** tasks operating simultaneously
- Potential speedup = **Number pipe stages**
- Unbalanced lengths of pipe stages reduces speedup
- Time to "**fill**" pipeline and time to "**drain**" it reduces speedup

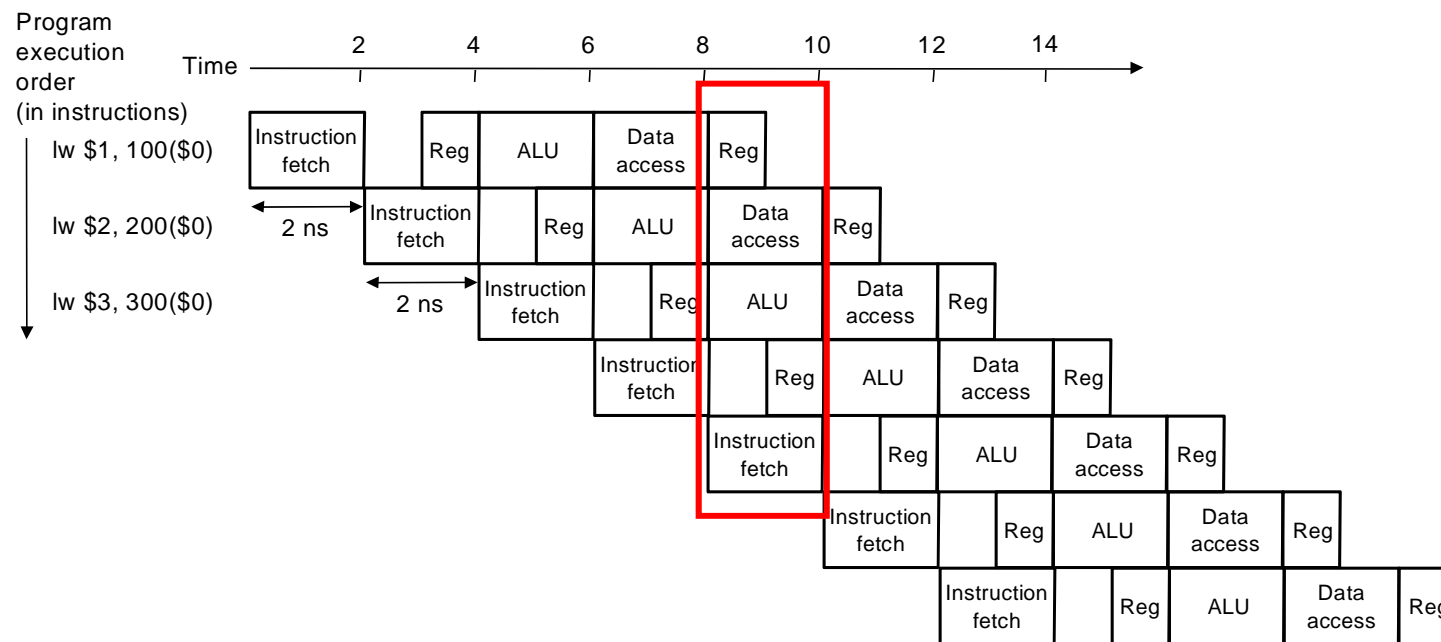
Pipelining

- Improve performance by increasing instruction throughput



Ideal speedup is number of stages in the pipeline. Do we achieve this?

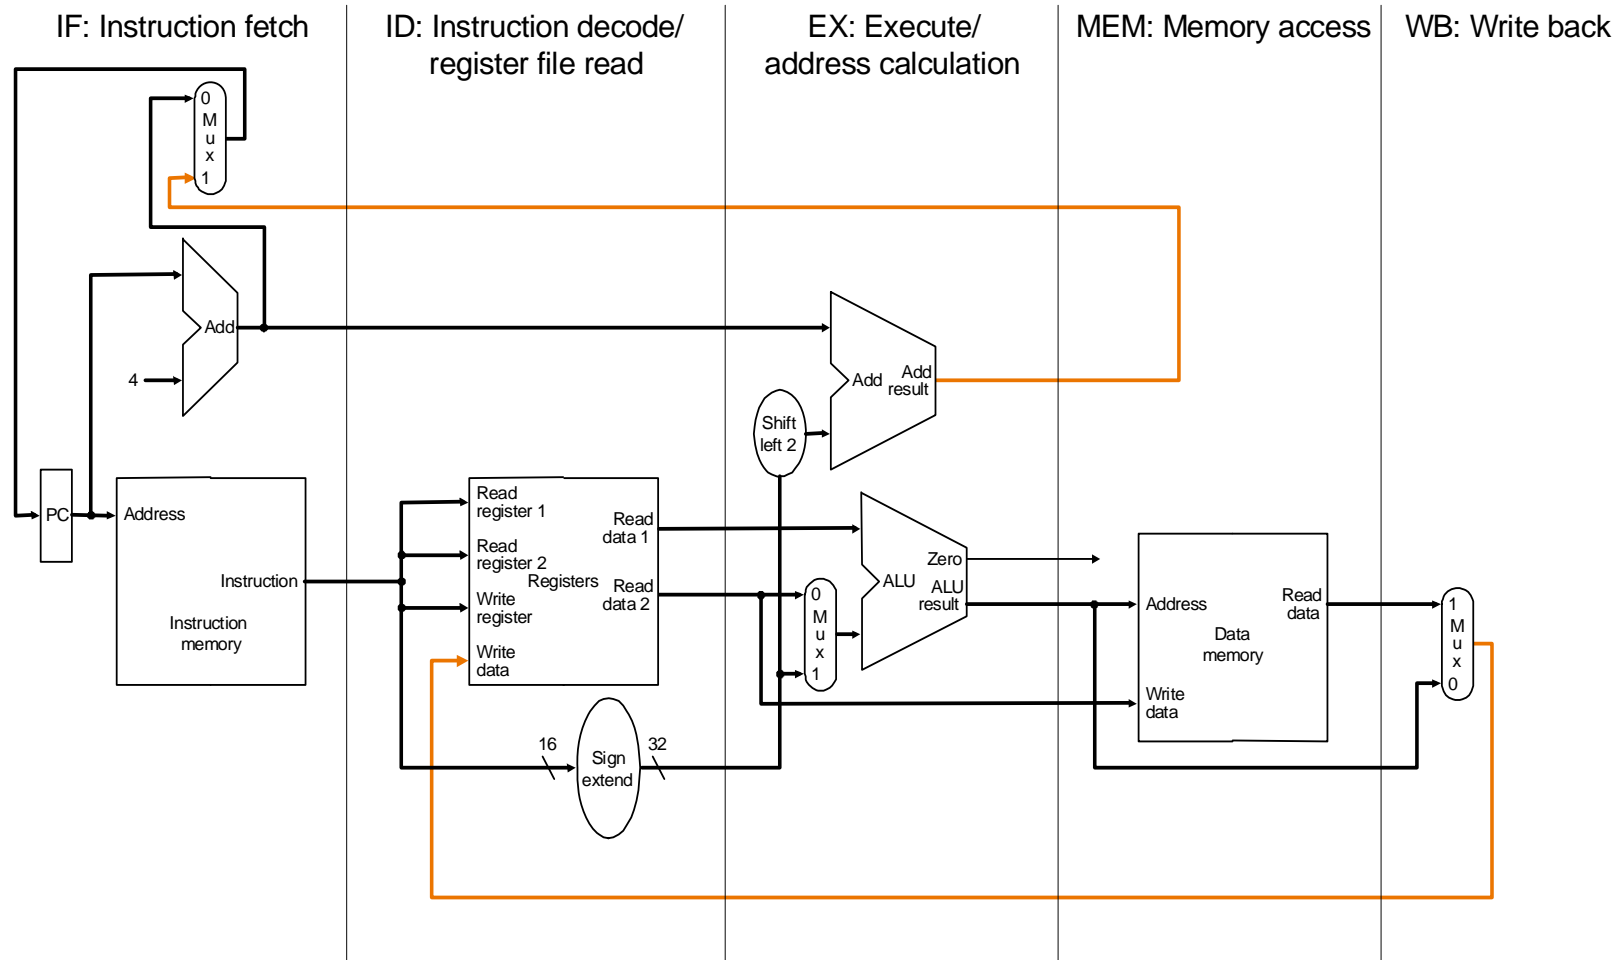
Ideal Pipelining



Pipelining

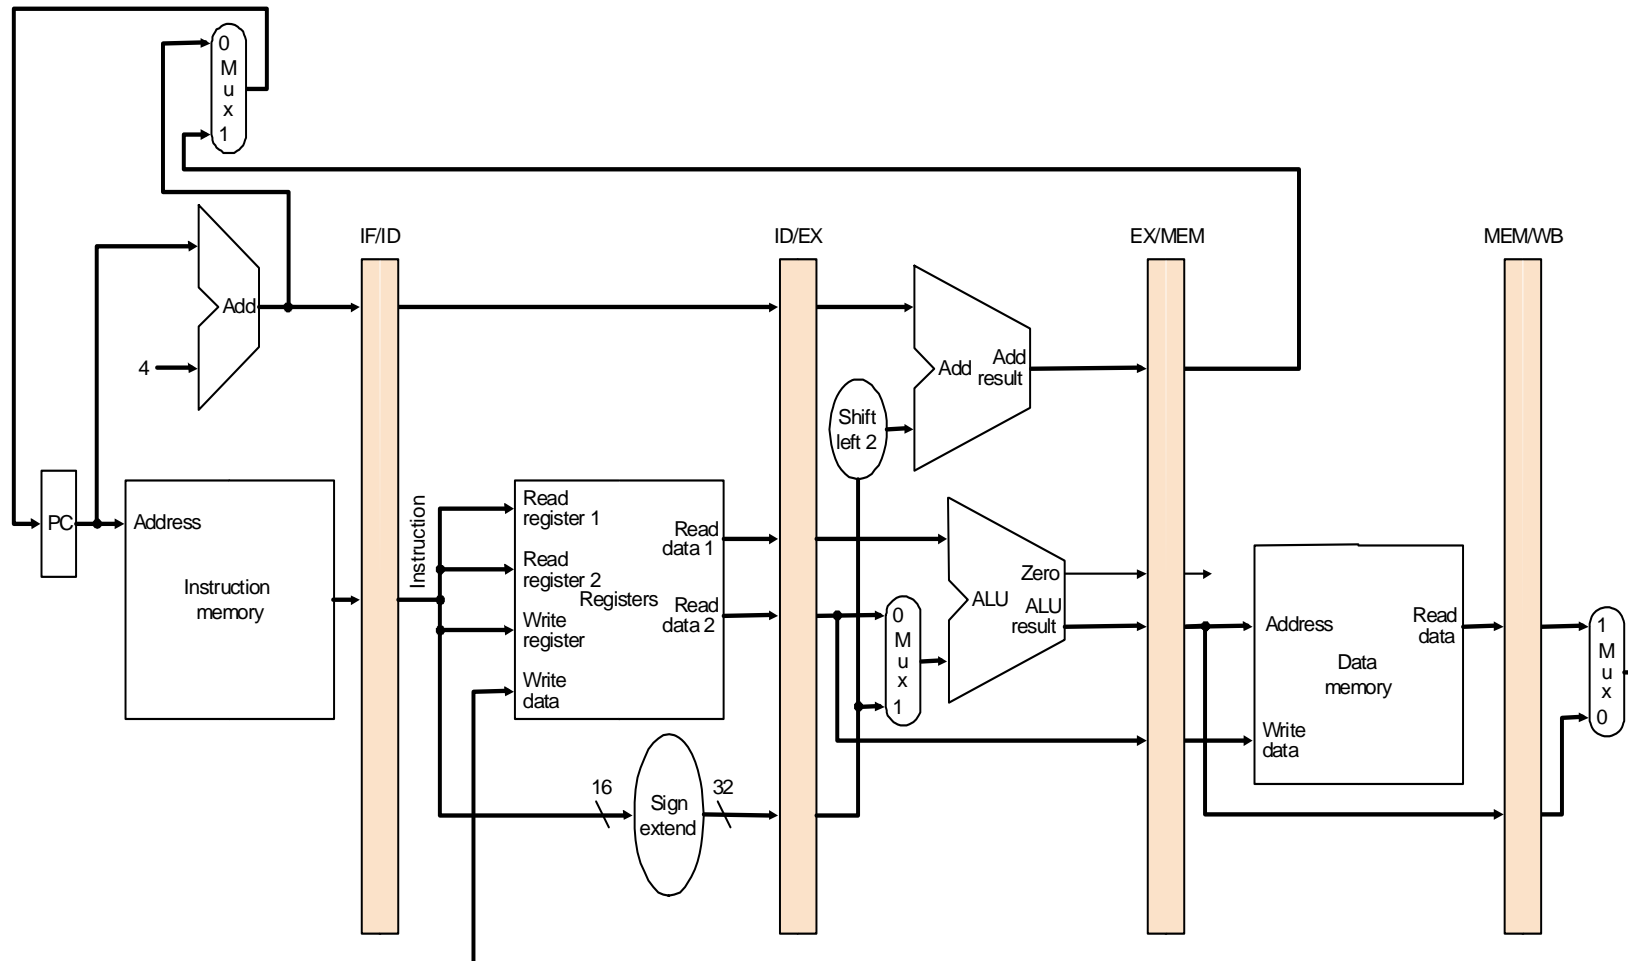
- What makes it easy
 - All instructions are the same length
 - Simple instruction formats
 - Memory operands appear only in loads and stores
- What makes it hard?
 - structural hazards: suppose we had only one memory
 - control hazards: need to worry about branch instructions
 - data hazards: an instruction depends on a previous instruction
- We'll build a simple pipeline and look at these issues
- We'll talk about modern processors and what really makes it hard:
 - exception handling
 - trying to improve performance with out-of-order execution, etc.

Basic Idea




□ What do we need to add to actually split the datapath into stages?

Pipelined Datapath



Graphically Representing Pipelines

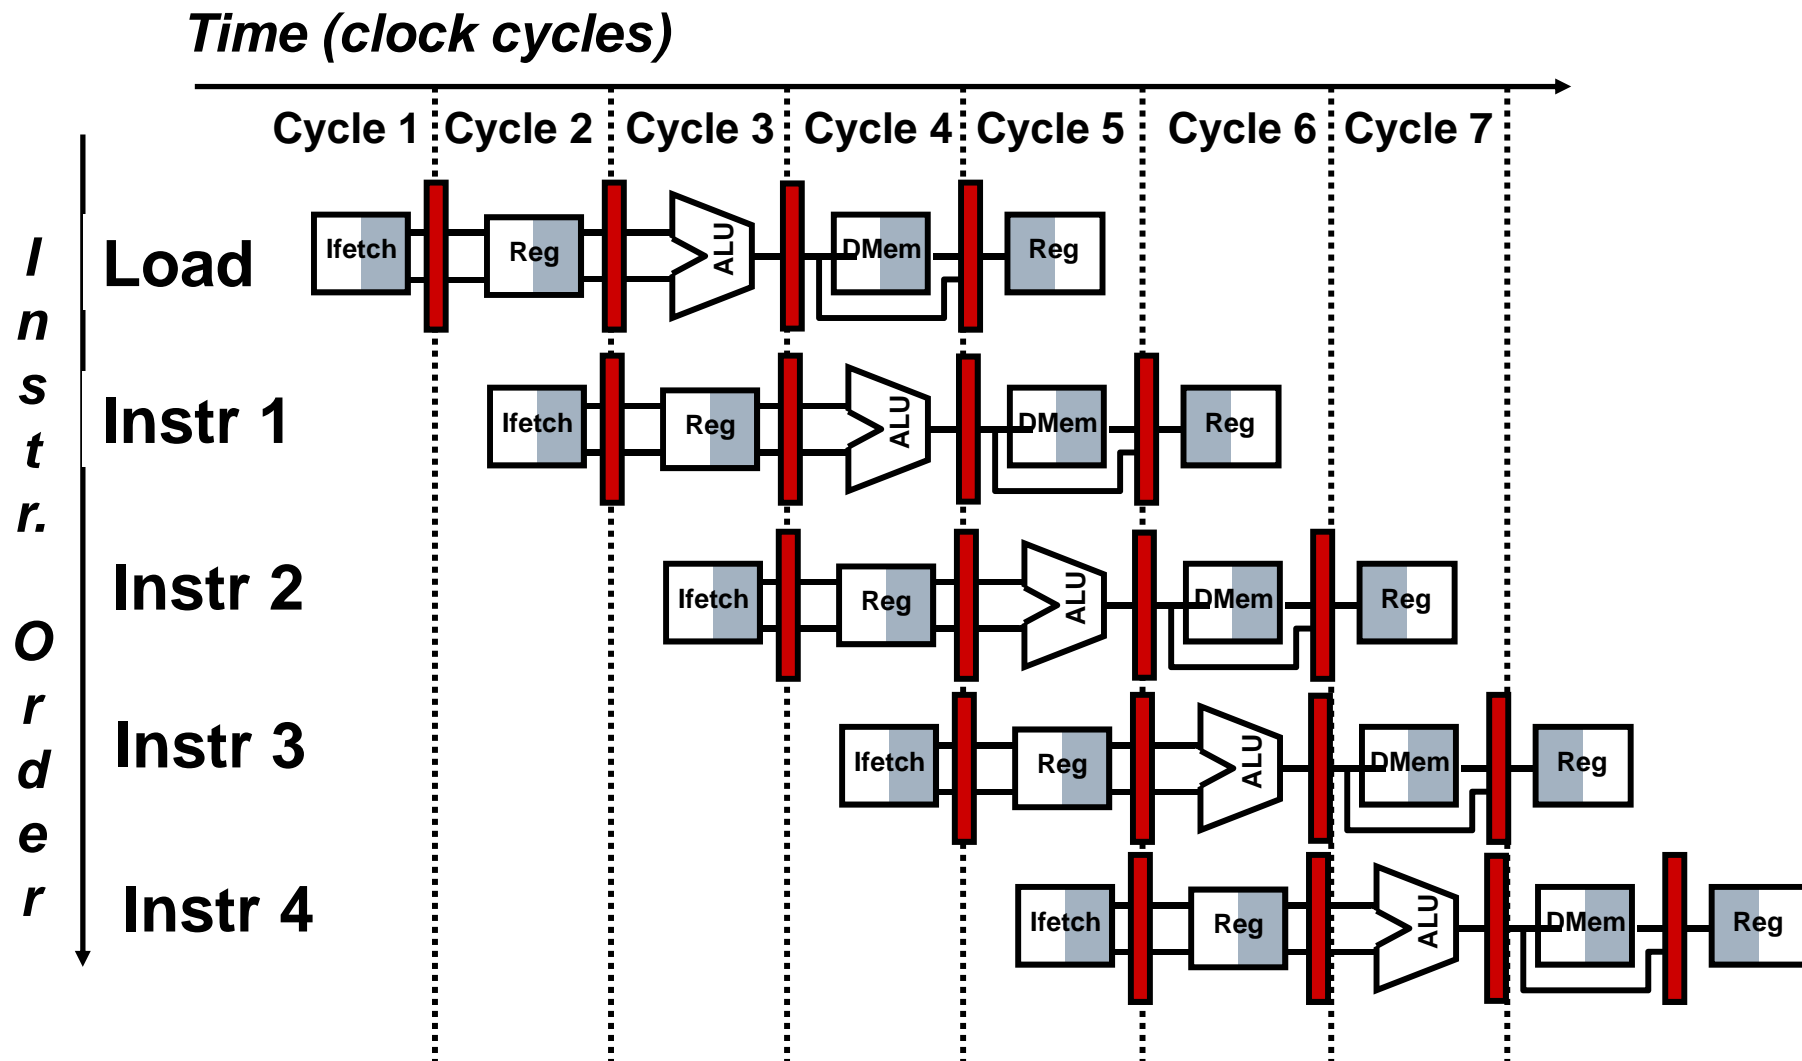


- Shading indicates the unit is being used by the instruction
- Shading on the right half of the register file (ID or WB) or memory means the element is being read in that stage
- Shading on the left half means the element is being written in that stage

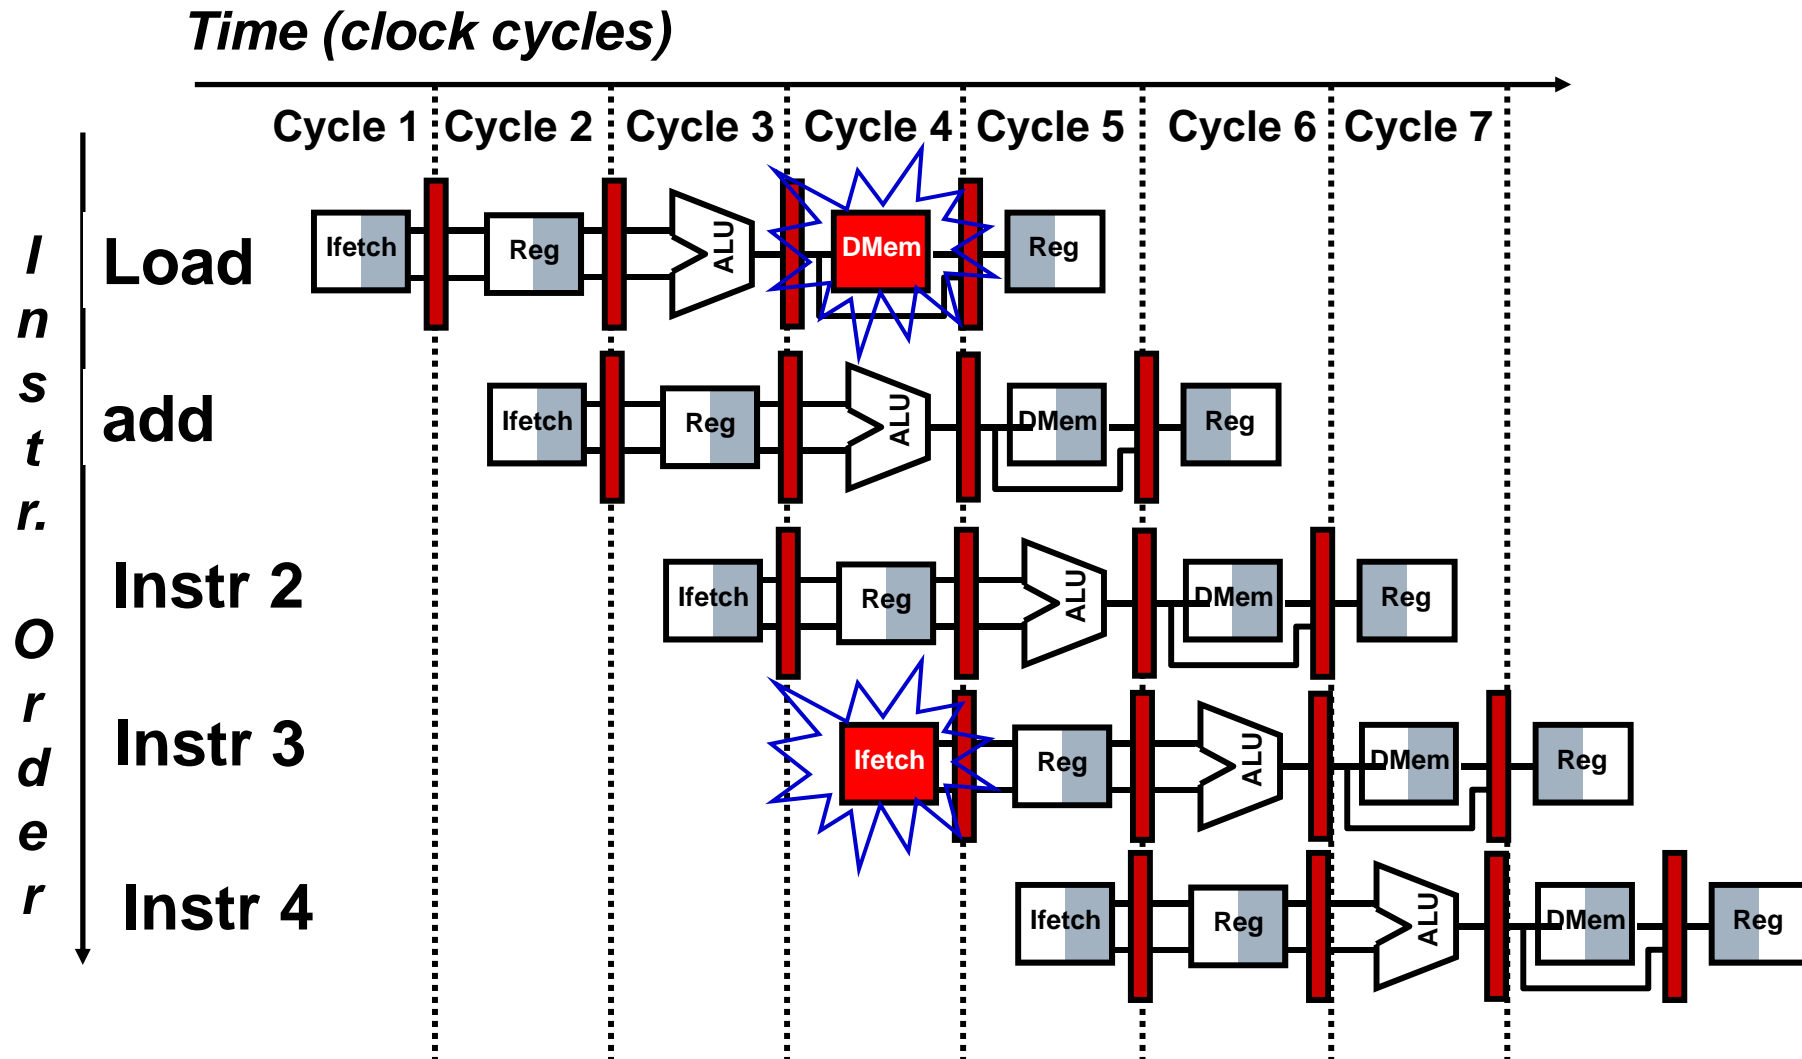
Pipelining is not quite that straightforward !

- Limits to pipelining: **Hazards** prevent next instruction from executing during its designated clock cycle
 - Structural hazards: HW cannot support this combination of instructions
 - Data hazards: Instruction depends on result of prior instruction still in the pipeline
 - Control hazards: Caused by delay between the fetching of instructions and decisions about changes in control flow (branches and jumps).


Single Memory Port



Single Memory Port / Structural Hazard

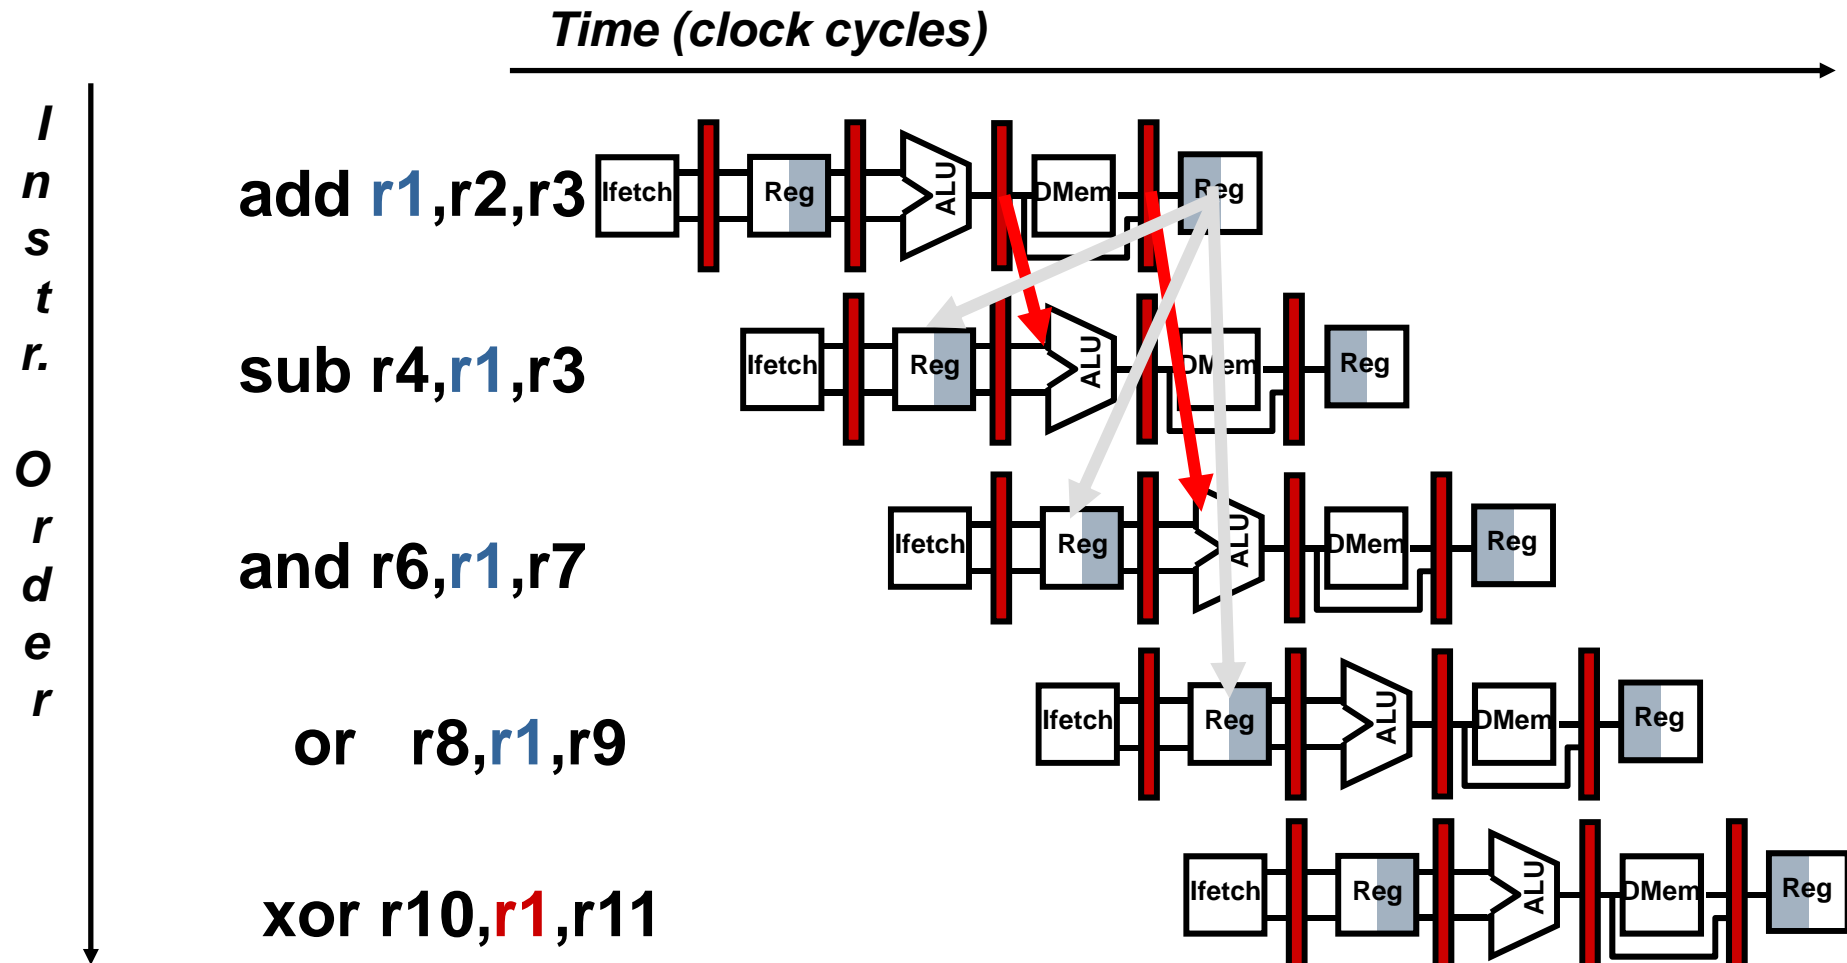


Data Hazard




- ❑ Data is not ready for the subsequent dependent instruction
- ❑ Pipeline stall, typically referred to as “bubble”
- ❑ Performance is penalized
- ❑ One solution — forwarding (or bypassing)


Forwarding to Avoid Data Hazard



Reducing Data Hazard: Forwarding (Bypassing)

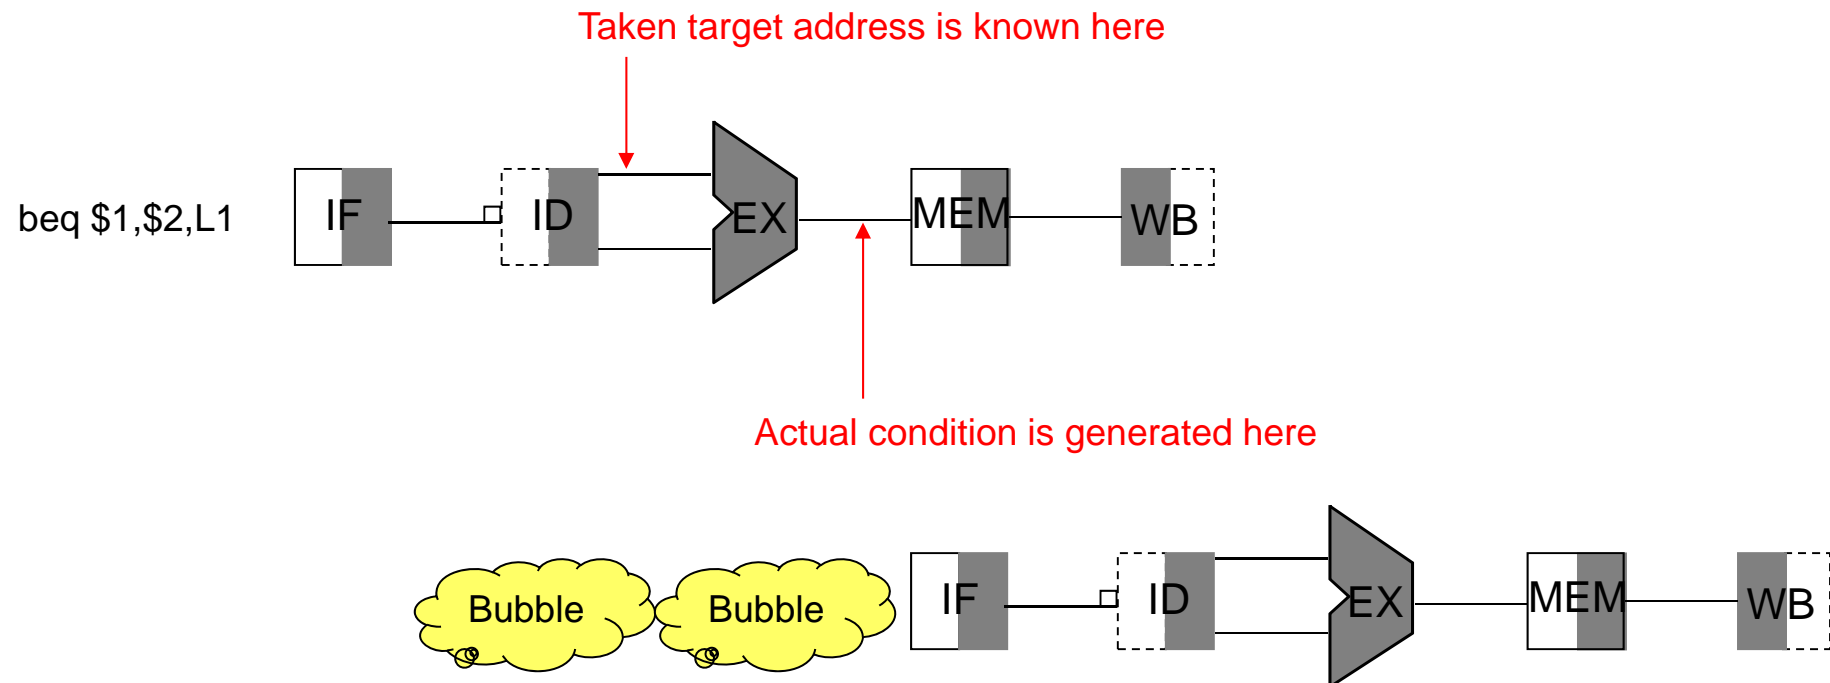


Load-to-Use Data Hazard




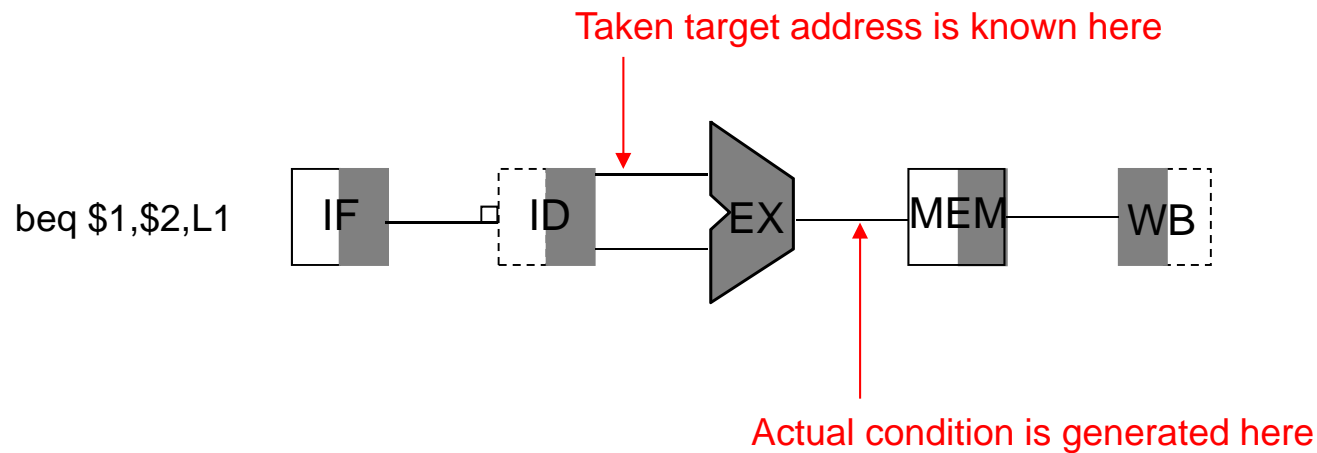
- This bubble can be hidden by proper instruction scheduling
- Hardware interlock is needed to install the pipeline stall

Control Hazard



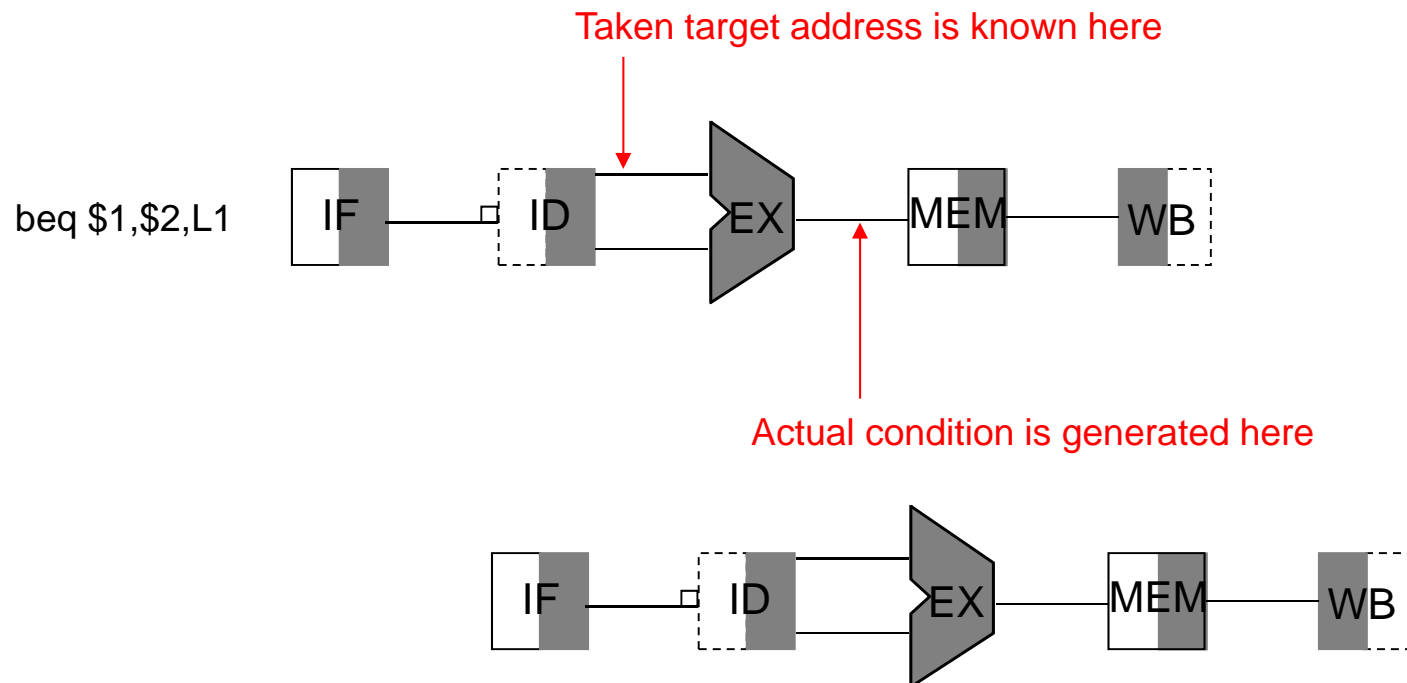
- Solution — branch prediction

Control Hazard



- Branch prediction — Predict “branch direction”

Control Hazard



- Branch prediction — Predict “branch target” and “branch direction”
- Static Branch Prediction
 - Use opcode
 - Use branch offset (+/-), e.g. backward taken, forward not taken,
- Dynamic Branch Predictor

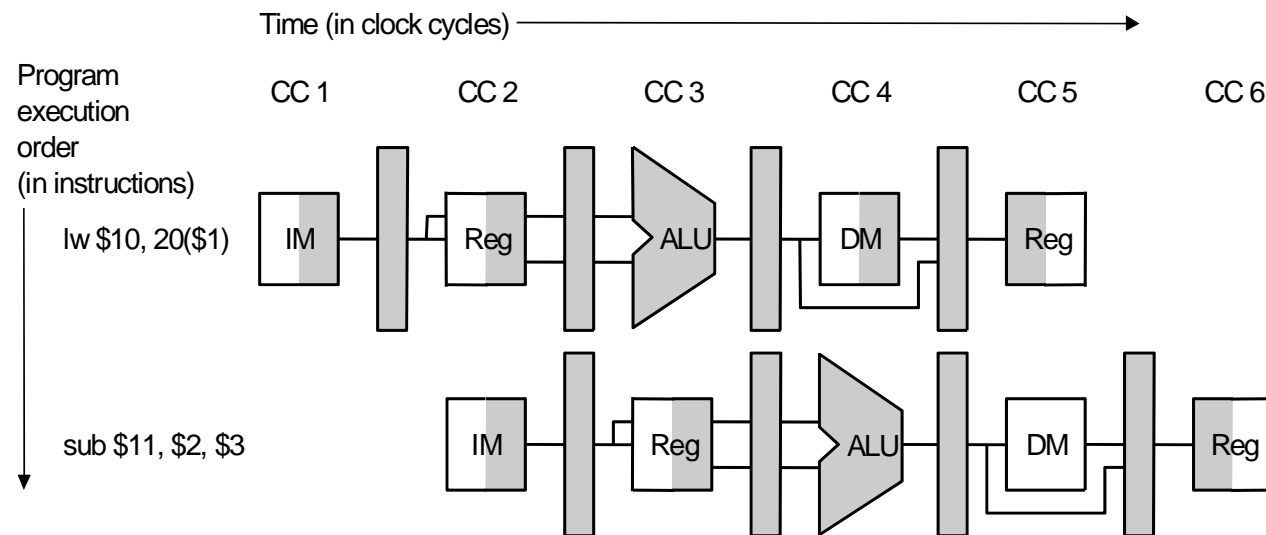
Delay Slot (MIPS)

- Expose pipeline into the ISA design
- Load and jump/branch entail a “delay slot” (e.g. delayed branch)
 - No need for load delay slot in R4000 and after for hardware interlock was implemented
- The instruction right after the jump or branch is executed before the jump/branch

```
jal    function_A
add    $4, $5, $6    ; executed before jmp
lw     $12, 8($4)    ; executed after return
```

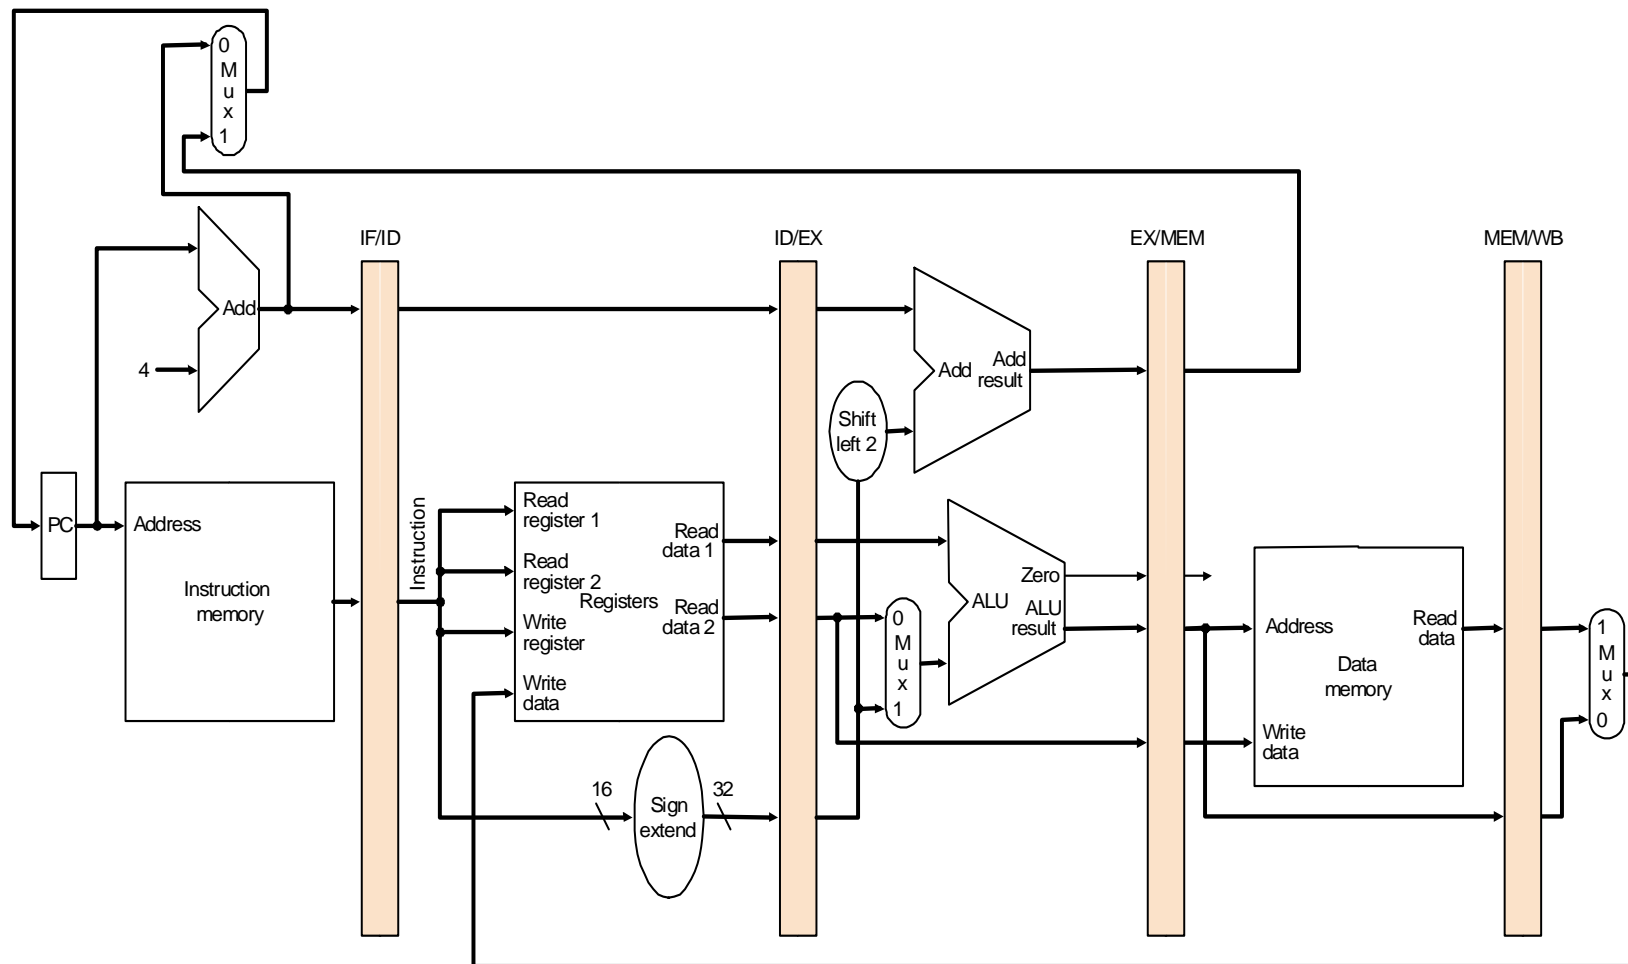
- Jump/branch and the delay slot instruction are considered “indivisible”
- In the delay slot, the compiler needs to schedule
 - A useful instruction (either before the jmp, or after the jmp w/o side effect)
 - otherwise a NOP

Graphically Representing Pipelines

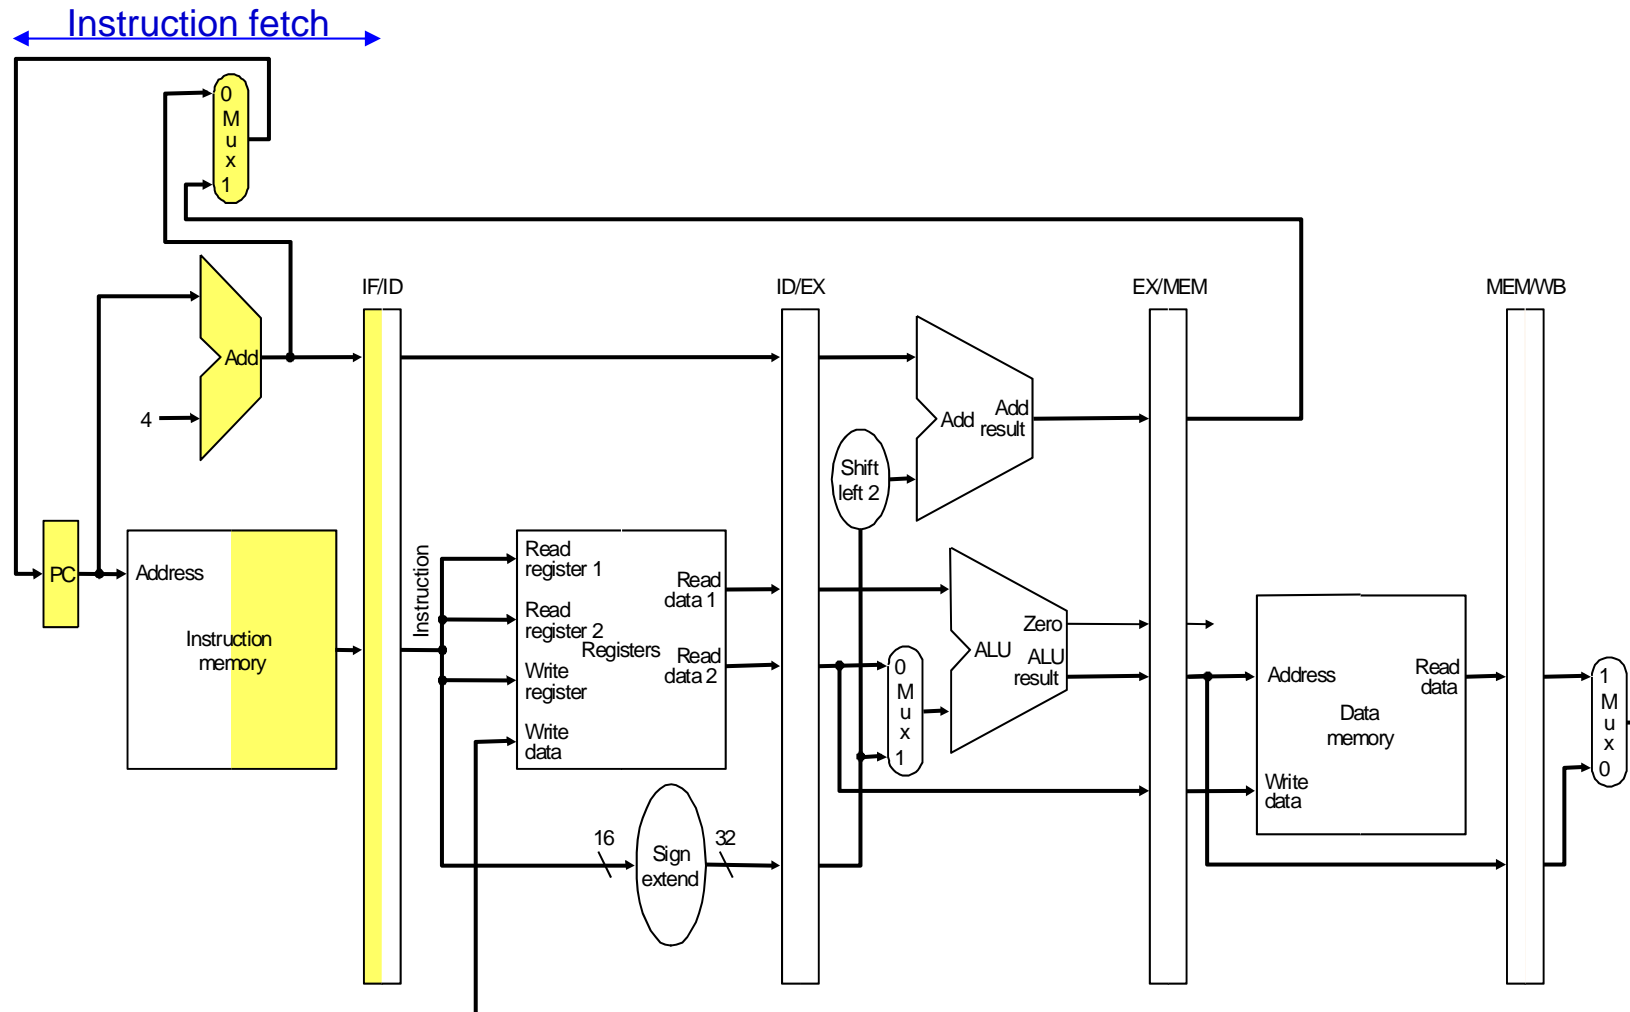


- Can help with answering questions like:
 - how many cycles does it take to execute this code?
 - what is the ALU doing during cycle 4?
 - use this representation to help understand datapaths

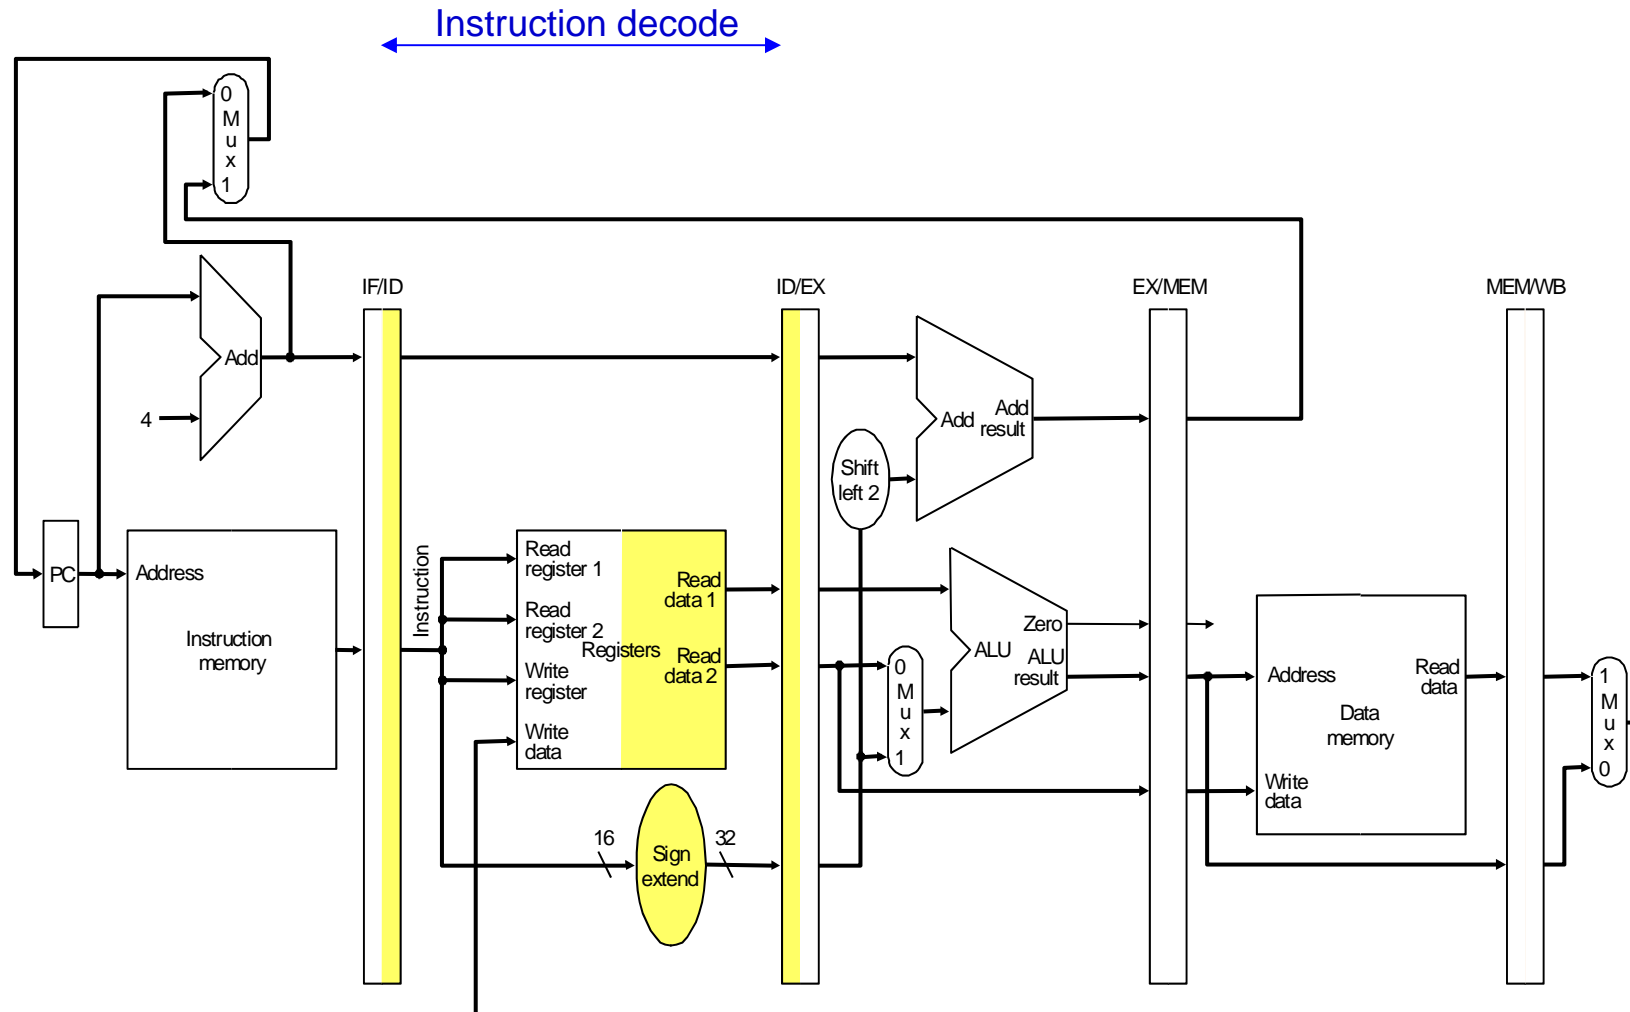
Pipelined Datapath



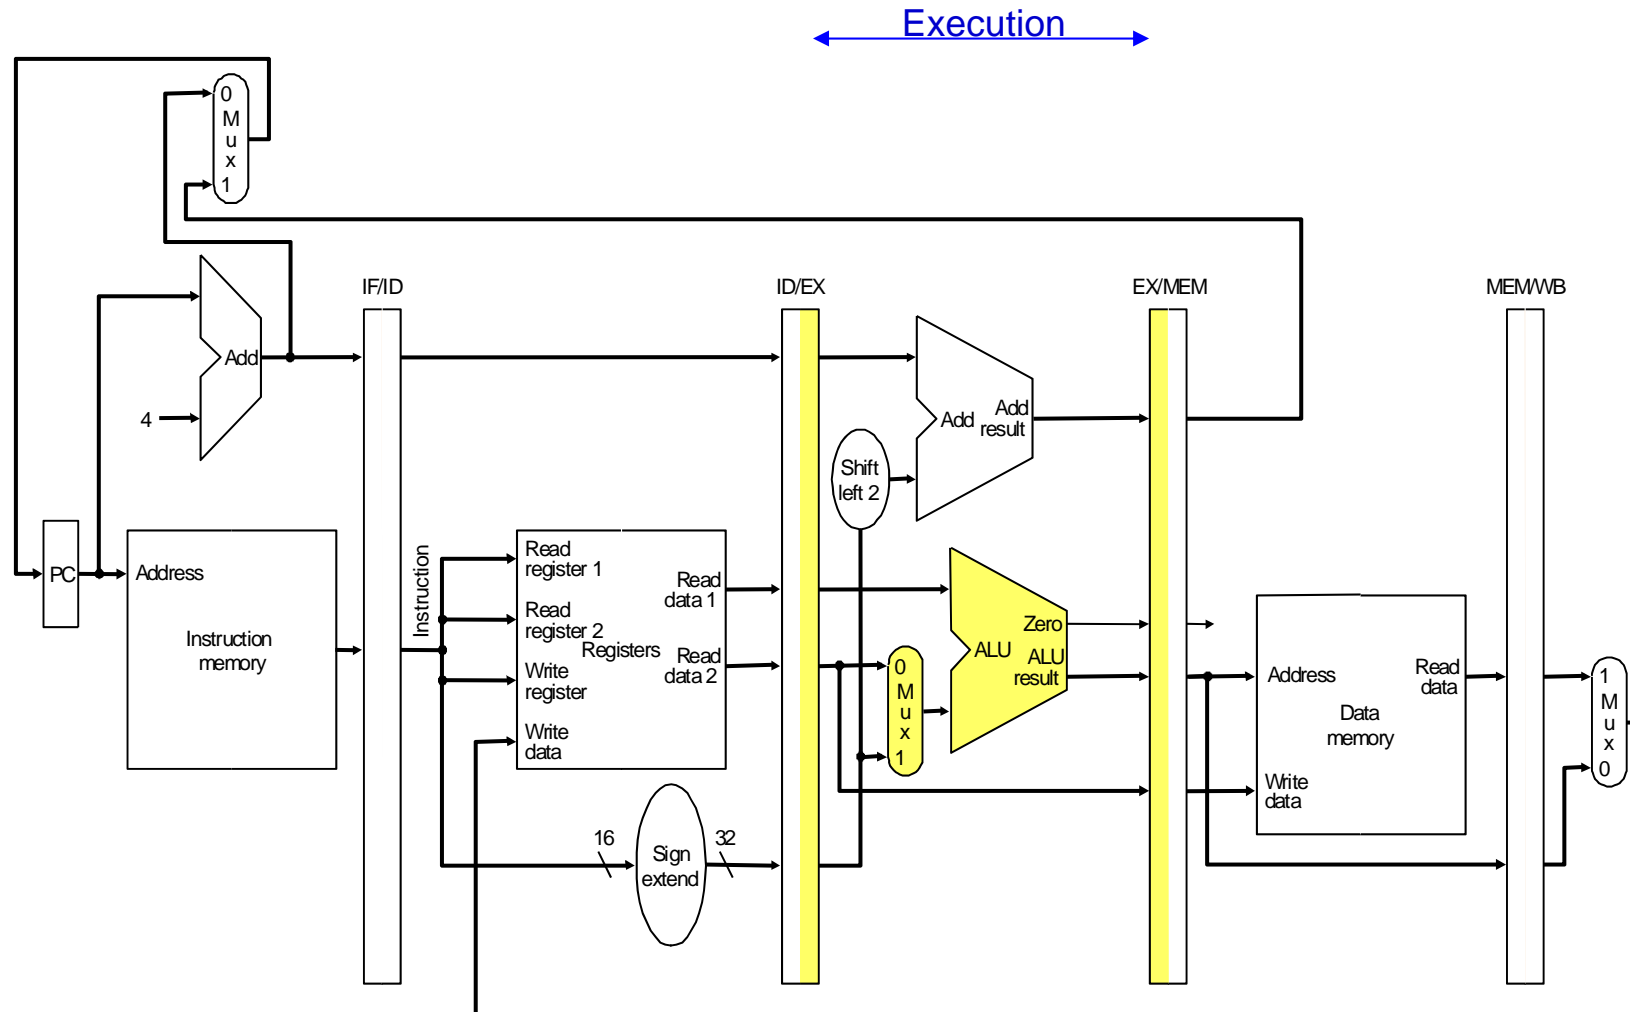
Example for lw instruction: Instruction Fetch (IF)



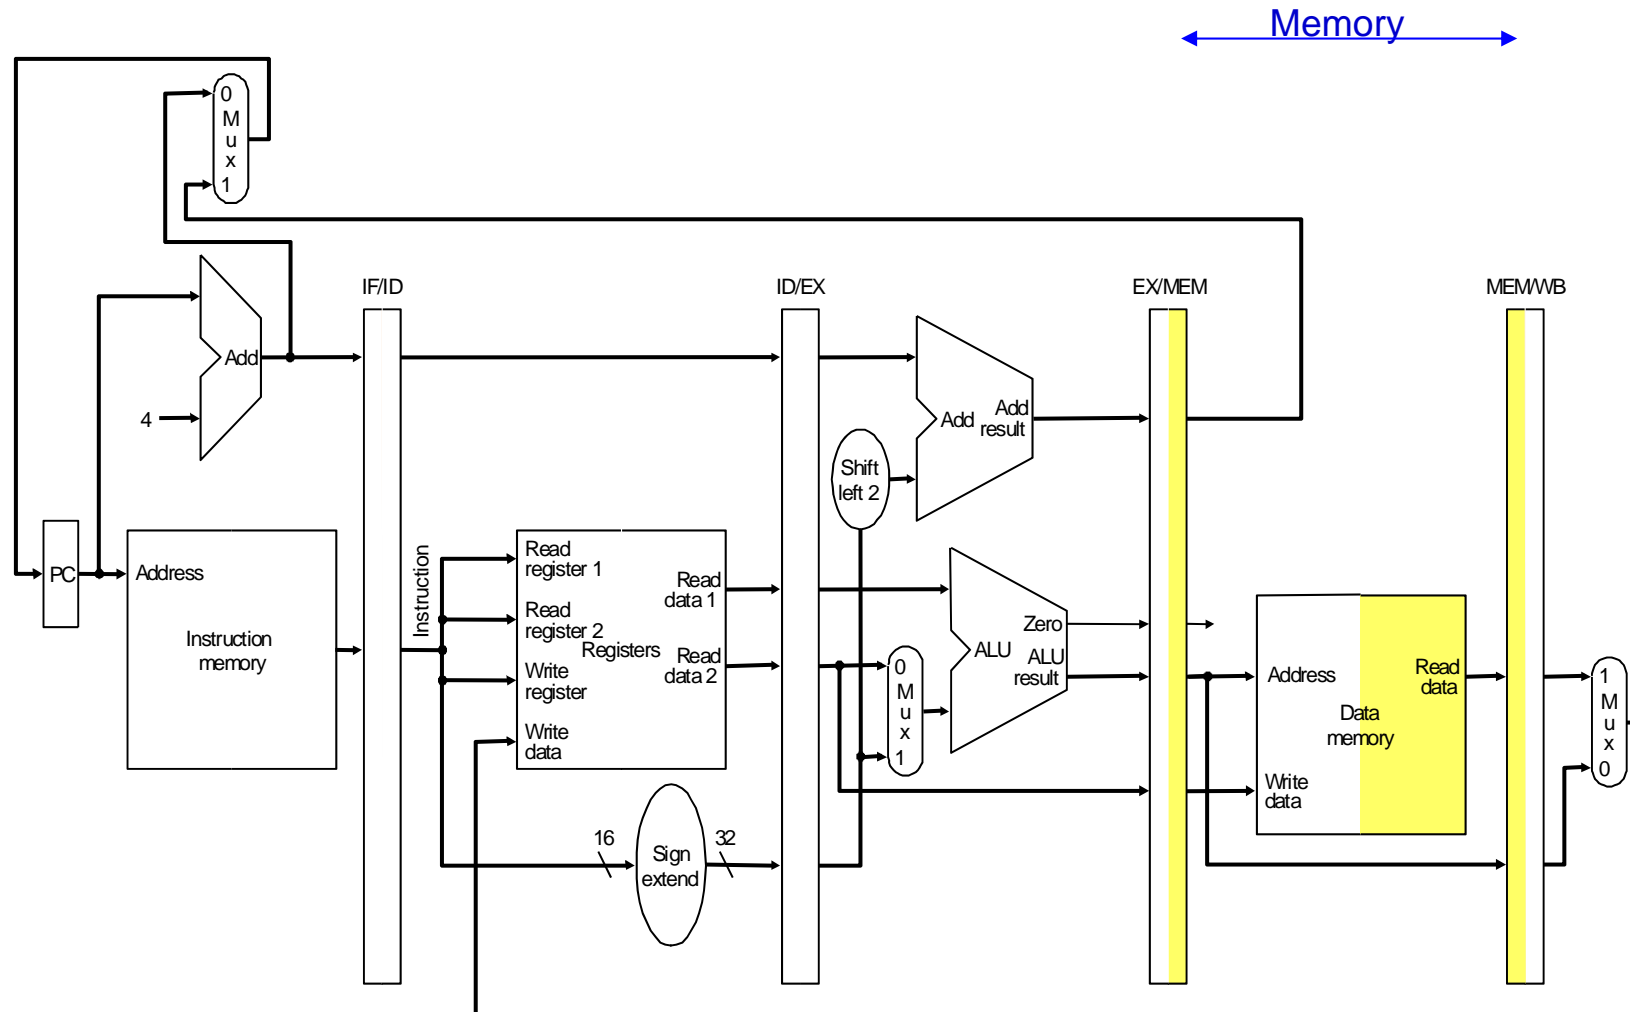
Example for lw instruction: Instruction Decode (ID)



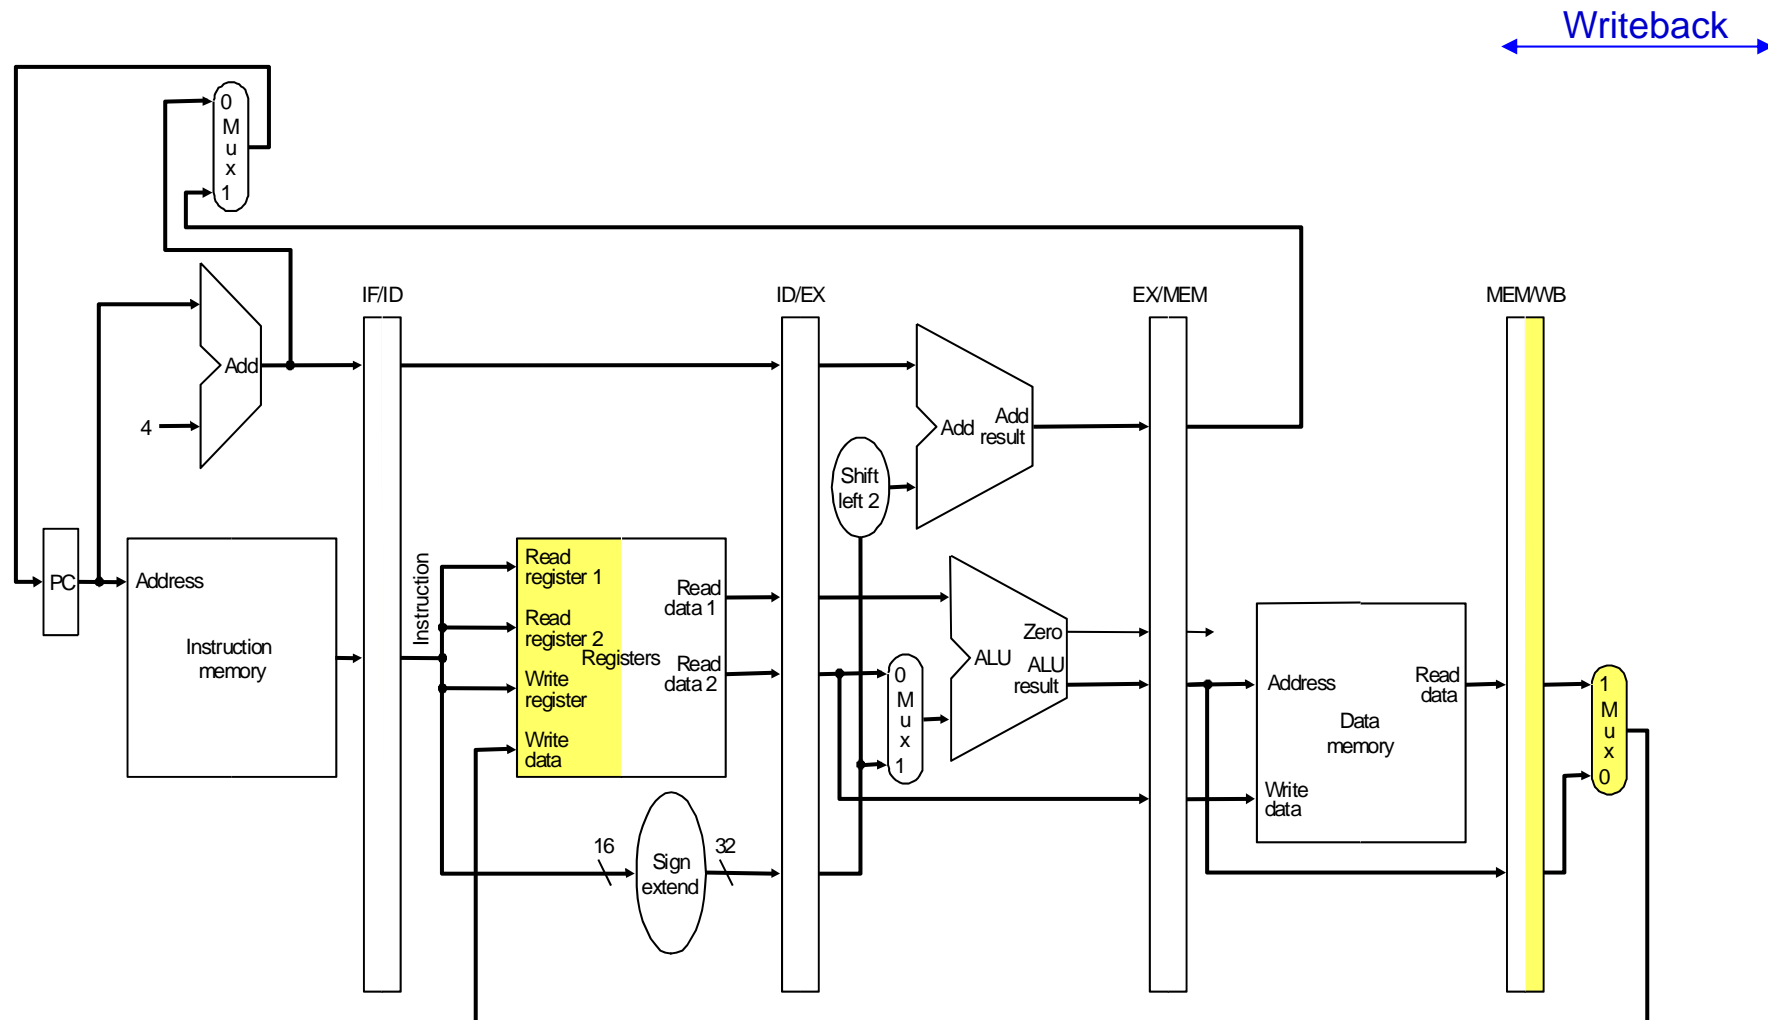
Example for lw instruction: Execution (EX)



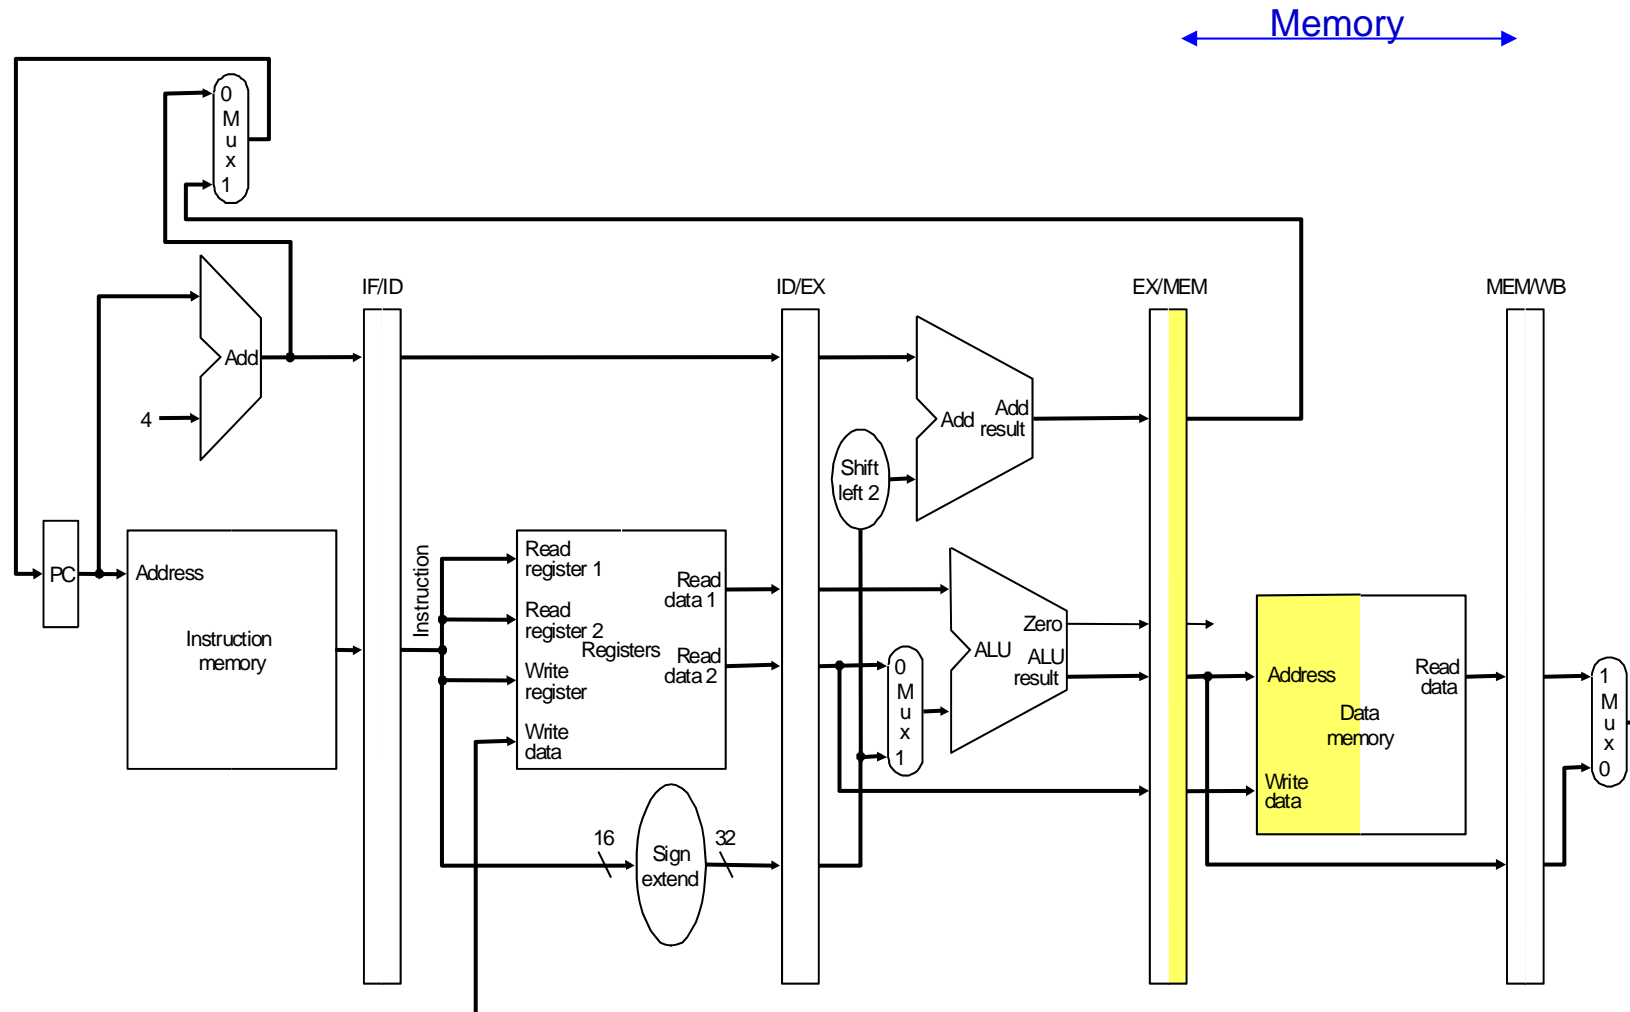
Example for lw instruction: Memory (MEM)



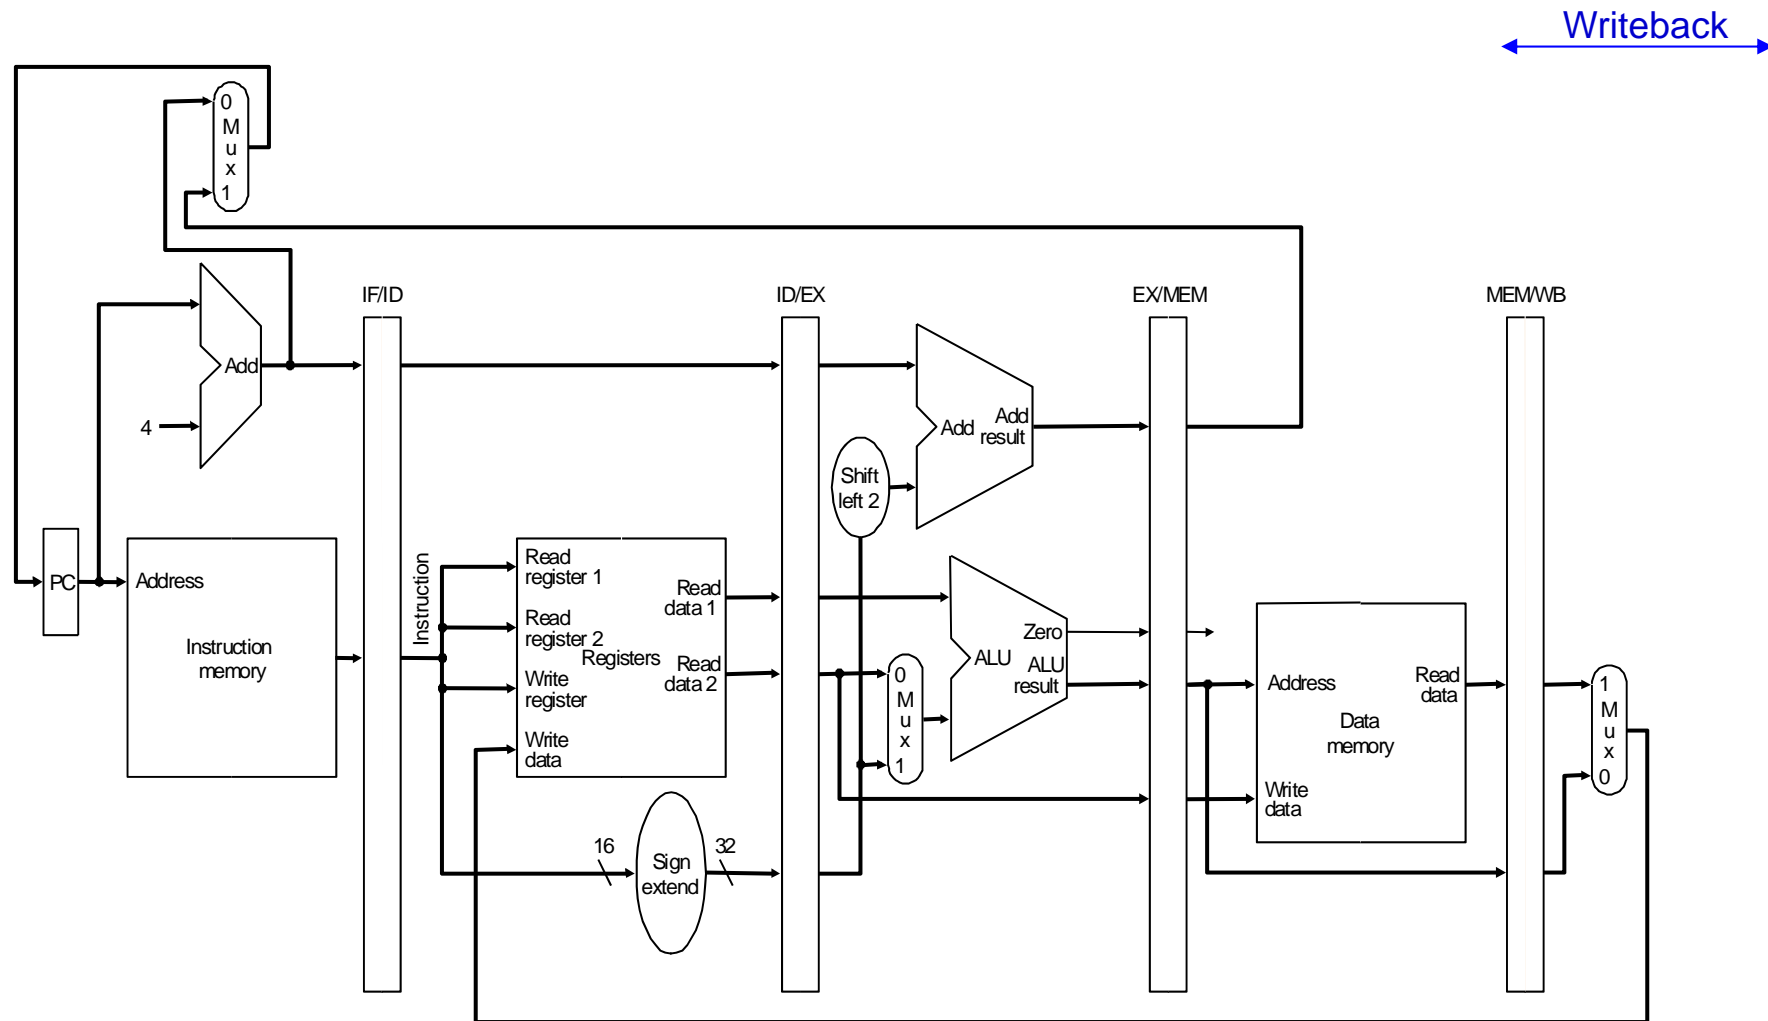
Example for lw instruction: Writeback (WB)



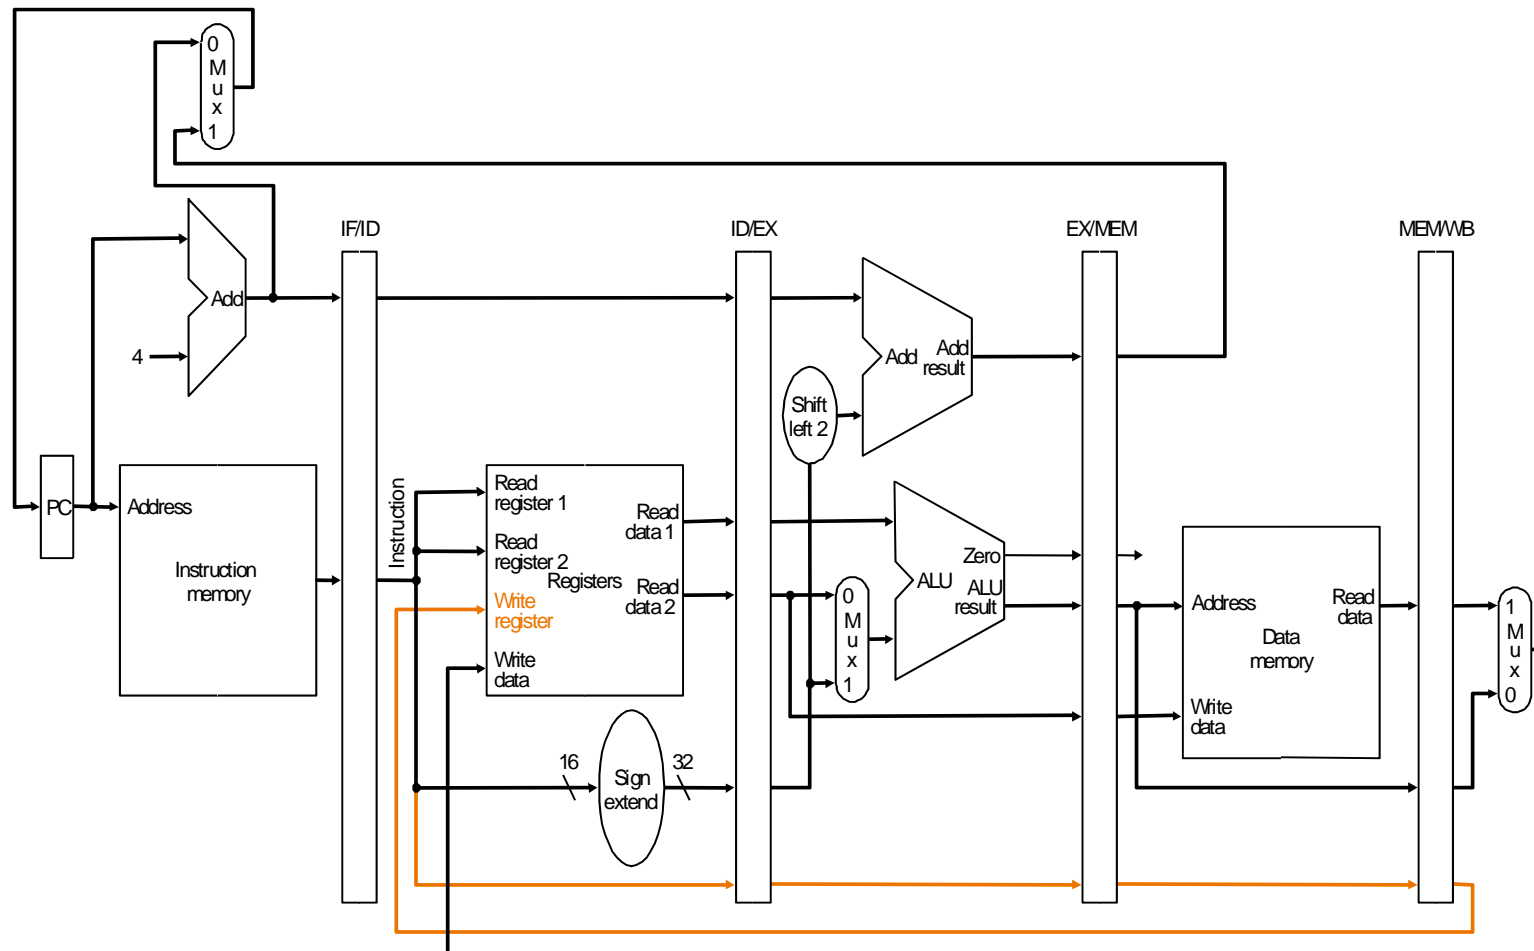
Example for sw instruction: Memory (MEM)



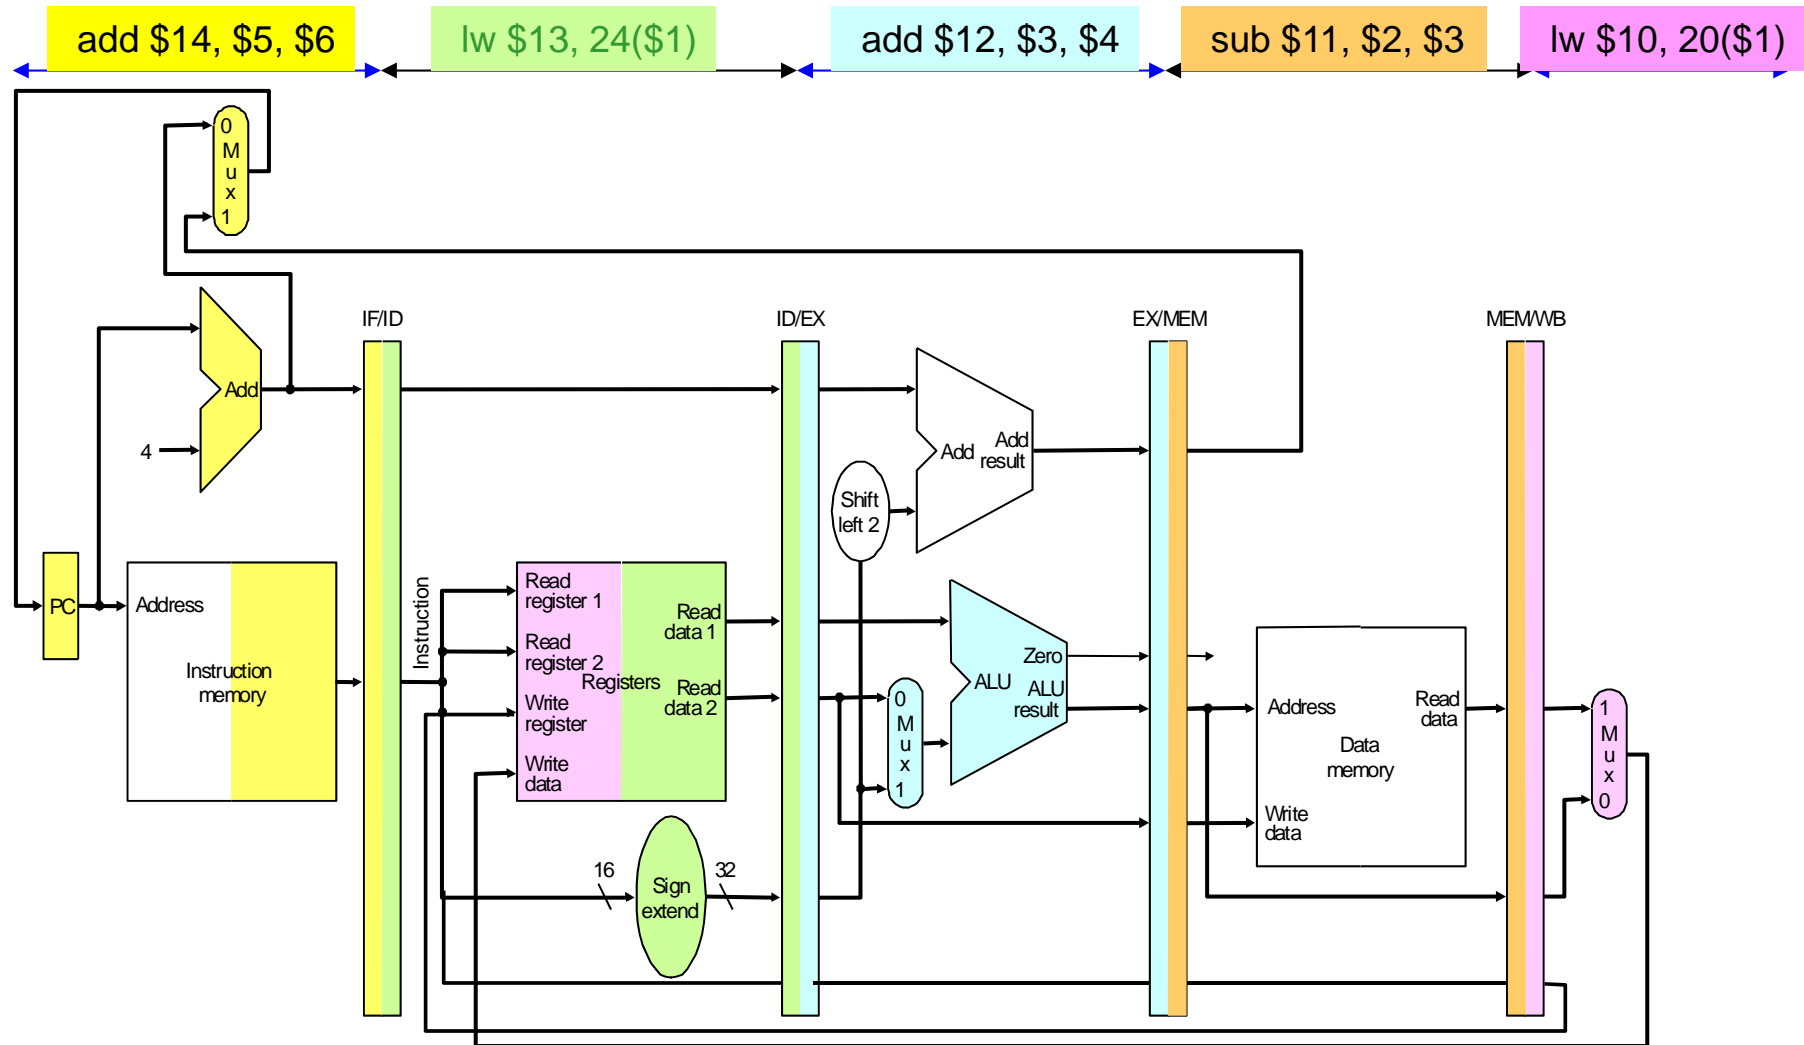
Example for sw instruction: Writeback (WB). do nothing



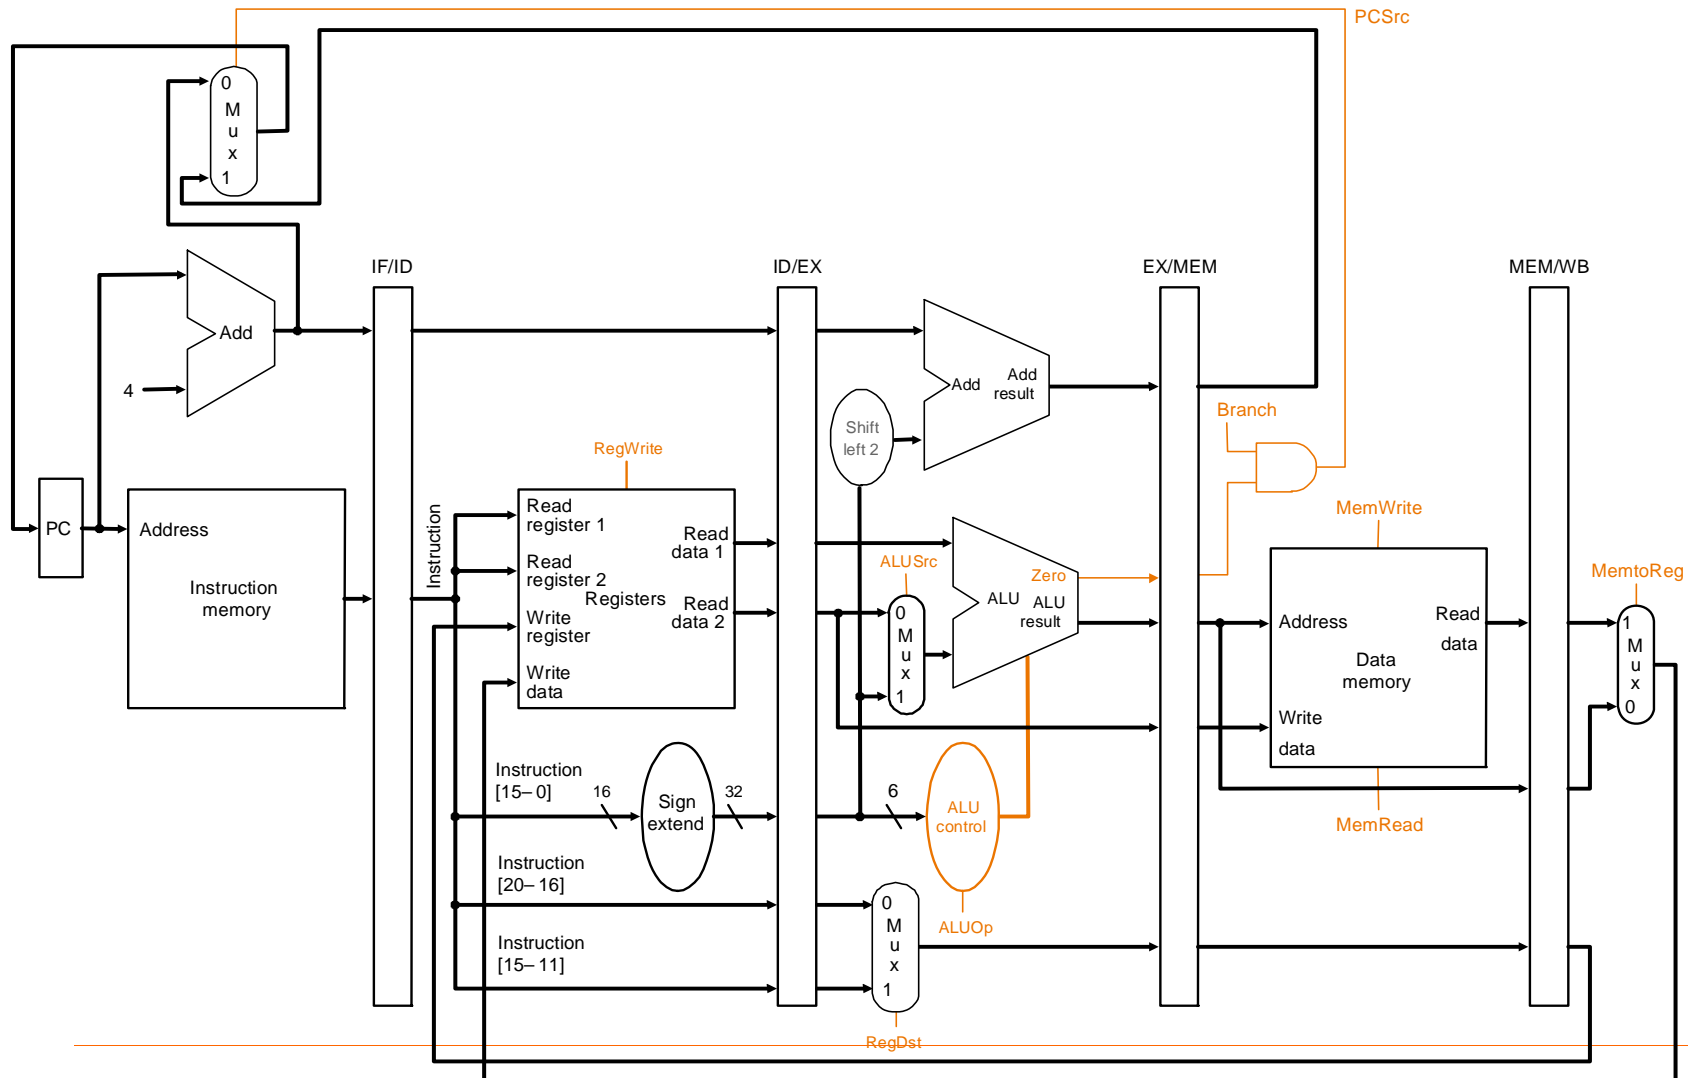
Corrected Datapath (for lw)



Pipelining Example



Pipeline Control




Pipeline control

- ☐ We have 5 stages. What needs to be controlled in each stage?
 - Instruction Fetch and PC Increment
 - Instruction Decode / Register Fetch
 - Execution (4 lines)
 - ☐ RegDst
 - ☐ ALUOp[1:0]
 - ☐ ALUSrc
 - Memory Stage (3 lines)
 - ☐ Branch
 - ☐ MemRead
 - ☐ MemWrite
 - Write Back (2 lines)
 - ☐ MemtoReg
 - ☐ RegWrite (note that this signal is in ID stage)

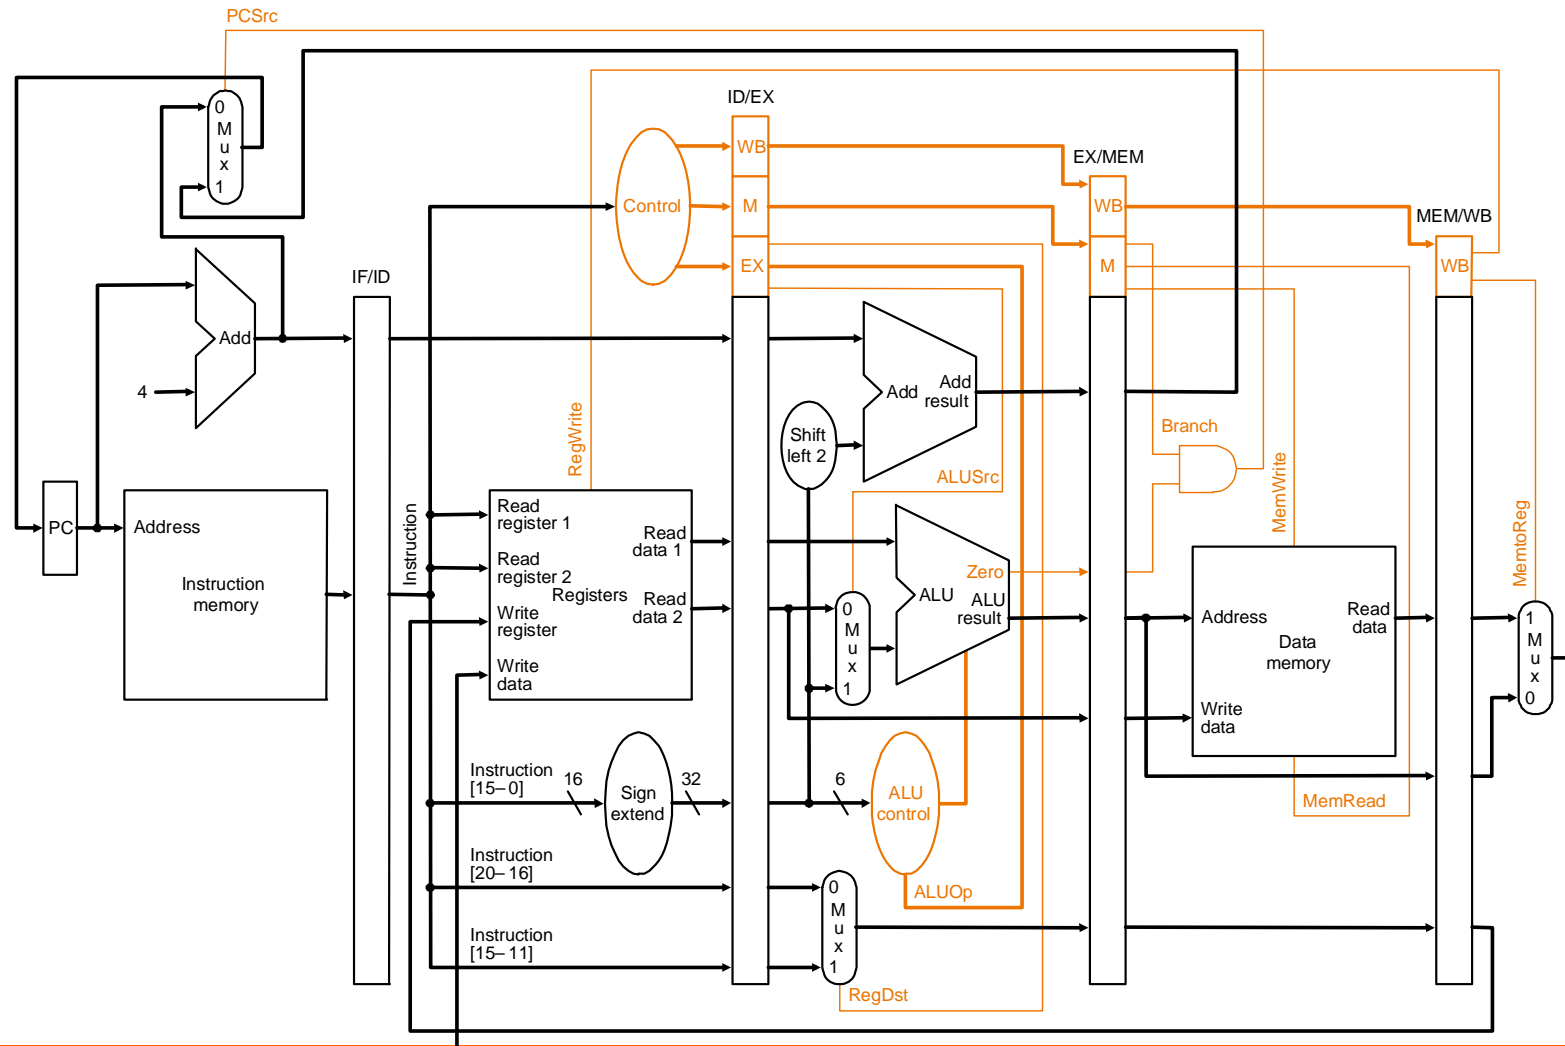
Pipeline Control

- Extend pipeline registers to include control information (created in ID)
- Pass control signals along just like the data

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X



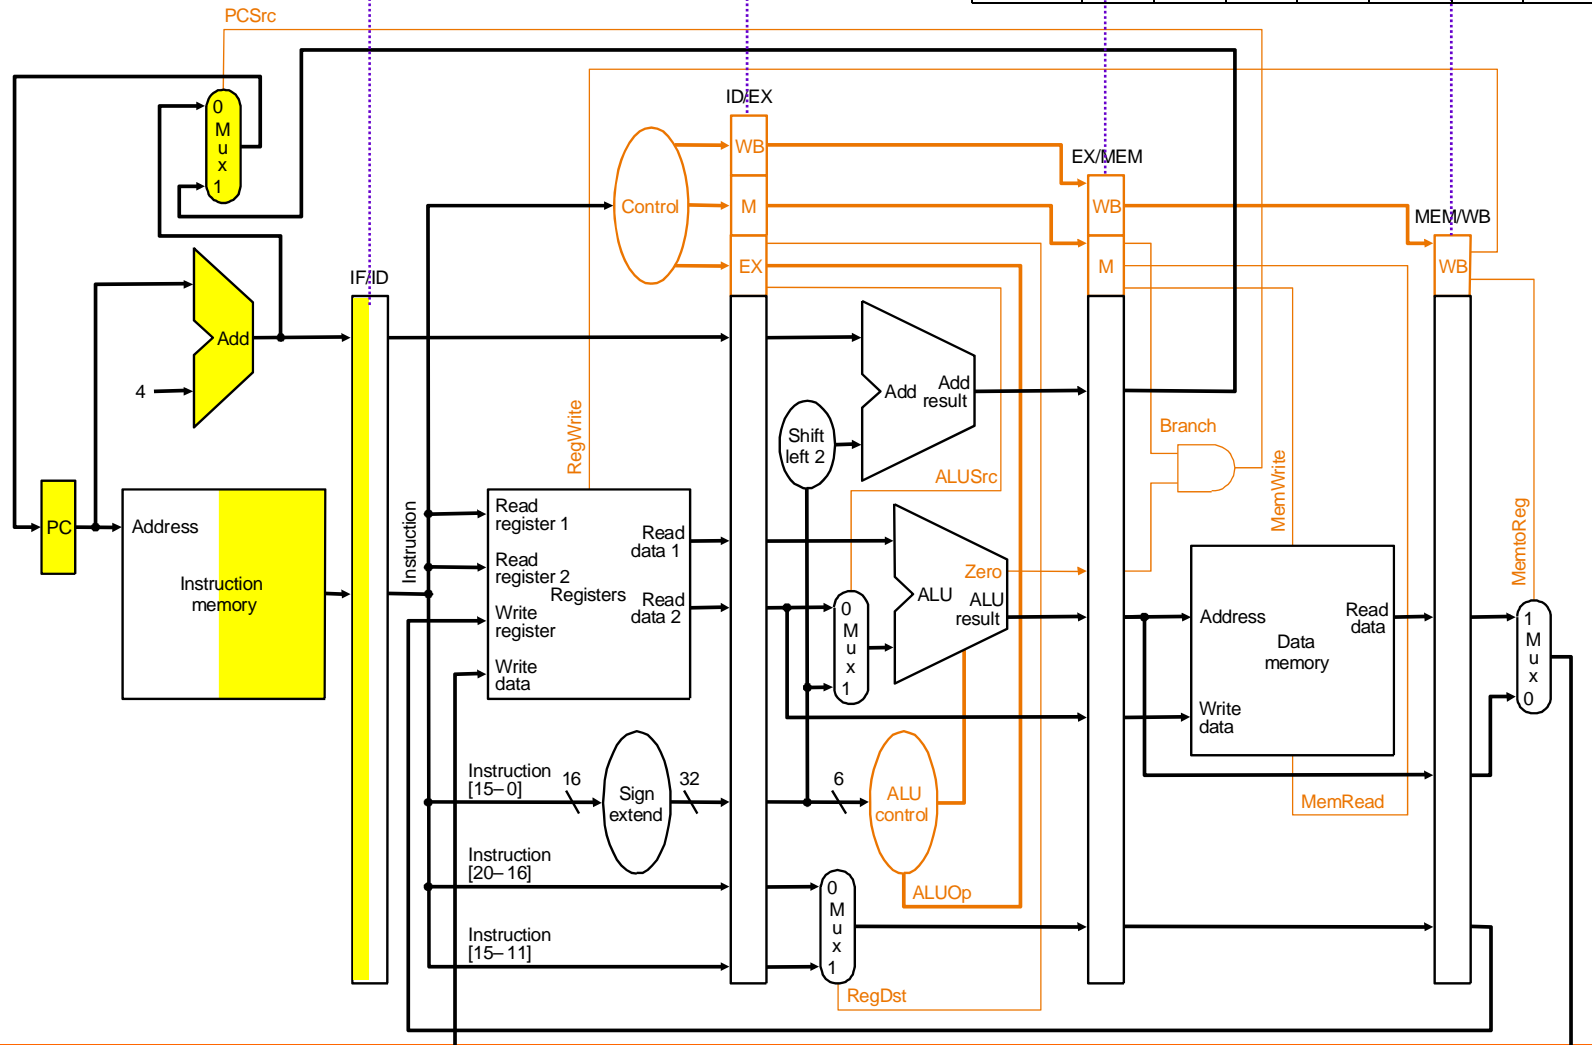
Datapath with Control



Datapath with Control

IF: lw \$10, 9(\$1)

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

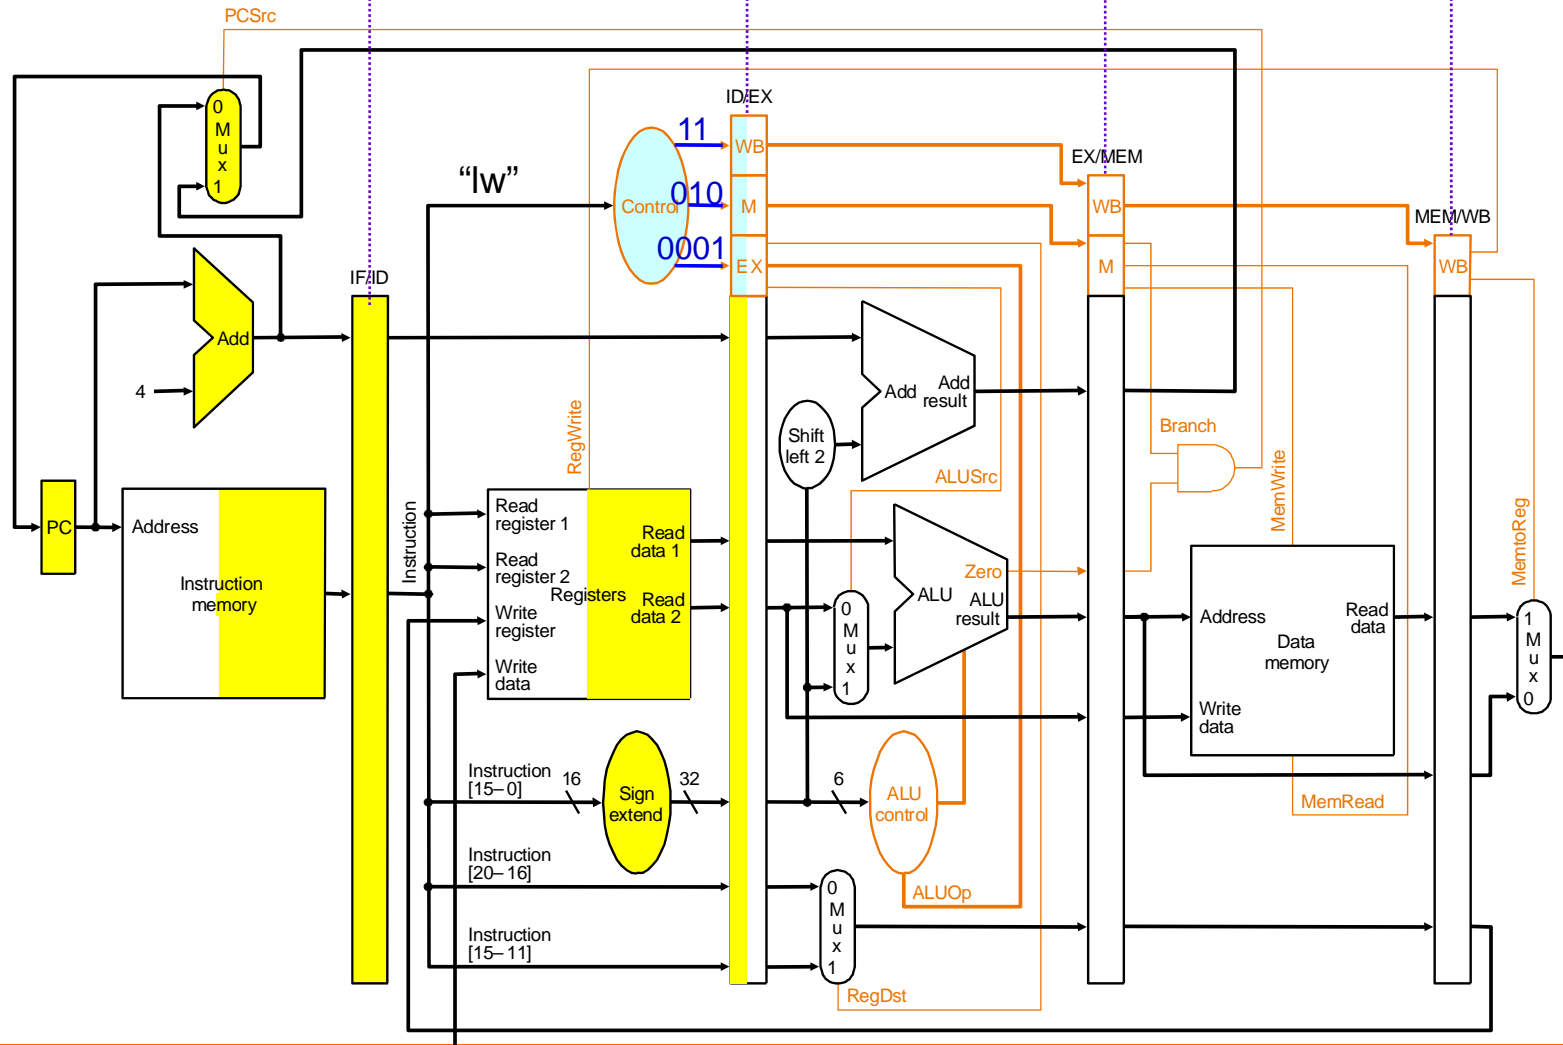


Datapath with Control

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

IF: sub \$11, \$2, \$3

ID: lw \$10, 9(\$1)



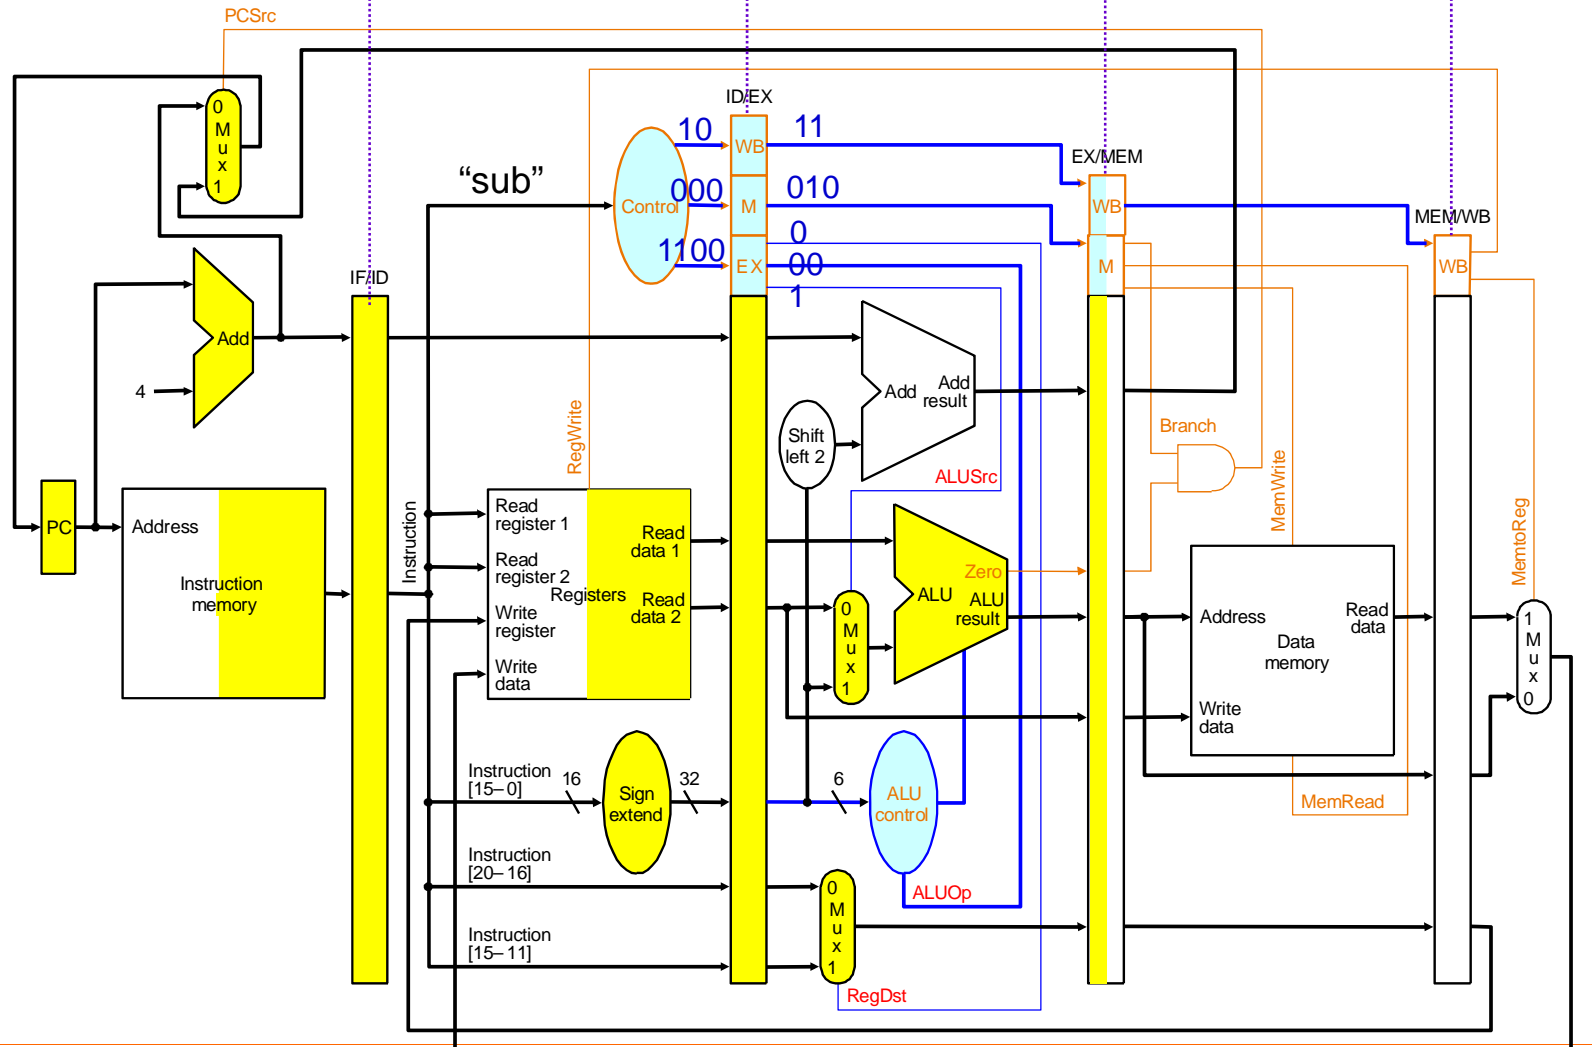
Datapath with Control

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
reg. lw	X	0	1	0	1	0	0	0	X

IF: and \$12, \$4, \$5

ID: sub \$11, \$2, \$3

EX: lw \$10, 9(\$1)



Datapath with Control

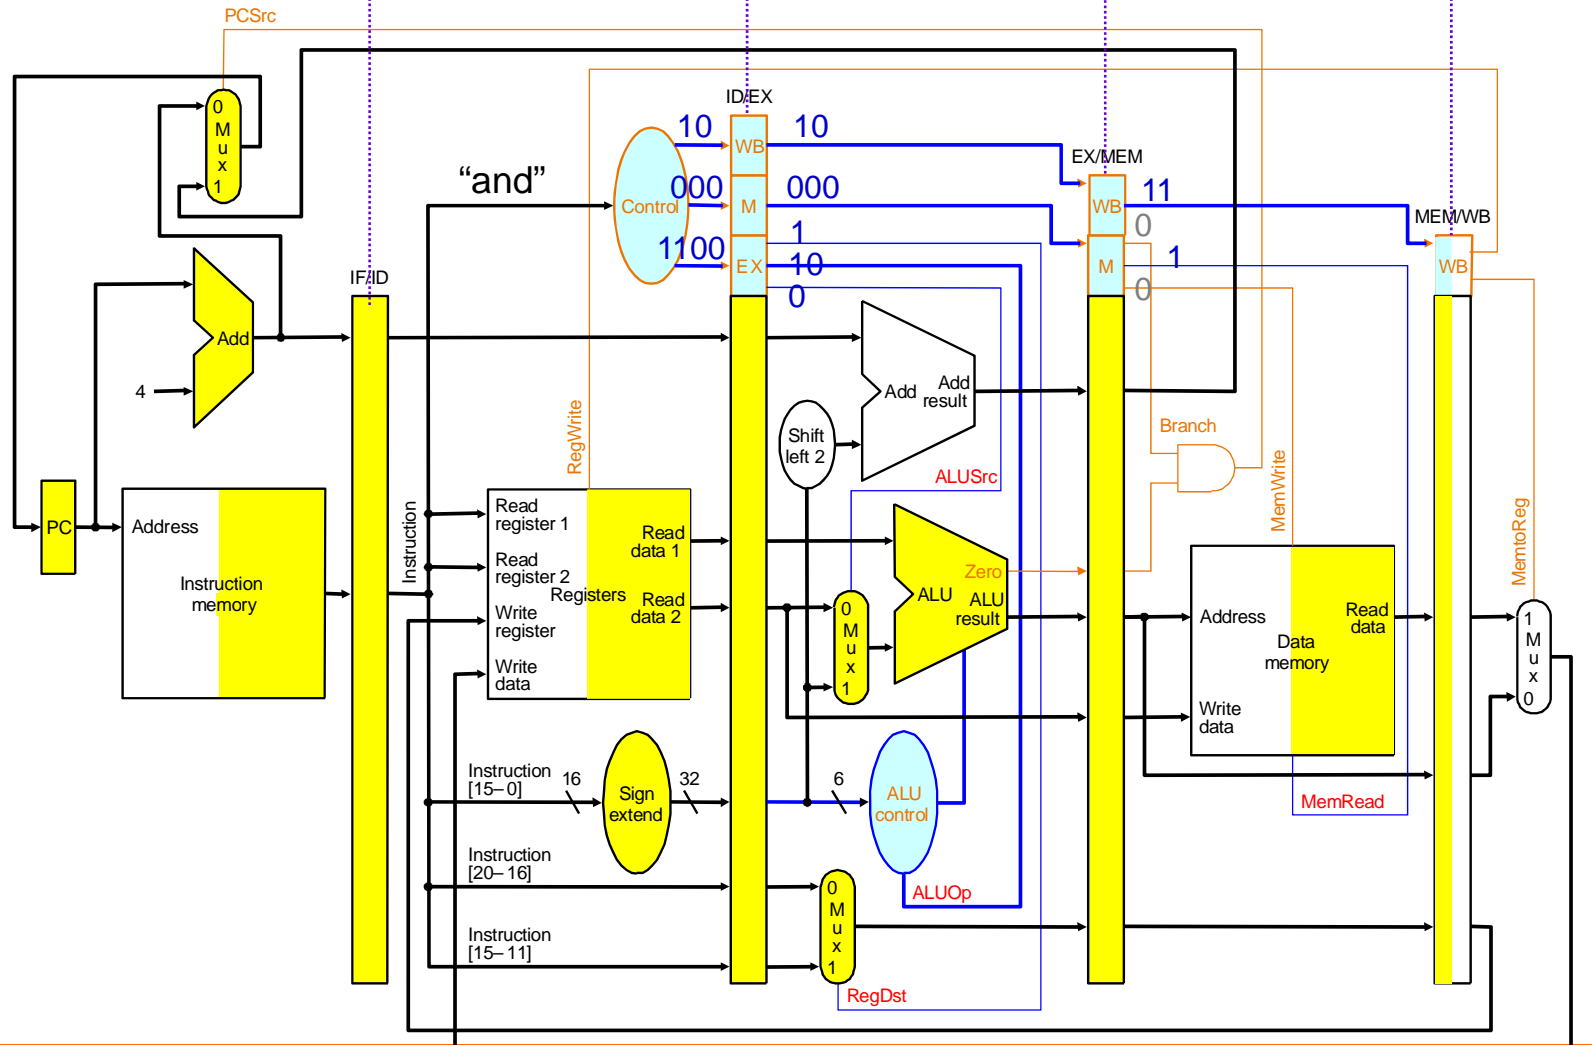
Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	0	0	1	0	0	0	X

IF: or \$13, \$6, \$7

ID: and \$12, \$4, \$5

EX: sub \$11, \$2, \$3

MEM: lw \$10, 9(\$1)



Datapath with Control

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	X	1	0	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

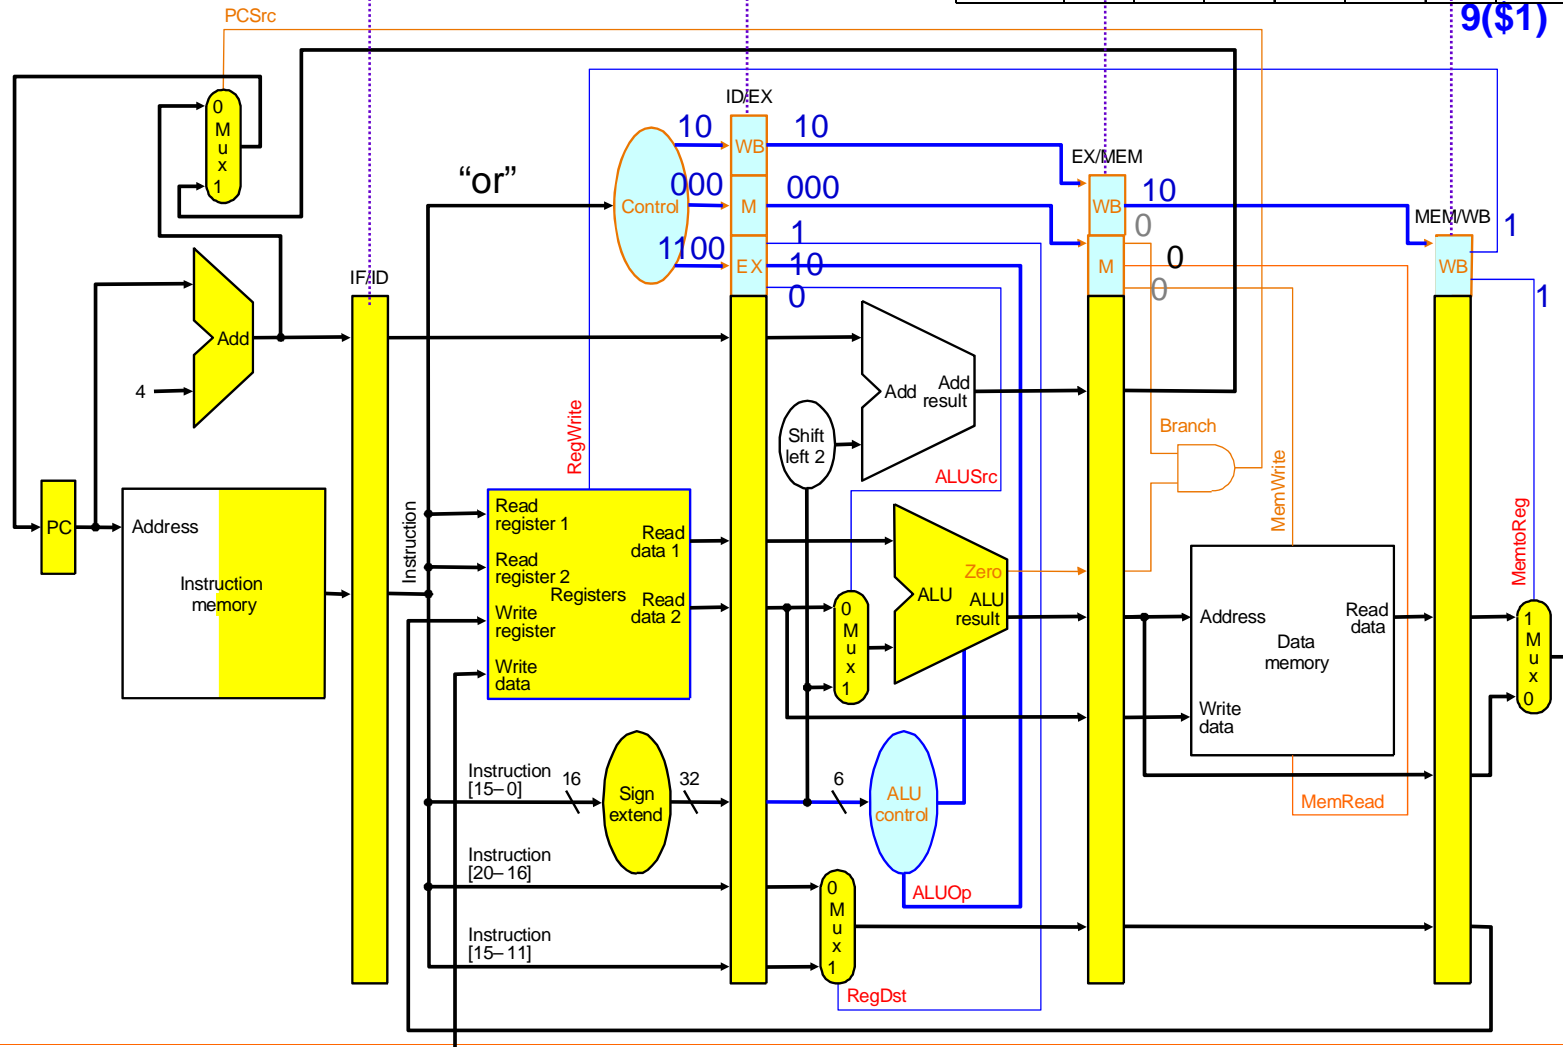
IF: add \$14, \$8, \$9

ID: or \$13, \$6, \$7

EX: and \$12, \$4, \$5

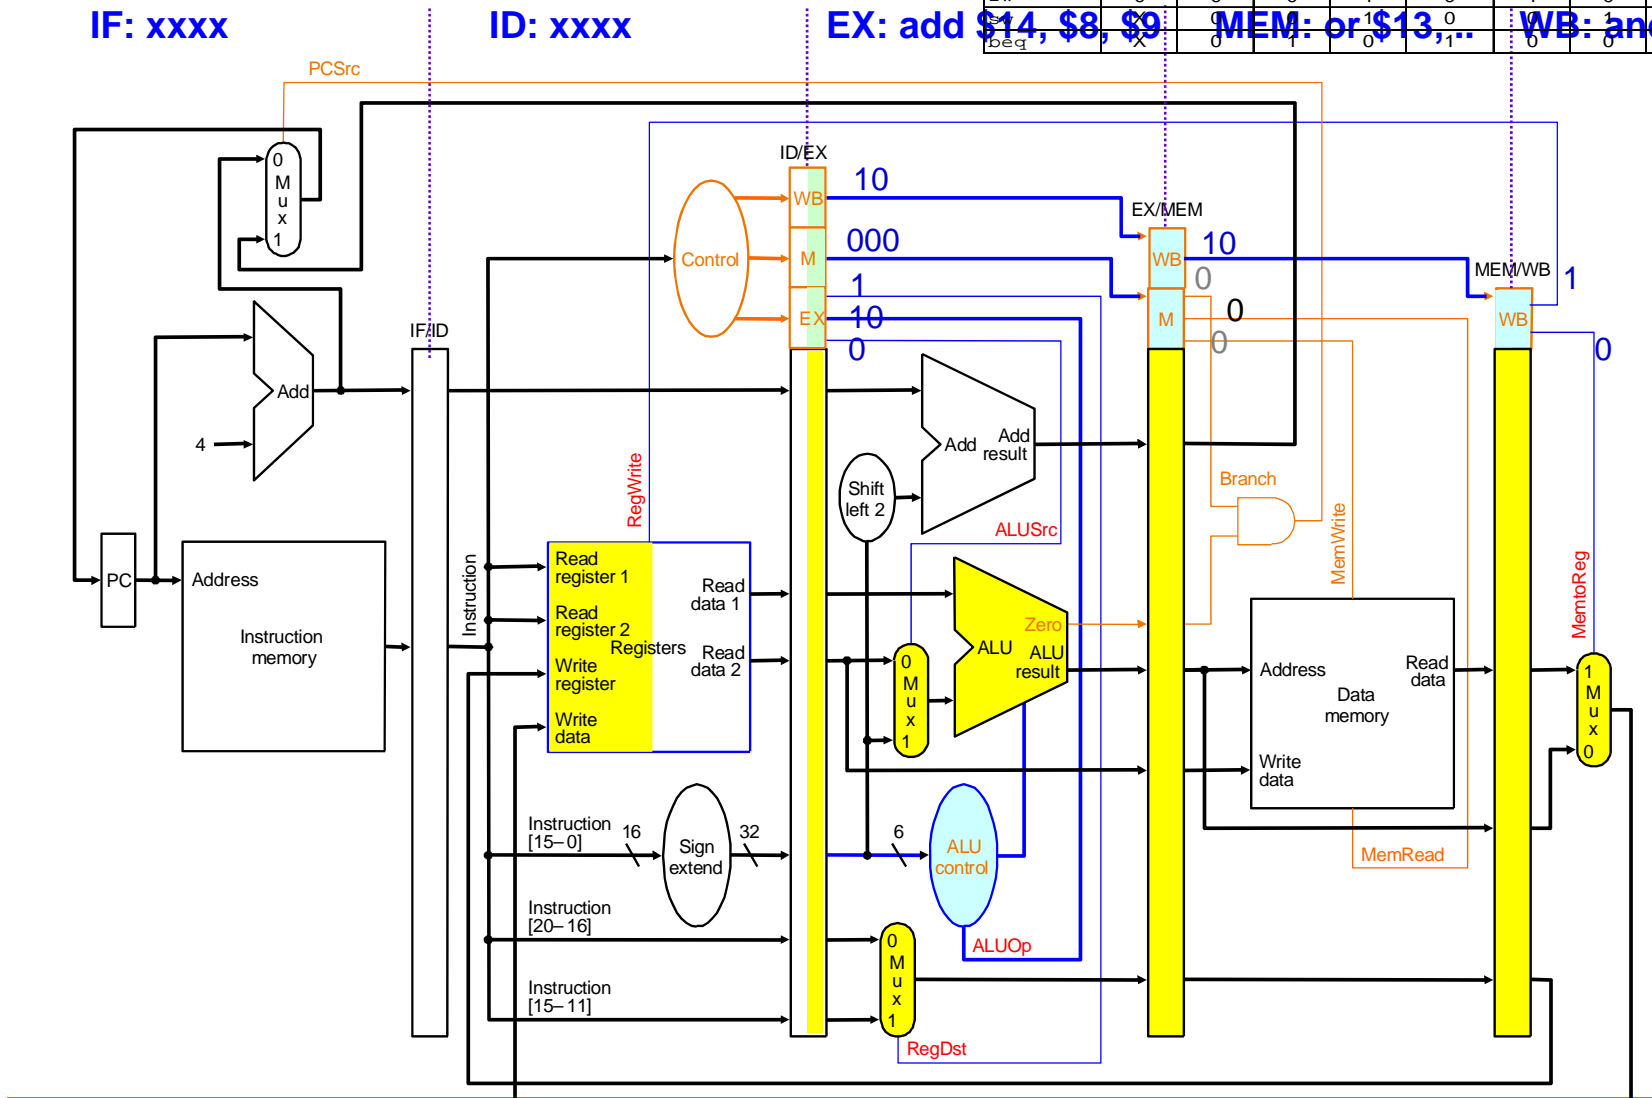
MEM: sub \$11, ...

WB: lw \$10, 9(\$1)



Datapath with Control

Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	0	0	1	0	0	0	1	0	0
beq	0	0	0	0	1	0	0	0	0



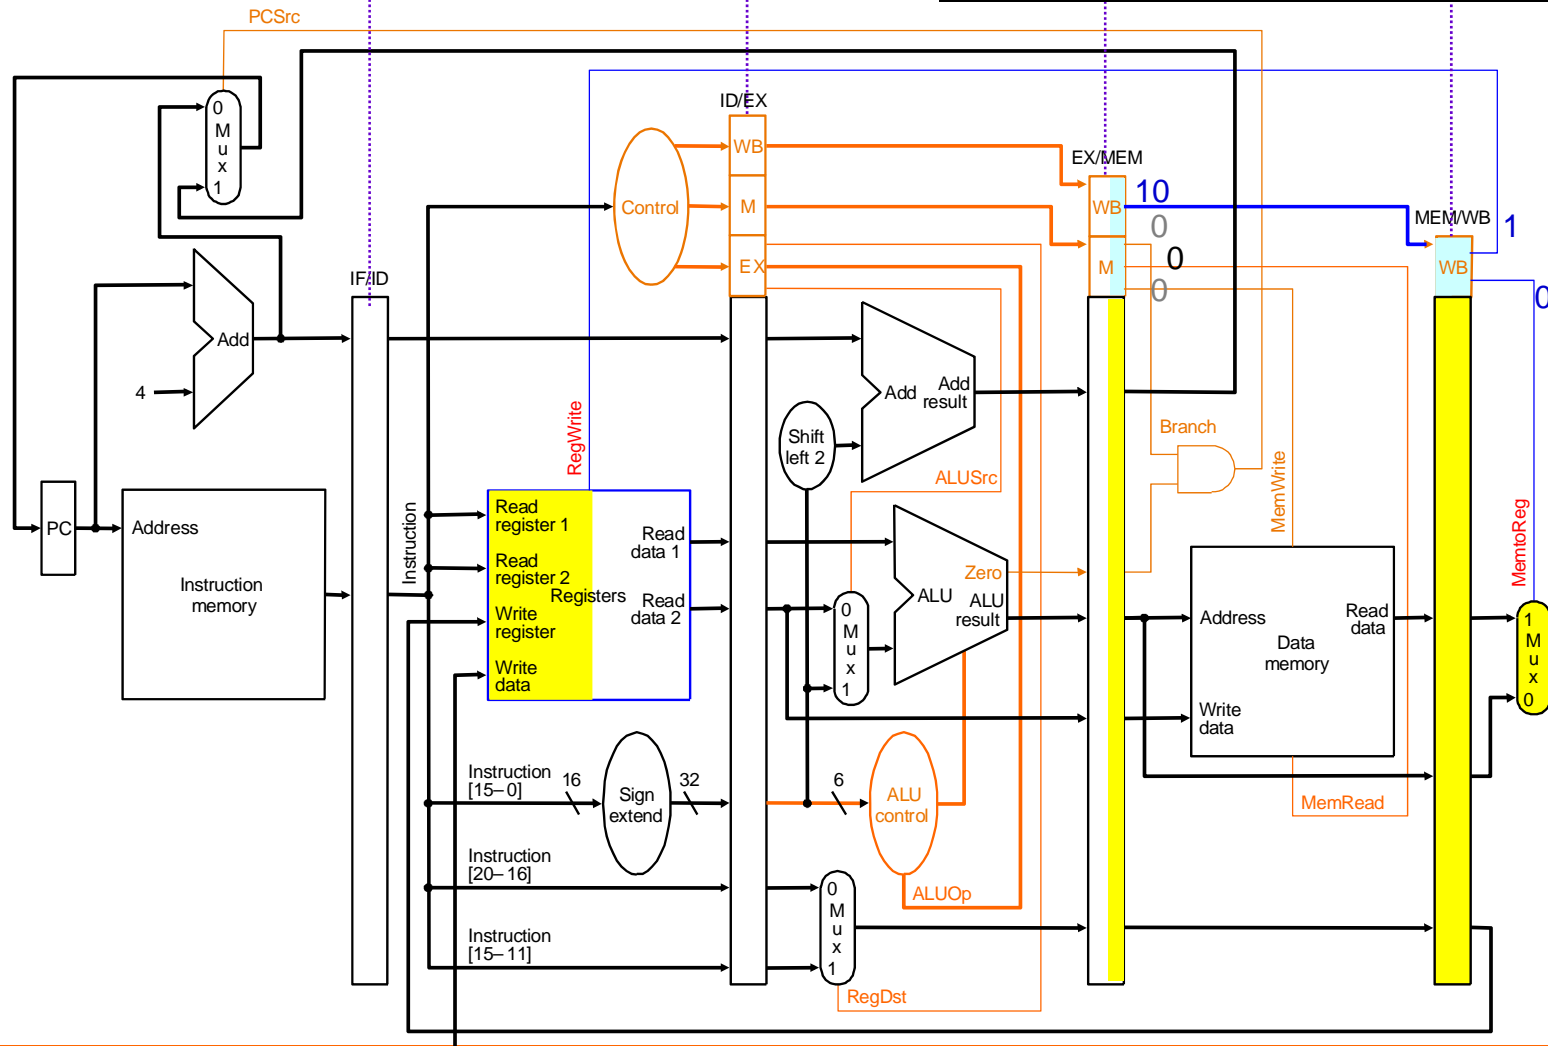
IF: xxxx

ID: xxxx

EX: xxxx

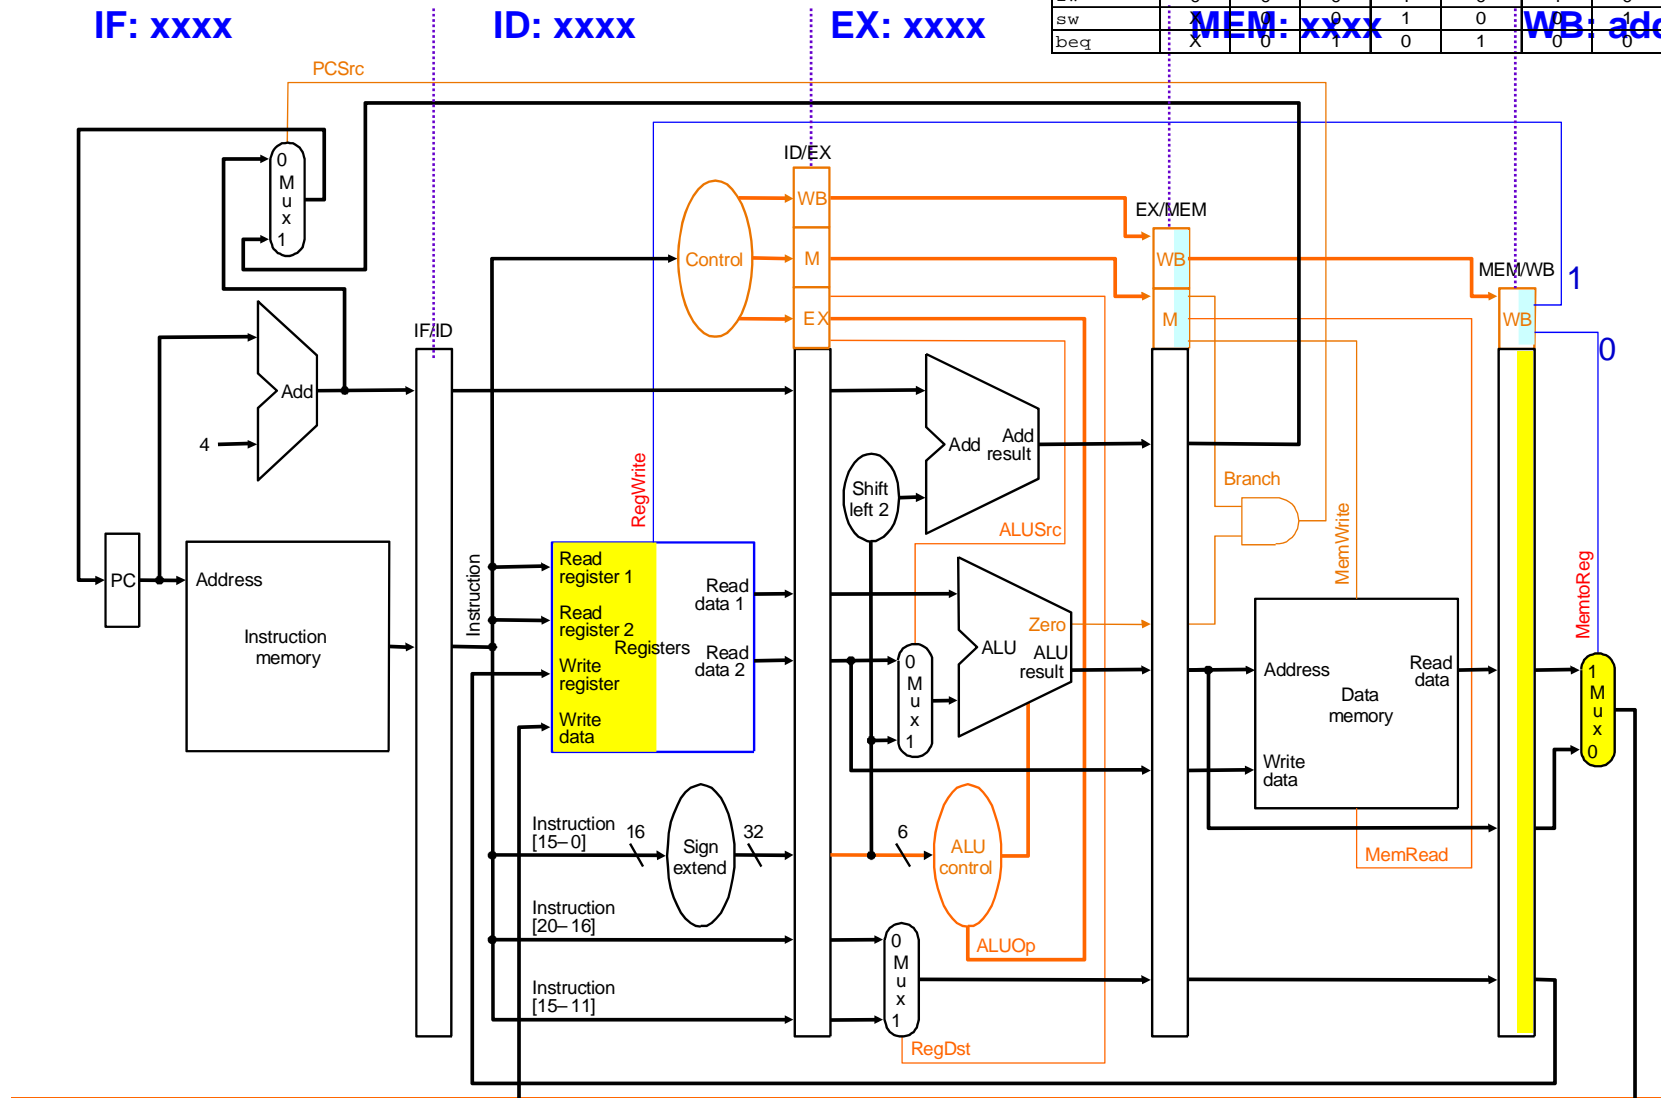
		MEM:	0	add \$14,	0
			1		1

WB: or \$13...



Datapath with Control

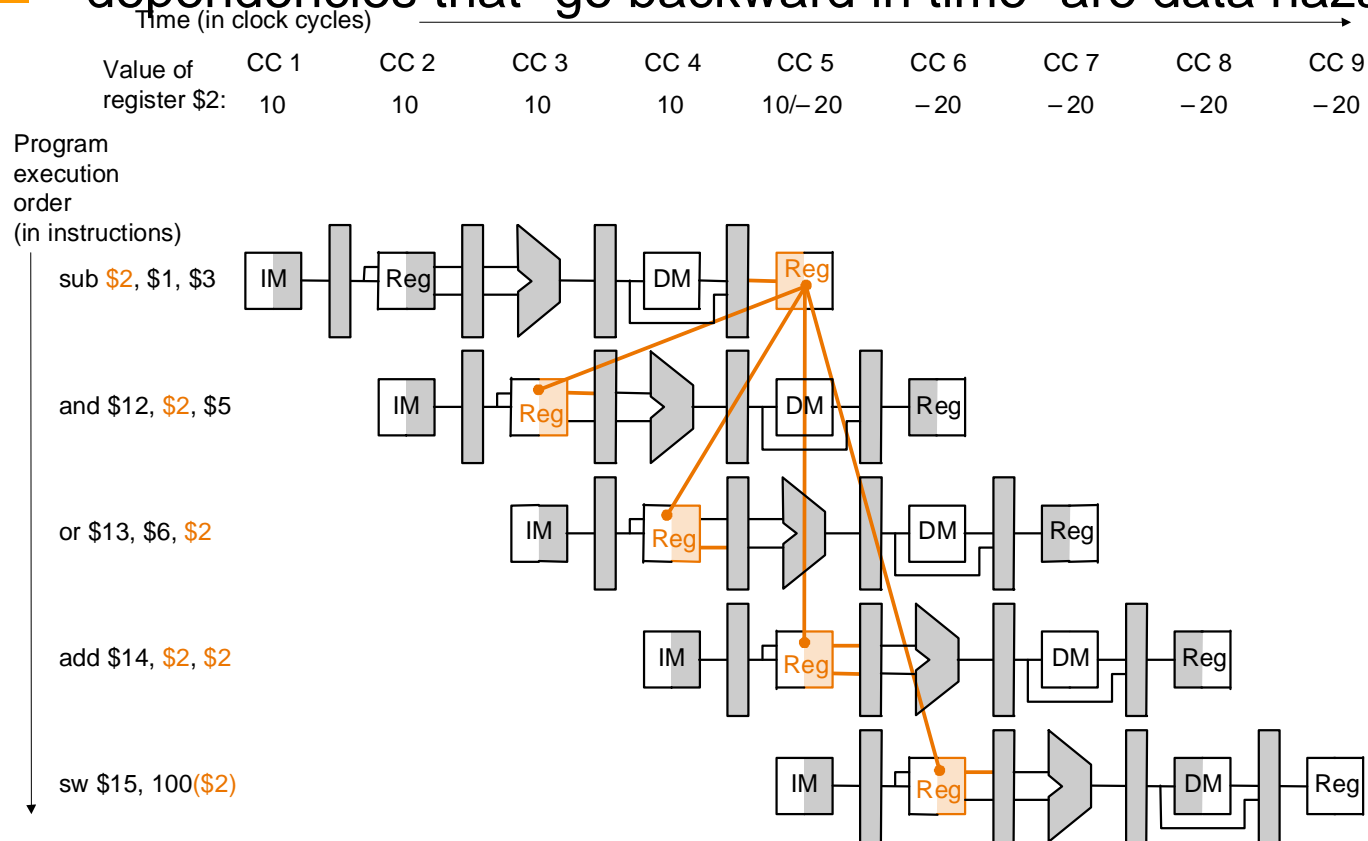
Instruction	Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	1	0	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X



Dependencies

□ Problem with starting next instruction before first is finished

■ dependencies that “go backward in time” are data hazards



Software Solution


- Have compiler guarantee no hazards
- Where do we insert the “nops” ?

```
sub    $2, $1, $3
and    $12, $2, $5
or     $13, $6, $2
add    $14, $2, $2
sw     $15, 100($2)
```

- Problem: this really slows us down!

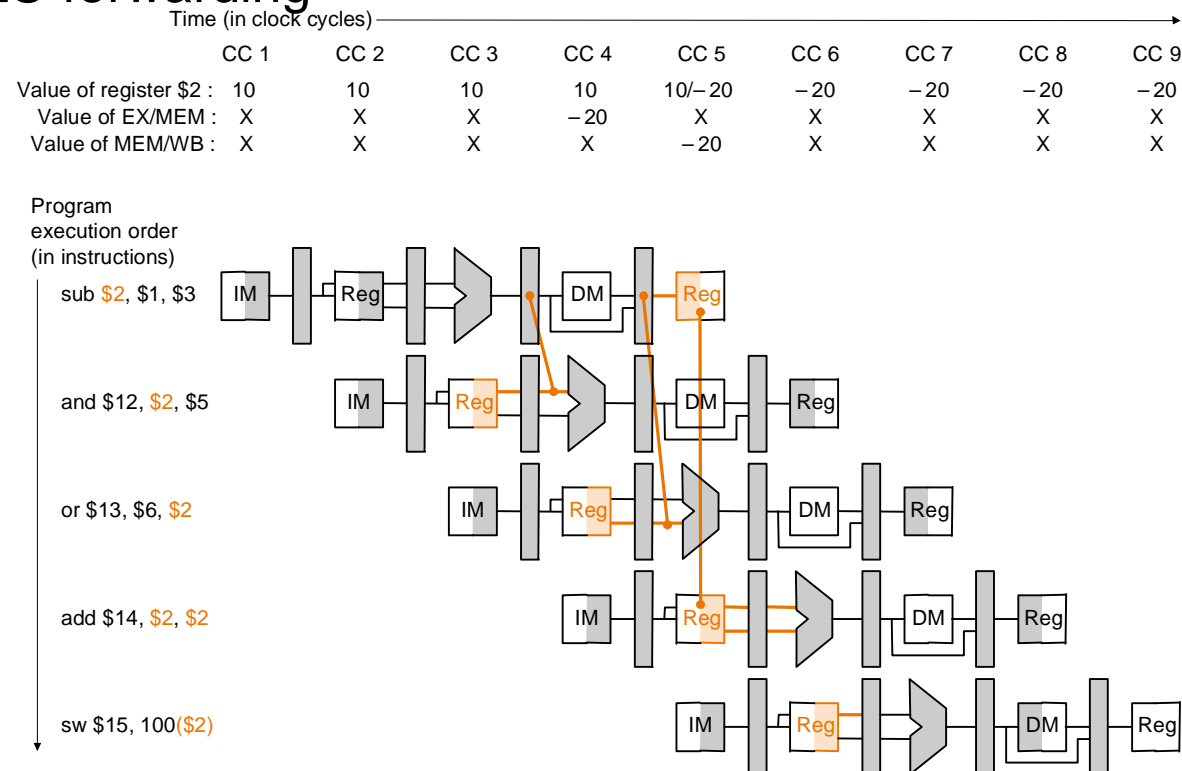
Or Pipeline Stalls

bubble

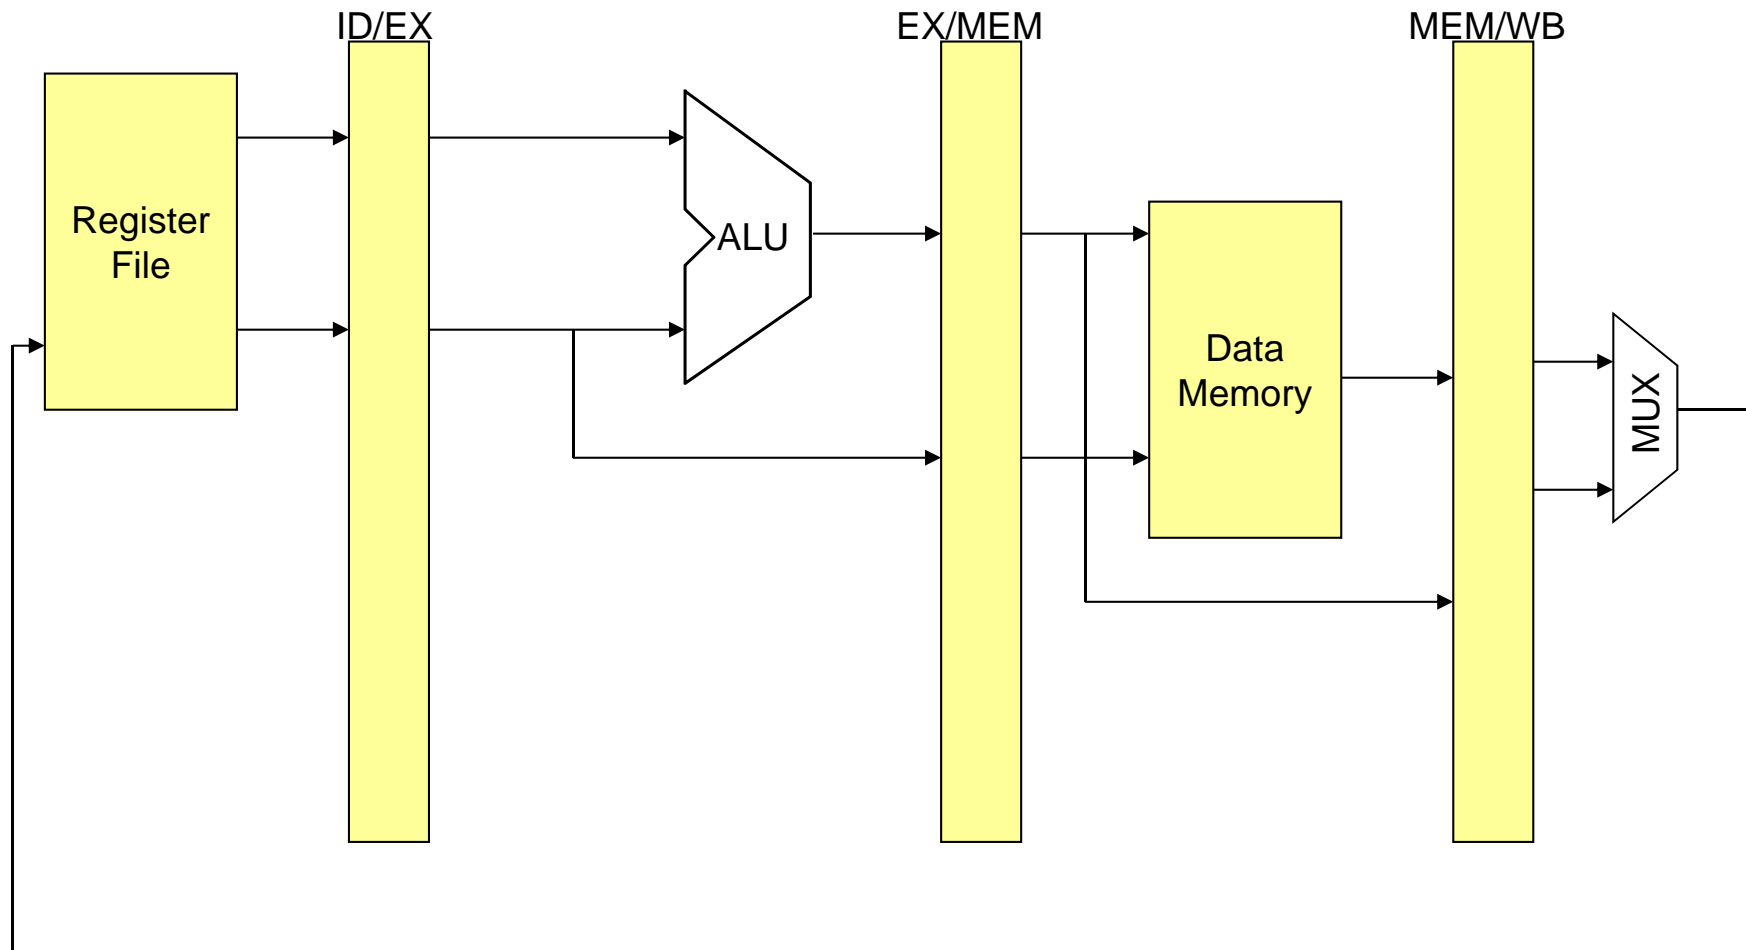


Forwarding

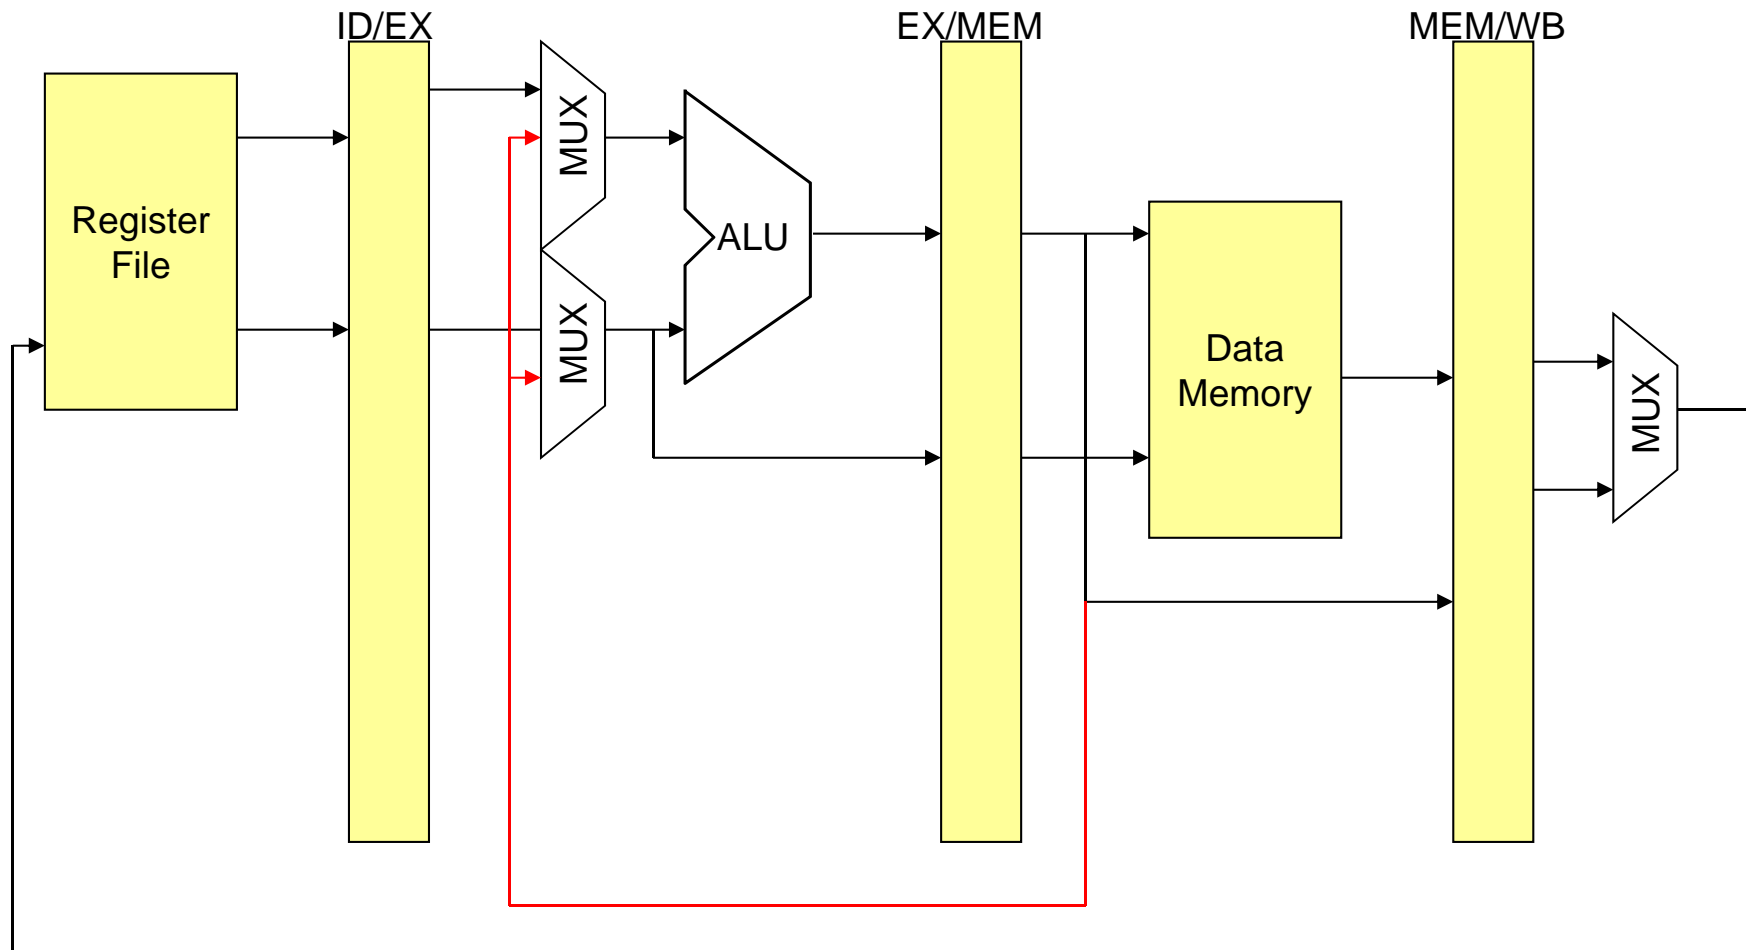
- Use temporary results, don't wait for them to be written
 - register file forwarding to handle read/write to same register
 - ALU forwarding



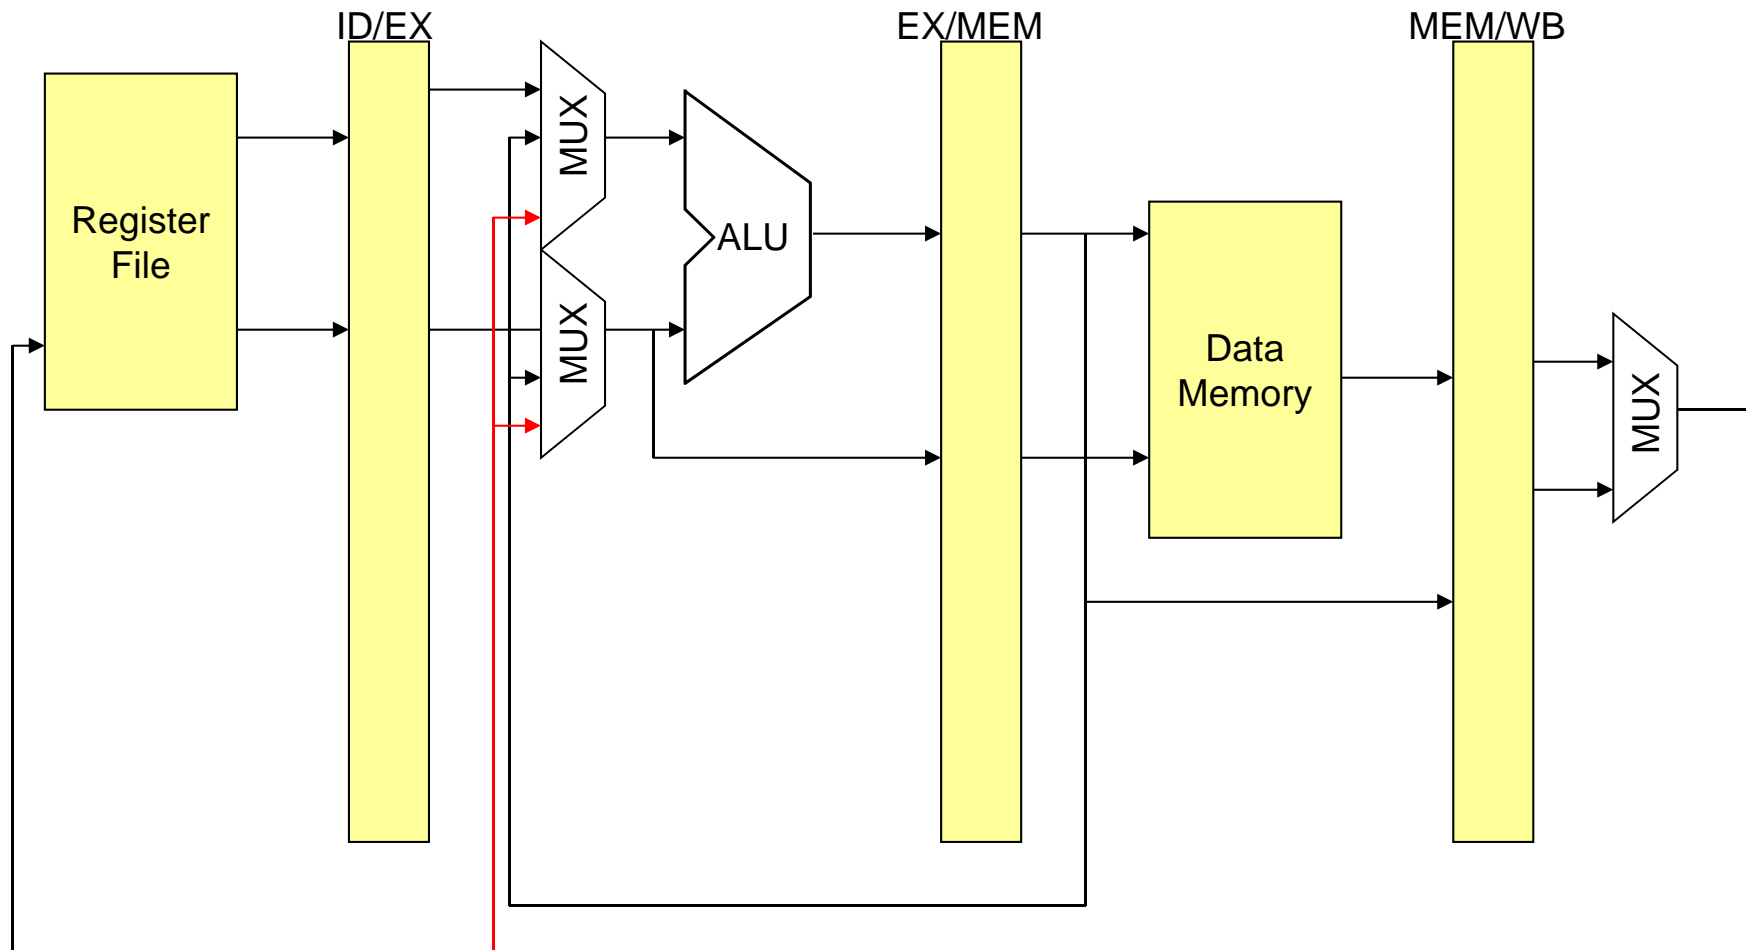
Forwarding (simplified)



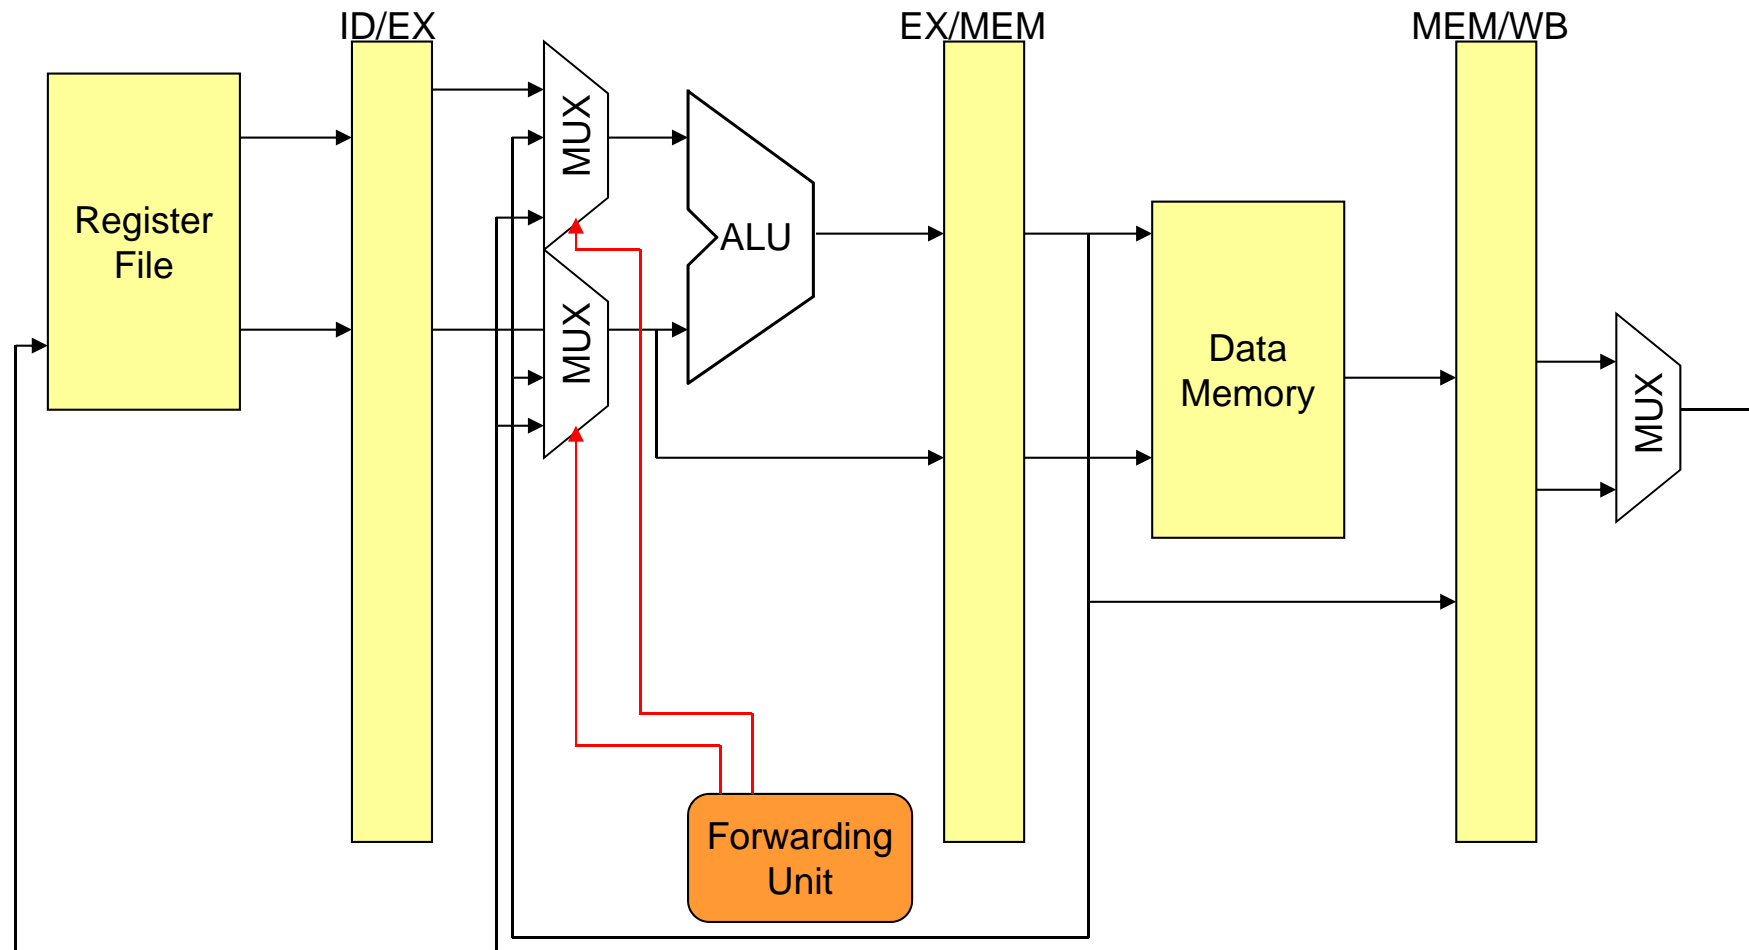
Forwarding (from EX/MEM)



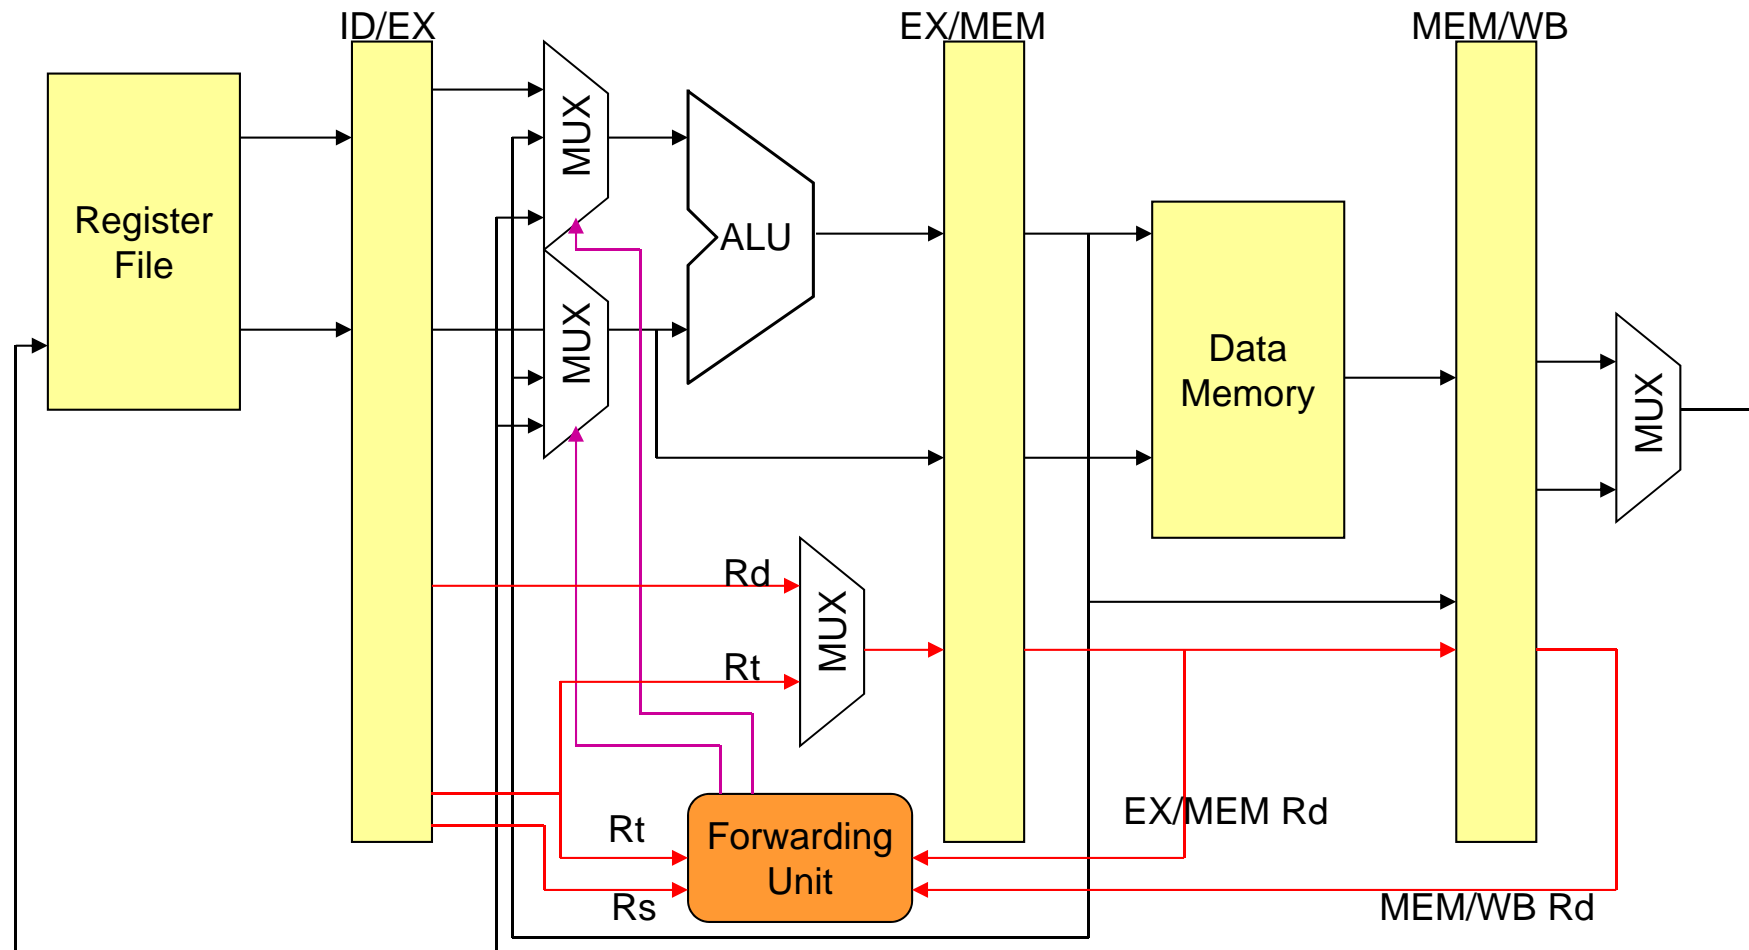
Forwarding (from MEM/WB)



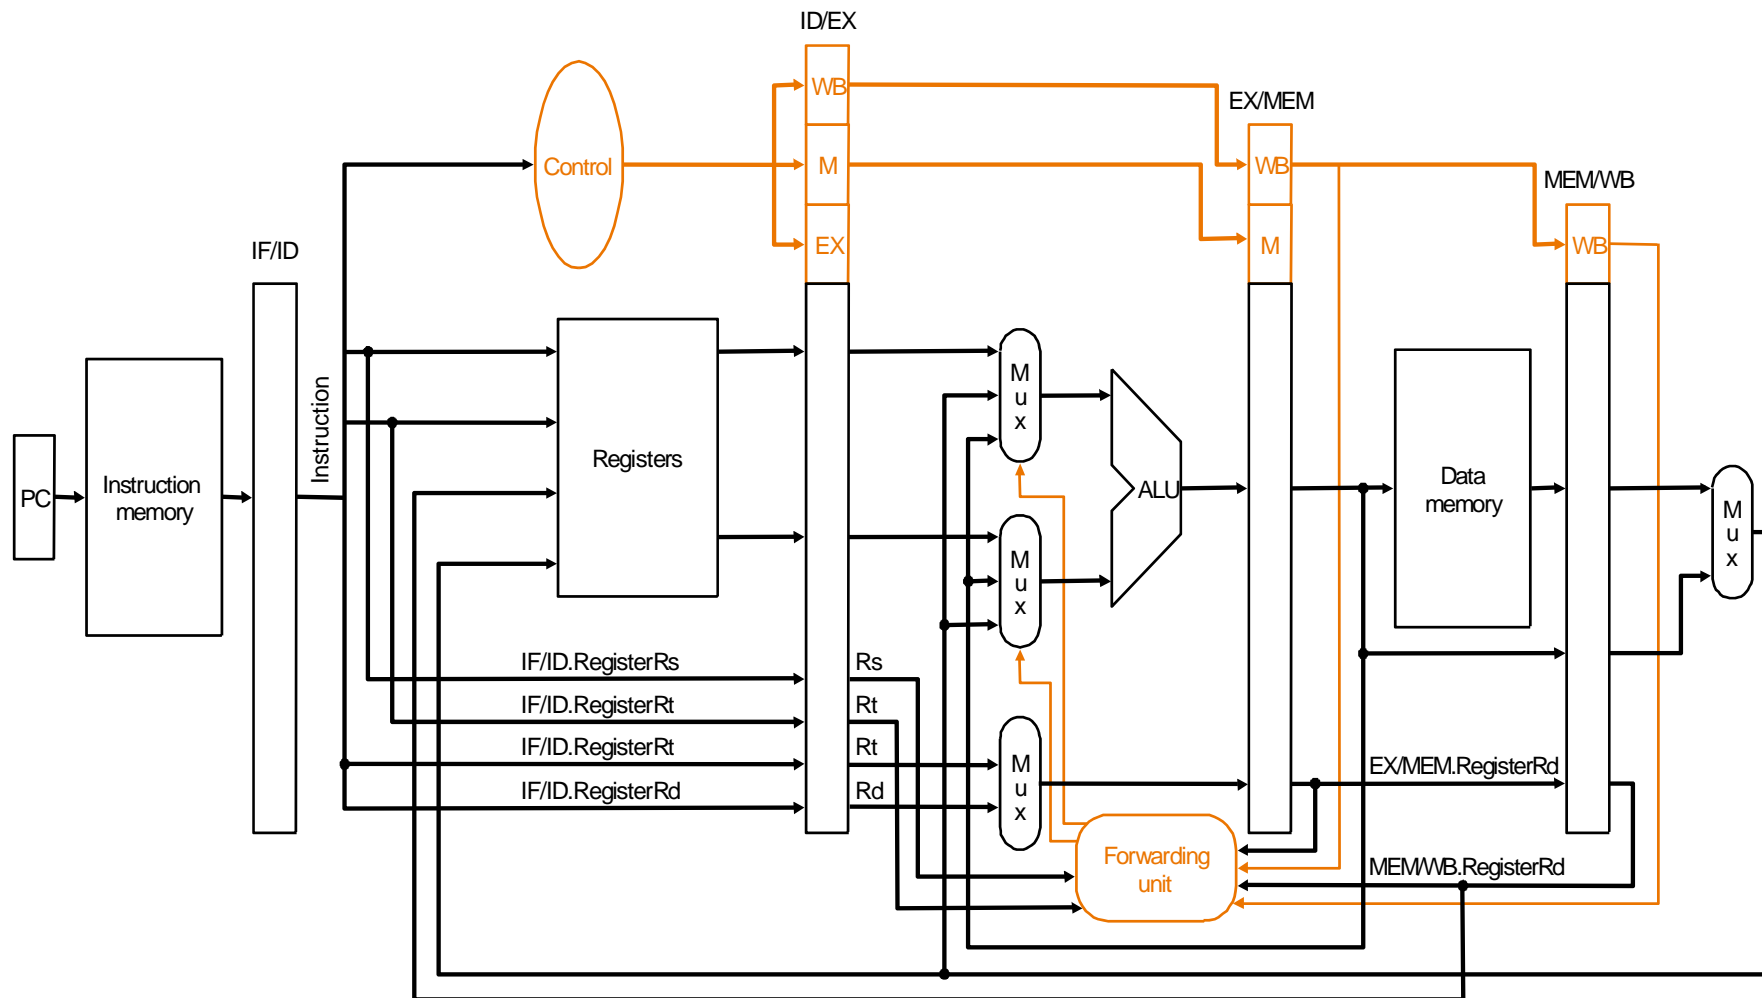
Forwarding (operand selection)



Forwarding (operand propagation)

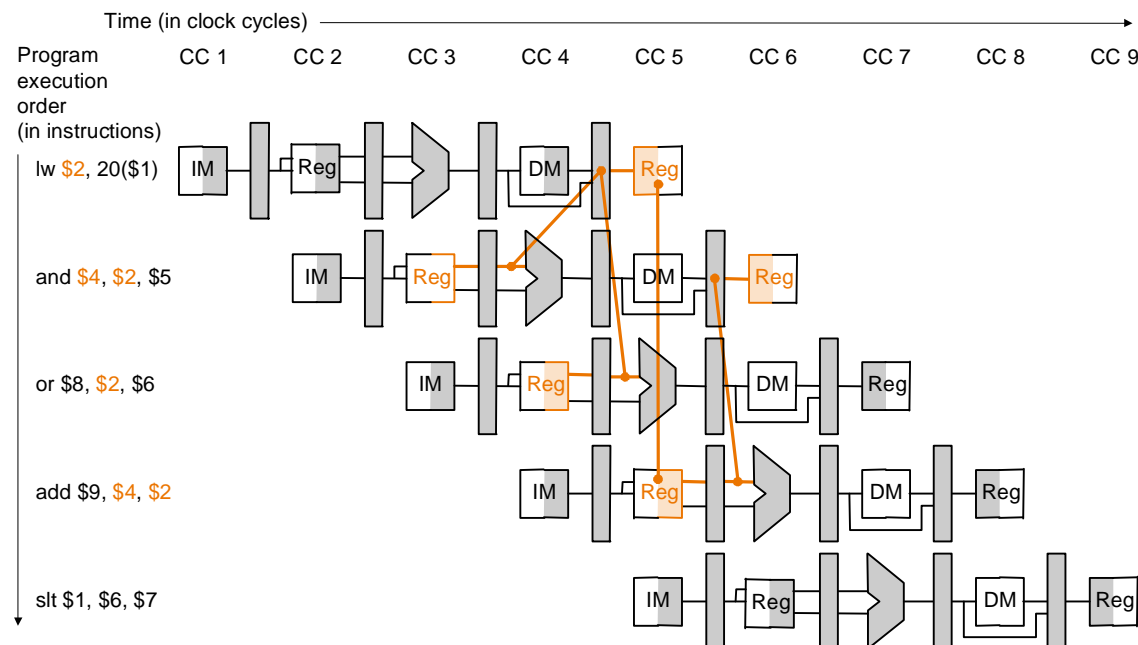


Forwarding



Can't always forward

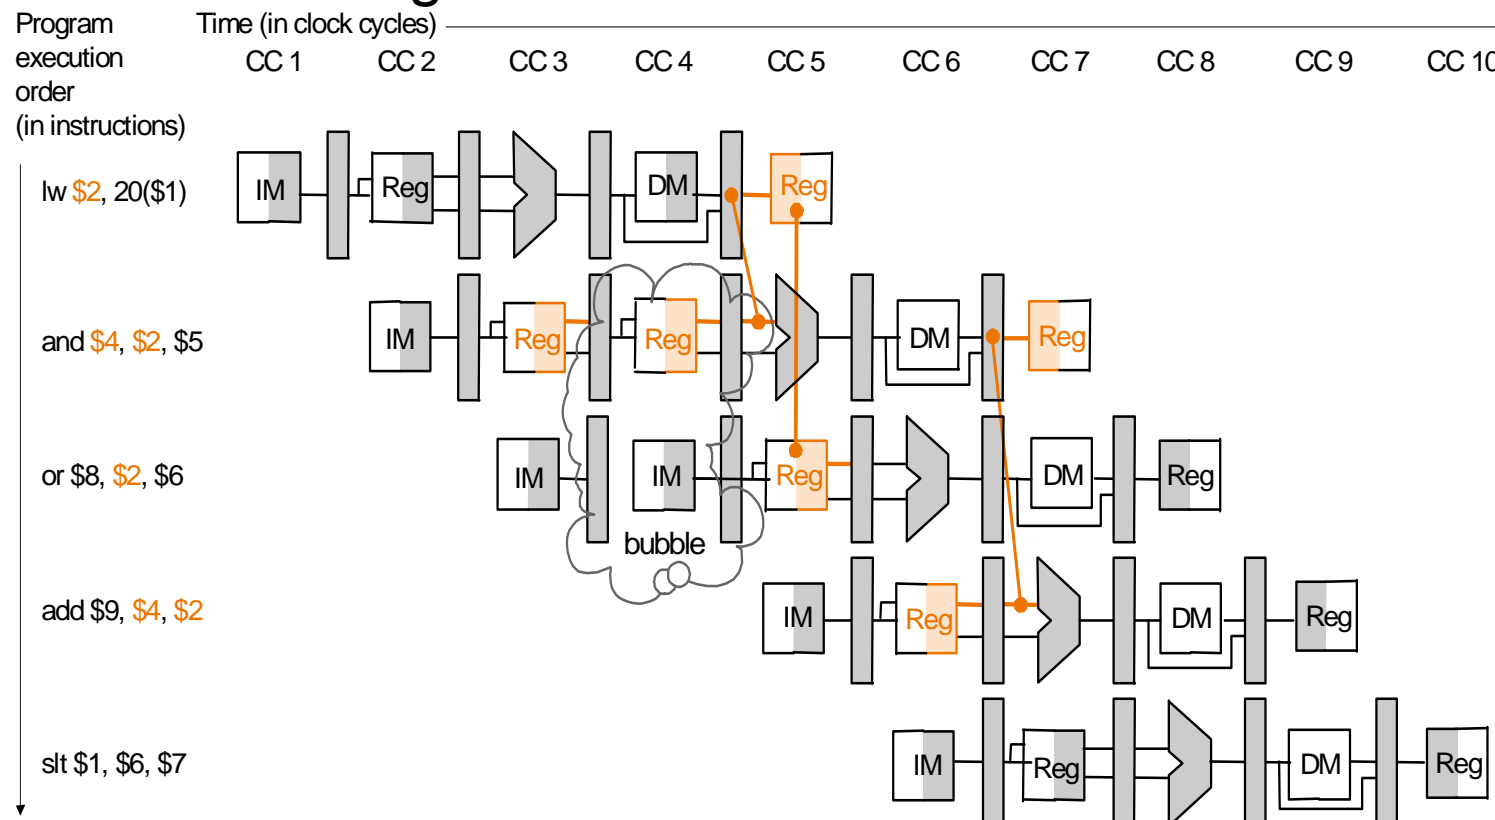
- Load word can still cause a hazard:
 - an instruction tries to read a register following a load instruction that writes to the same register.



- Thus, we need a hazard detection unit to “stall” the load instruction

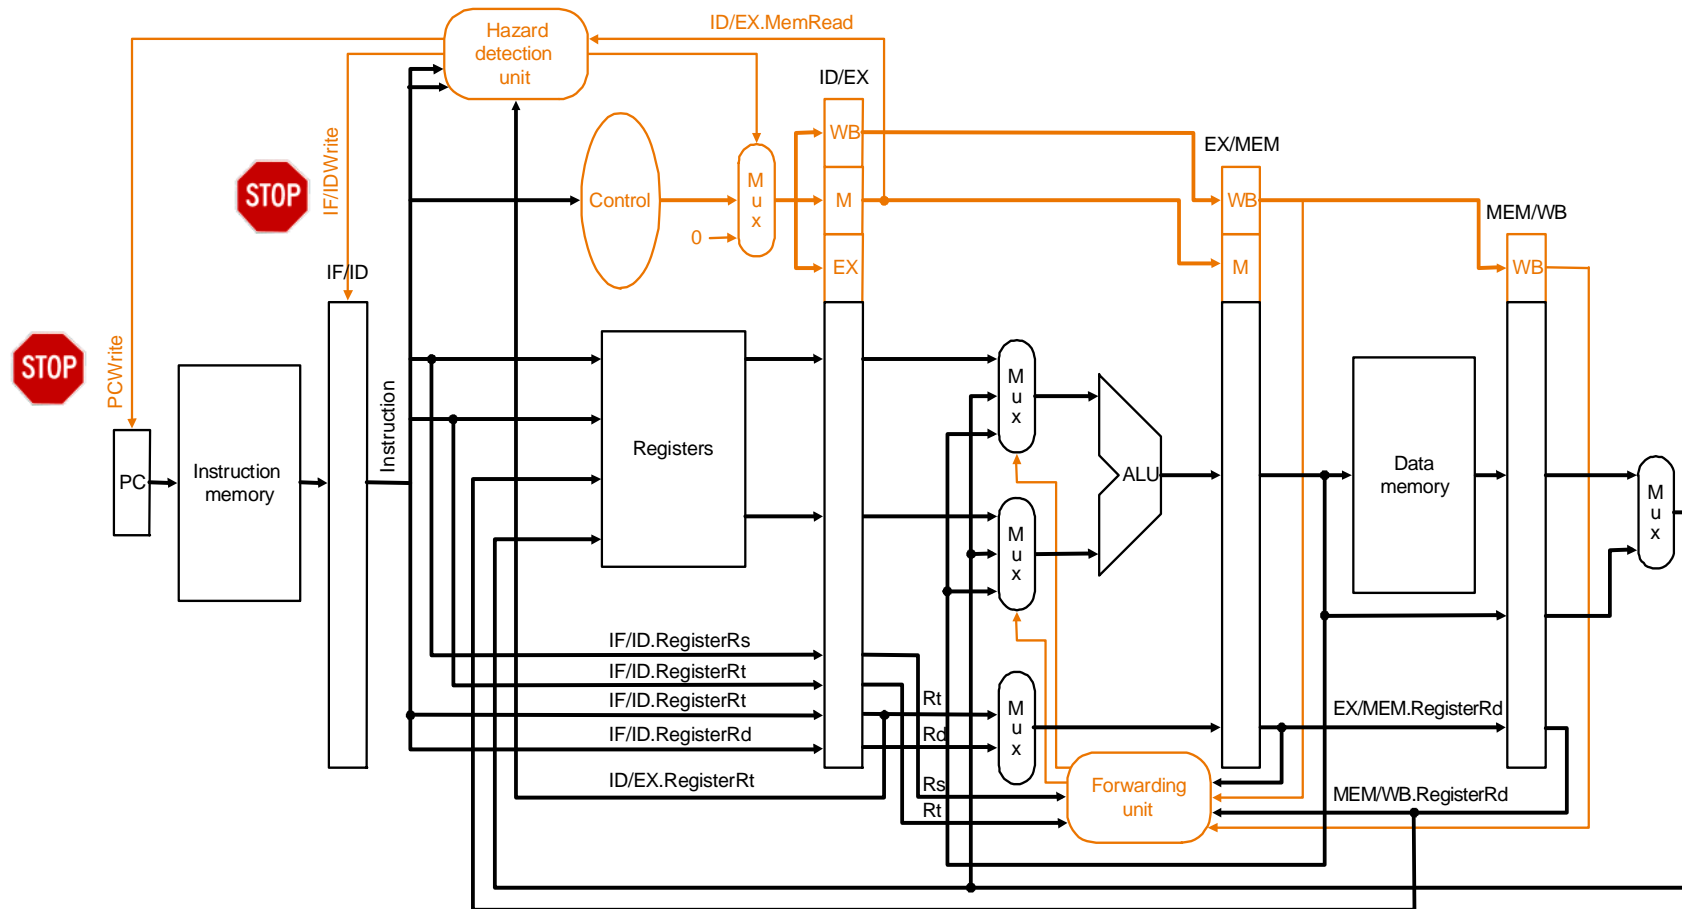
Stalling

- We can stall the pipeline by keeping an instruction in the same stage



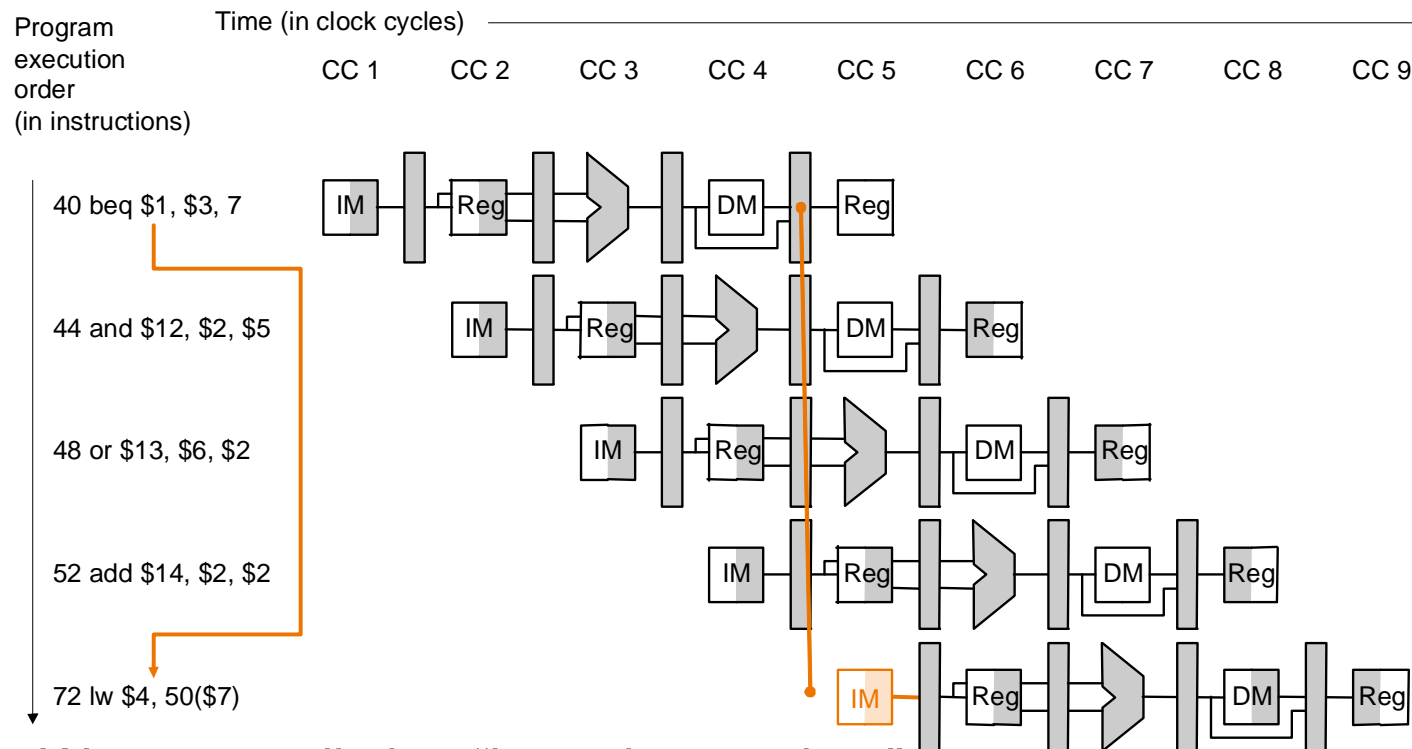
Hazard Detection Unit

- Stall by letting an instruction that won't write anything go forward
- Stall the pipeline if ID/EX is a load, and ($rt=IF/ID.rs$ or $rt=IF/ID.rt$)



Branch Hazards

- When we decide to branch, other instructions are in the pipeline!
- Assume: branch is not taken
- When assumption failed, flush 3 instructions




- We are predicting “branch not taken”
- need to add hardware for flushing instructions if we are wrong


Alleviate Branch Hazards

- ❑ Move branch compare to ID stage of the pipeline
- ❑ Add adder to calculate branch target in ID stage
- ❑ Add IF.flush signal that zeros the instruction (or squash) in IF/ID pipeline register
- ❑ Reduce penalty to 1 cycle

Flushing Instructions



Flushing Instructions (cycle N)




and \$12, \$2, \$5

```
beq $1, $3, L2
```



Flushing Instructions (cycle N+1)



Improving Performance

- Try and avoid stalls! E.g., reorder these instructions:

```
lw $t0, 0($t1)
lw $t2, 4($t1)
sw $t2, 0($t1)
sw $t0, 4($t1)
```

- Superscalar: start more than one instruction in the same cycle
- Most all processors are now pipelined and Superscalar

Dynamic Scheduling

- The hardware performs the “scheduling”
 - hardware tries to find instructions to execute
 - out of order execution is possible
 - speculative execution and dynamic branch prediction
- All modern processors are very complicated
 - DEC Alpha 21264: 9 stage pipeline, 6 instruction issue
 - PowerPC and Pentium: branch history table
 - Compiler technology important